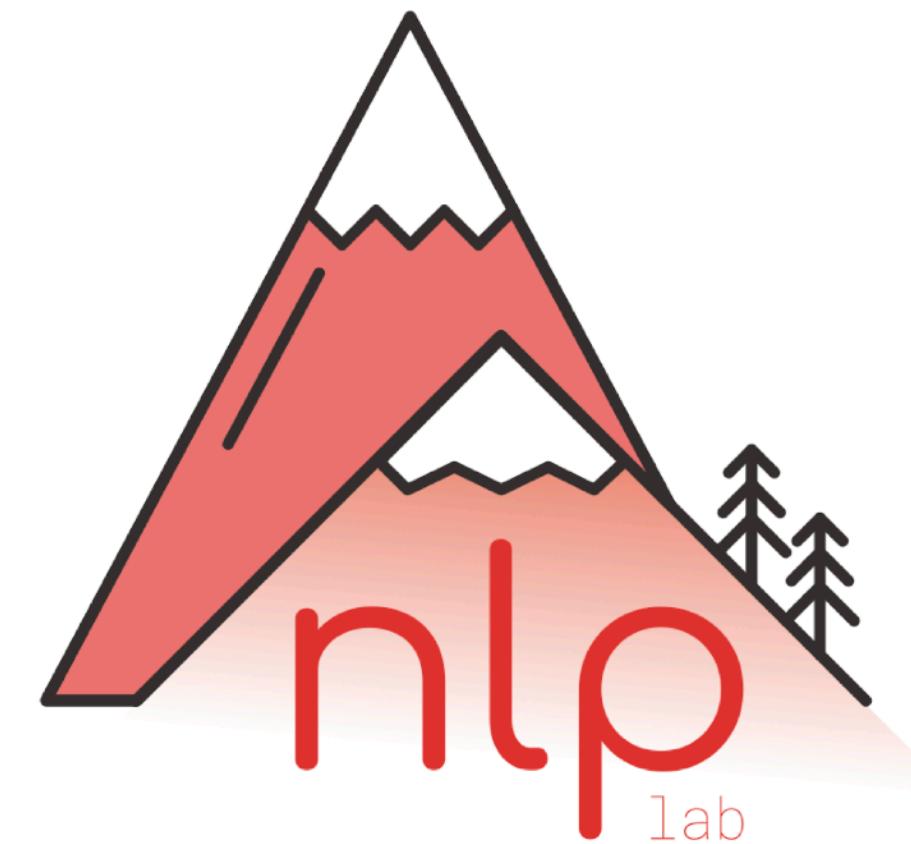


Classic Language Models

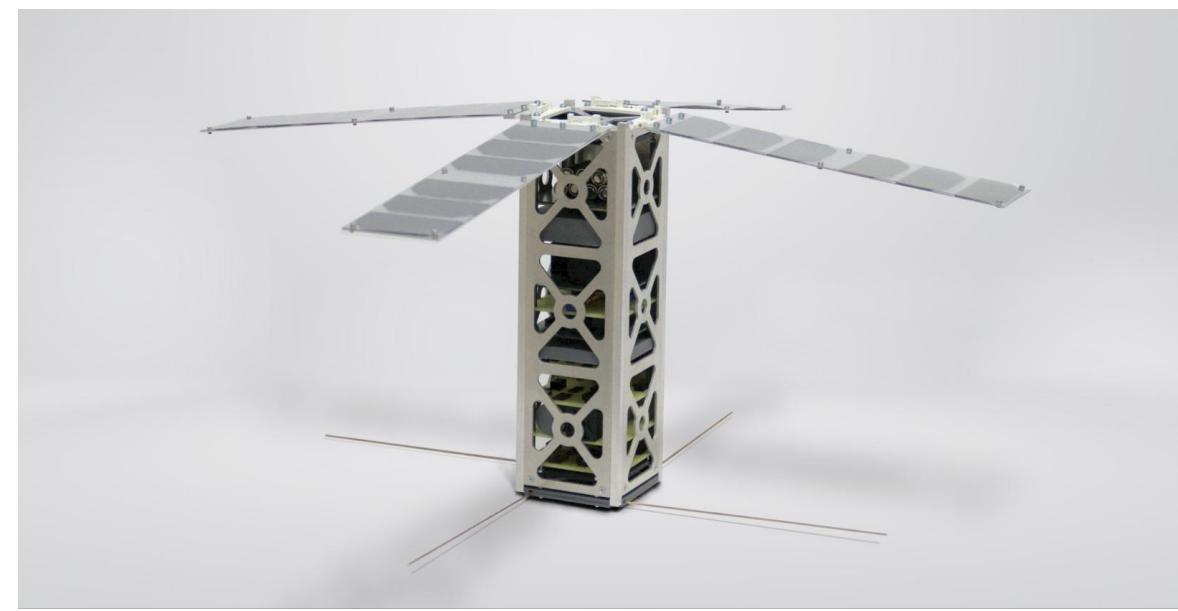
Antoine Bosselut



S/C: 3U CubeSat

Payloads: GNSS and CubeSatTOF

Objective: Measure chemical composition of the exosphere



Supported by the ESA FYS Design Booster 

Next launch: Summer 2027

Build space-grade onboard software

- Own a subsystem (e.g. power system)
- Interface with flight hardware
- Implement control & telemetry logic
- Write and review critical code



Scan and apply!

Or build the digital twin

- Simulate the spacecraft
- Inject failures
- Stress-test mission software

Workflow

Implementation → Unit Tests → Peer Review
→ Simulation → Flight



SAFE AI LAUSANNE

We are interested in making
advanced artificial intelligence
beneficial for humanity.





ICBD

IC BOOST DAY

7th edition

>> TALKS >> DISCUSSIONS >> CV PRESENTATION
>> NETWORKING LUNCH >> SPEED NETWORKING
>> MOCK CODING INTERVIEWS >> CLOSING APERITIF

>> TALKS

DISTRIBUTED ALGORITHMS

INTRODUCTION TO CRYPTOGRAPHY:
5D OVERVIEW OF THE DOMAIN

DATA ARCHITECTURE IN DEFENSE

DESIGNING FOR DISTRUST:
THE IMPOSSIBLE TRADE-OFFS
AND MORE...

>> DISCUSSIONS

ACADEMIA VS INDUSTRY

WOMEN IN IT

RESEARCH AND DEVELOPMENT

HOW TO TRAIN YOUR LLM

PHDS' EXPERIENCE

Wed. 11th of March

10am to 7pm in BC

For IC students in master or
third year bachelor

>> SCAN



or go to clic.epfl.ch/ICBD

Announcements

- **IC Boost Day, March 11th 2026:** Class cancelled
 - Video Lectures will be released early instead.
 - Watch on your own time (ideally by March 12th guest lecture)
- **Midterm:** Tentative Date is **April 15, 2026**
 - Time will be during class period: **11 AM - 1 PM**
 - Spread across multiple rooms on campus
 - If you have a **good** reason for not being available on these dates, send an e-mail to:
nlp-cs552-spring2026-ta-team@groupes.epfl.ch
 - **More details on midterm content in the coming weeks**

Today's Outline

- **Part 1:** Language models introduction, count-based language models, evaluating language models, smoothing
- **Part 2:** Fixed-context Language Models, Recurrent Neural Networks 1

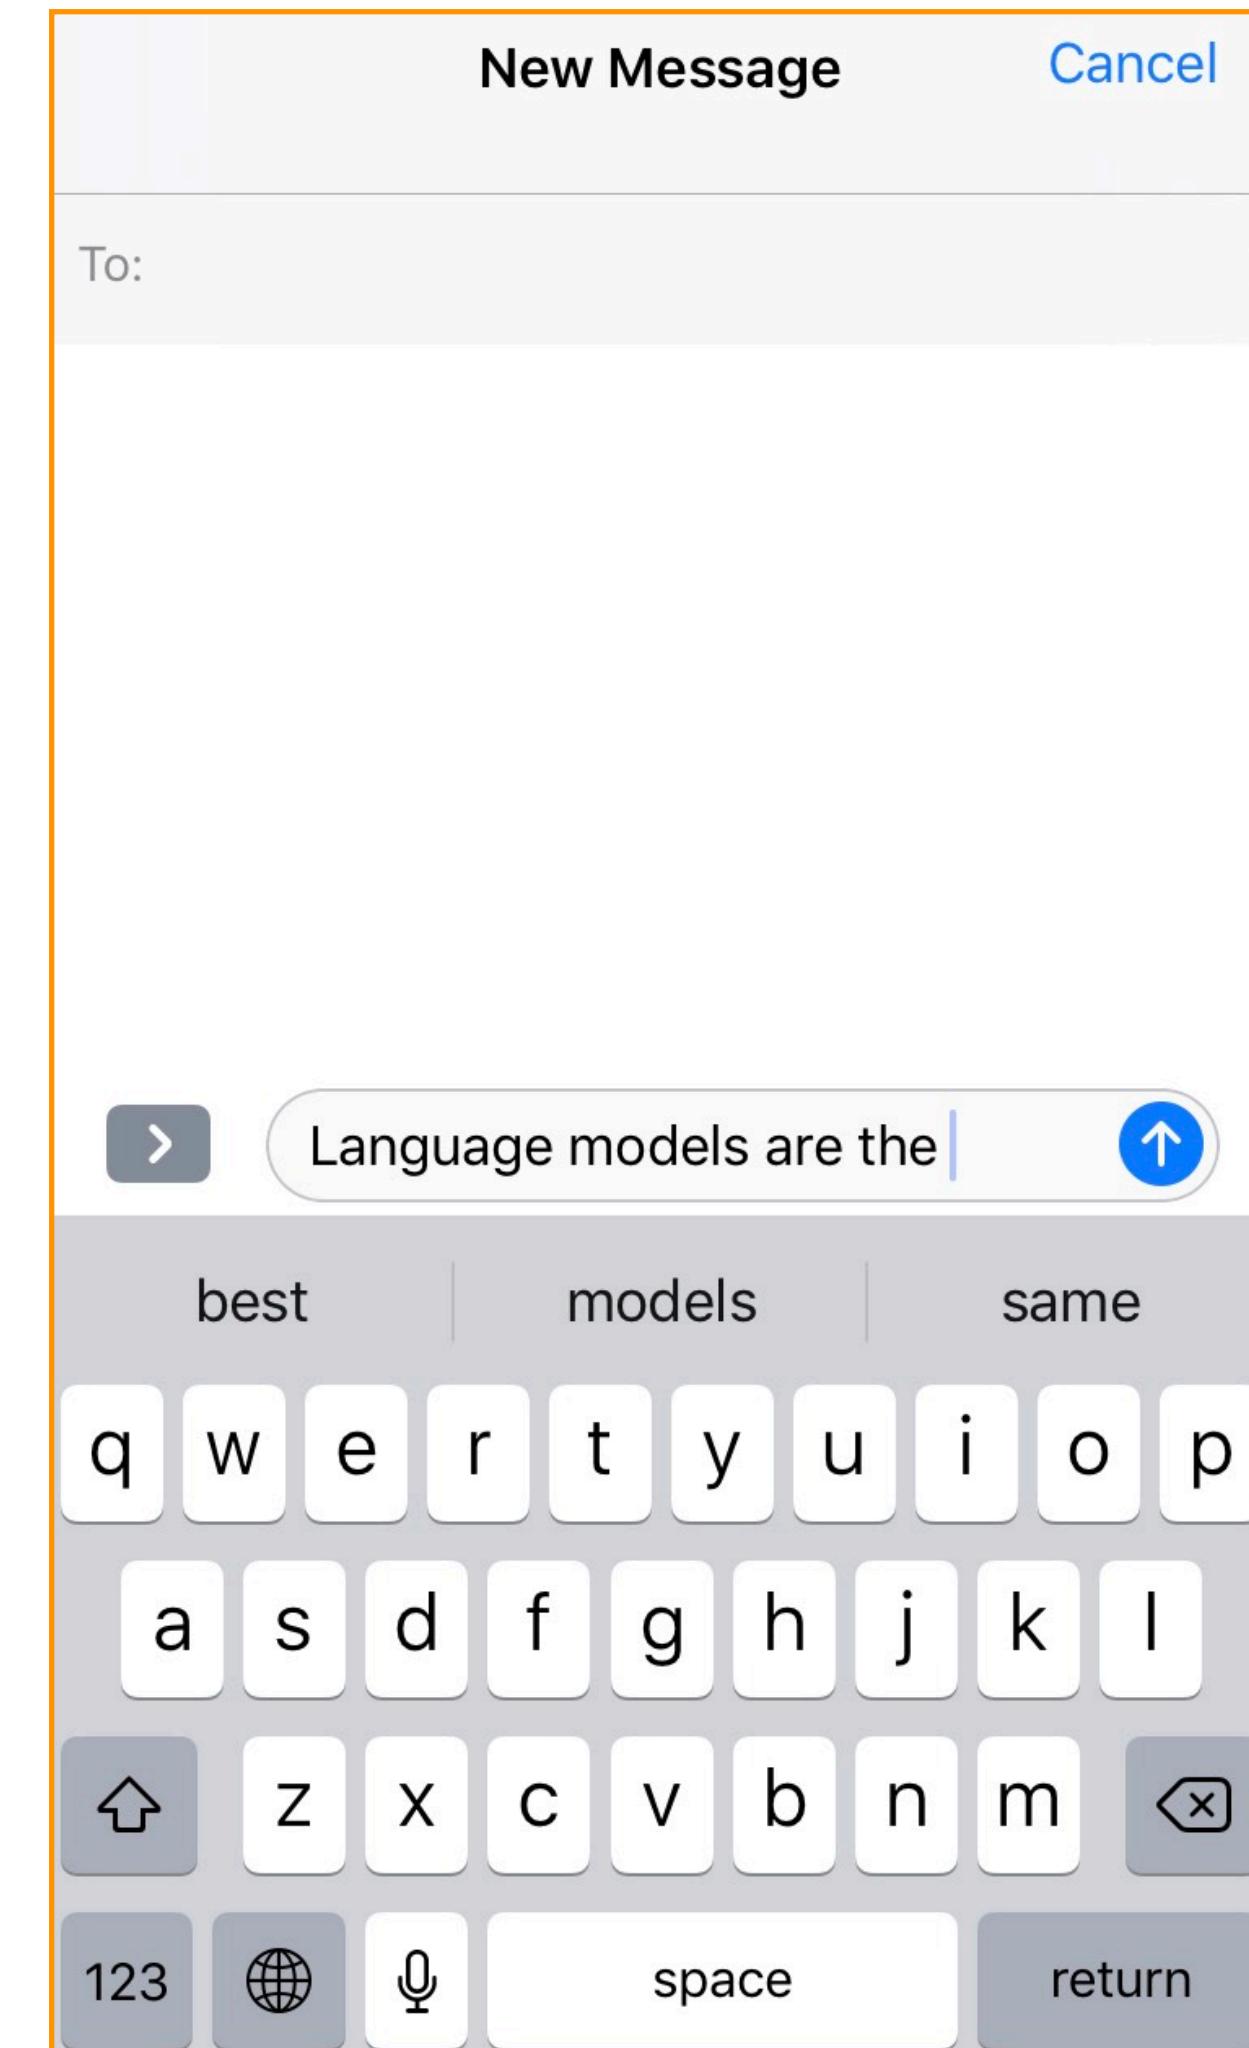
What is a language model?

- A language model is a **probabilistic model of a sequence of tokens**
 - How likely is a given phrase/sentence/paragraph/document?
- A sequence is modelled as a joint distribution of tokens $w_1, w_2, w_3, \dots, w_n$:

$$P(w_1, w_2, w_3, \dots, w_n)$$

- Language models estimate this probability!

Language models are ubiquitous



Applications of LMs

- Predicting words is important in many situations

Given an empty input

- **Classical:** Machine translation, text generation

$$P(\text{a smooth finish}) > P(\text{a flat finish})$$

- **Classical:** Speech recognition/Spell checking

$$P(\text{high school principal}) > P(\text{high school principle})$$

- **Modern:** Information extraction, Question answering, Classification, etc.

What is a language model?

- A language model is a **probabilistic model of a sequence of tokens**
 - How likely is a given phrase/sentence/paragraph/document?
- A sequence is modelled as a joint distribution of tokens $w_1, w_2, w_3, \dots, w_n$:

$$P(w_1, w_2, w_3, \dots, w_n)$$

- Language models estimate this probability!

Chain rule

- How should we compute this joint probability over words in a sequence?

$$\begin{aligned} P(w_1, w_2, w_3, \dots, w_n) &= P(w_1) \times P(w_2 | w_1) \times P(w_3 | w_1, w_2) \\ &\quad \times \dots \times P(w_n | w_1, w_2, \dots, w_{n-1}) \end{aligned}$$

$$P(w_1, w_2, w_3, \dots, w_n) = \prod_i P(w_i | \boxed{w_1, w_2, \dots, w_{i-1}}) \rightarrow \text{“prefix”}$$

Chain rule

- How should we compute this joint probability over words in a sequence?

$$\begin{aligned} P(w_1, w_2, w_3, \dots, w_n) &= P(w_1) \times P(w_2 | w_1) \times P(w_3 | w_1, w_2) \\ &\quad \times \dots \times P(w_n | w_1, w_2, \dots, w_{n-1}) \end{aligned}$$

$$P(w_1, w_2, w_3, \dots, w_n) = \prod_i P(w_i | w_1, w_2, \dots, w_{i-1}) \rightarrow \text{“prefix”}$$

Sentence: “the cat sat on the mat”

$$\begin{aligned} P(\text{the cat sat on the mat}) &= P(\text{the}) * P(\text{cat}|\text{the}) * P(\text{sat}|\text{the cat}) \\ &\quad * P(\text{on}|\text{the cat sat}) * P(\text{the}|\text{the cat sat on}) \\ &\quad * P(\text{mat}|\text{the cat sat on the}) \end{aligned}$$

Count-based Modeling

Maximum likelihood estimate (MLE)

$$P(\text{sat}|\text{the cat}) = \frac{\text{count}(\text{the cat sat})}{\text{count}(\text{the cat})}$$

$$P(\text{on}|\text{the cat sat}) = \frac{\text{count}(\text{the cat sat on})}{\text{count}(\text{the cat sat})}$$

⋮

Note on Notation:

Word type: unique word in our vocabulary V

Token: instance of a word type in our corpus

- ▶ Estimate probabilities by counting occurrences of sequences in a **corpus** of text

Question

**What's going to be a problem with
counting sequences for any prefix?**

Count-based Modeling

Maximum likelihood estimate (MLE)

$$P(\text{sat}|\text{the cat}) = \frac{\text{count}(\text{the cat sat})}{\text{count}(\text{the cat})}$$

$$P(\text{on}|\text{the cat sat}) = \frac{\text{count}(\text{the cat sat on})}{\text{count}(\text{the cat sat})}$$

:

Note on Notation:

Word type: unique word in our vocabulary V

Token: instance of a word type in our corpus

- ▶ Estimate probabilities by counting occurrences of sequences in a **corpus** of text
- ▶ With a vocabulary of size V ,
 - number of sequences of length $n = V^n$
- ▶ Typical vocabulary ≈ 40000 words
 - even sentences of length $<= 11$ results in more than $4 * 10^{50}$ sequences!
(more than the number of atoms in the earth)

Very sparse representation aswell

Markov assumption

- Use only the recent past to predict the next word
- **Effect:** reduce the number of estimated parameters in exchange for modeling capacity
- 1st order Markov

$$P(\text{mat}|\text{the cat sat on the}) \approx P(\text{mat}|\text{the})$$

Bigram model

- 2nd order Markov

$$P(\text{mat}|\text{the cat sat on the}) \approx P(\text{mat}|\text{on the})$$

Trigram model

k^{th} order Markov

- Consider only the last k words for context

$$P(w_i | w_1 w_2 \dots w_{i-1}) \approx P(w_i | w_{i-k} \dots w_{i-1})$$

which implies the probability of a sequence is:

$$P(w_1 w_2 \dots w_n) \approx \prod_i P(w_i | w_{i-k} \dots w_{i-1})$$

(ignore $w_j \quad \forall j < 0$)

n-gram models

Unigram

$$P(w_1, w_2, \dots, w_n) = \prod_{i=1}^n P(w_i)$$

Bigram

$$P(w_1, w_2, \dots, w_n) = \prod_{i=1}^n P(w_i | w_{i-1})$$

and Trigram, 4-gram, and so on.

Larger n leads to more accurate and better language models
(but also higher costs)*

**Caveat: Assuming infinite data!*

Examples

Unigram

release millions See ABC accurate President of Donald Will
cheat them a CNN megynkelly experience @ these word
out- the

Bigram

Thank you believe that @ ABC news, Mississippi tonight
and the false editorial I think the great people Bill
Clinton . "

Trigram

We are going to MAKE AMERICA GREAT AGAIN!
#MakeAmericaGreatAgain <https://t.co/DjkdAzT3WV>

$$\arg \max_{(w_1, w_2, \dots, w_n)} P(w_1, w_2, \dots, w_n) = \arg \max_{(w_1, w_2, \dots, w_n)} \prod_{i=1}^n P(w_i | w_{i-k}, \dots, w_{i-1})$$

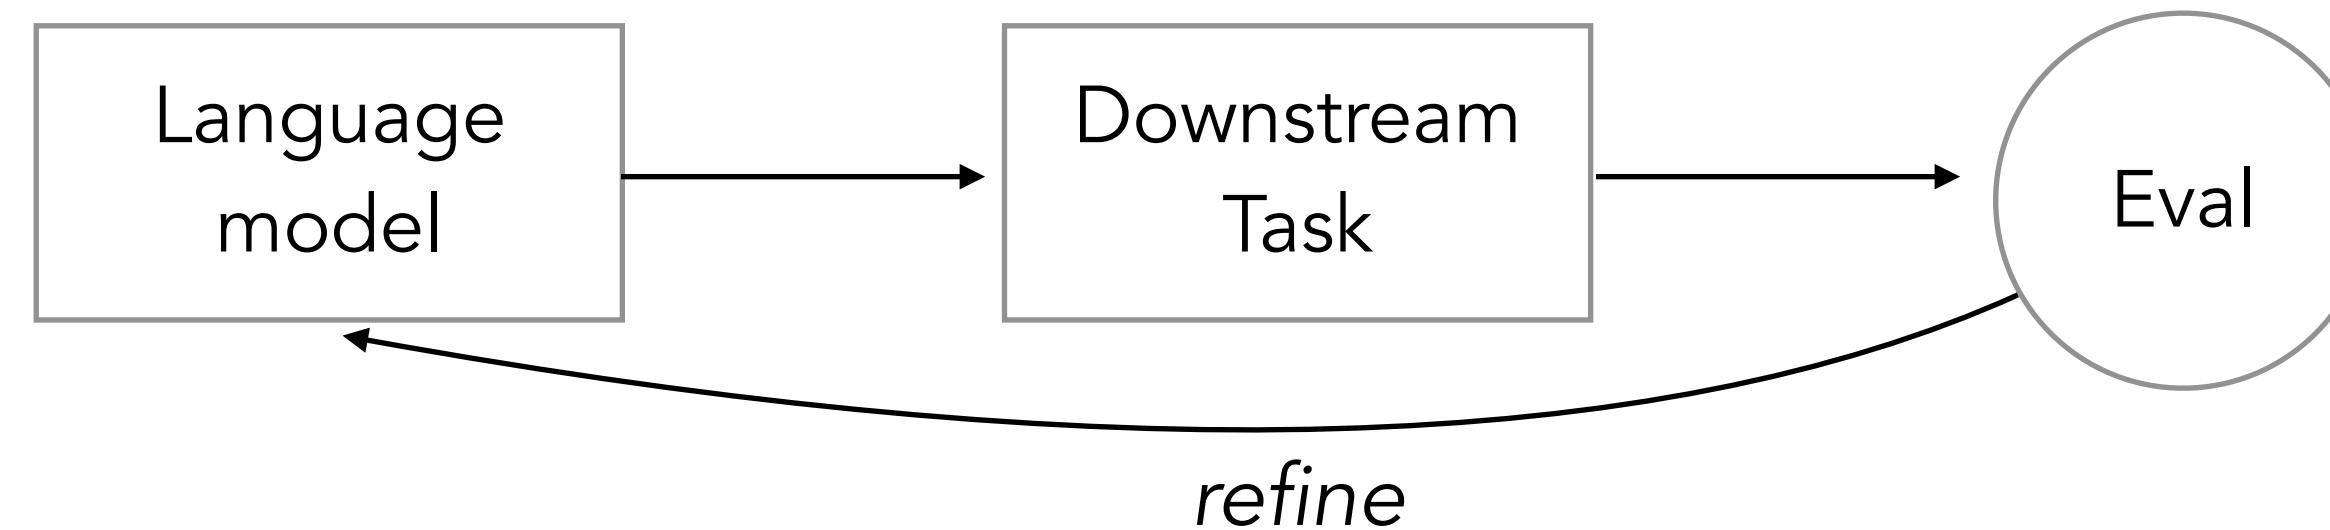
Evaluating language models

- A good language model should assign **higher probability** to typical, grammatically correct sentences
- Research process:
 - **Train** parameters on a suitable training corpus
 - ▶ Assumption: observed sentences are probably good sentences
 - **Test** on *different, unseen* corpus
 - ▶ **Training on any part of test set not acceptable! Why?**
 - **Evaluation metric**

Question

What evaluation metric should we use?

Extrinsic evaluation



- Train LM —> apply to task —> observe accuracy
- Directly optimized for downstream tasks
 - higher task accuracy —> better model
- Expensive, time consuming
- Hard to optimize downstream objective (indirect feedback)

Perplexity (ppl)

- Measure of how well a probability distribution (or LM) **predicts** a sample
- For a corpus S with sentences S^1, S^2, \dots, S^m

Negative log likelihood of the ‘ground truth’ sentences

$$\text{ppl}(S) = 2^x \quad \text{where} \quad x = -\frac{1}{m} \sum_{i=1}^m \log_2 P(S^i) \longrightarrow \text{Cross-Entropy}$$

- Unigram model: $x = -\frac{1}{W} \sum_{i=1}^m \sum_{j=1}^n \log_2 P(w_j^i)$ (since $P(S) = \prod_j P(w_j)$)

where W is the total number of words in test corpus

- Minimizing perplexity is the same as maximizing probability of corpus $P(S^1 S^2 \dots S^m)$

Intuition on perplexity

- If our n-gram model (with vocabulary V) has following probability:

$$P(w_i|w_{i-n}, \dots w_{i-1}) = \frac{1}{|V|} \quad \forall w_i$$

- What is the perplexity of the test corpus?

$$ppl = 2^{-\frac{1}{W} W \times \log_2(\frac{1}{|V|})} = |V|$$

$$\begin{aligned} ppl(S) &= 2^x \text{ where} \\ x &= -\frac{1}{m} \sum_{i=1}^m \log_2 P(S^i) \end{aligned}$$

- Any word is equally likely at next step!
- **Intuition:** Perplexity measures a model's uncertainty about the next word

Low perplexity means that the model assigns high prob to the correct next word.

Intuition on perplexity

- Perplexity is the exponentiated token-level negative log likelihood
 - **Why logs?**
- **Theory:** perplexity measured using a base of 2
 - **Practice:** doesn't matter what the base is, e used often
 - **Exercise:** derive why!

$$\text{ppl}(S) = 2^x \text{ where } x = -\frac{1}{m} \sum_{i=1}^n \log_2 P(S^i)$$

Comprehension Question:

How good is a language model with $\text{ppl} > 1$?

It is worse than predicting a random word every time

Perplexity as a metric

| Pros | Cons |
|----------------------------------------------------------|-----------------------------------------------|
| Easy to compute | Requires domain match between train and test |
| standardized | might not correspond to end task optimization |
| directly useful, easy to use to correct sentences | log 0 undefined |
| nice theoretical interpretation - matching distributions | can be 'cheated' by predicting common tokens |
| | size of test set matters |
| | can be sensitive to low prob tokens/sentences |

Question

What are major shortcomings of n-gram language models?

Exponential Decay of Counts

- What happens when we scale to longer contexts?

$P(w|to)$ *to* occurs 1M times in corpus

$P(w|go\ to)$ *go to* occurs 50,000 times in corpus

$P(w|to\ go\ to)$ *go to* occurs 1500 times in corpus

$P(w|want\ to\ go\ to)$ *want to go to:* only 100 occurrences

Exponential Decay of Counts

- What happens when we scale to longer contexts?

$P(w|to)$ *to* occurs 1M times in corpus

$P(w|go\ to)$ *go to* occurs 50,000 times in corpus

$P(w|to\ go\ to)$ *go to* occurs 1500 times in corpus

$P(w|want\ to\ go\ to)$ *want to go to:* only 100 occurrences

- Probability counts get very sparse, and we often want information from 5+ words away

Not all n-grams observed in training data

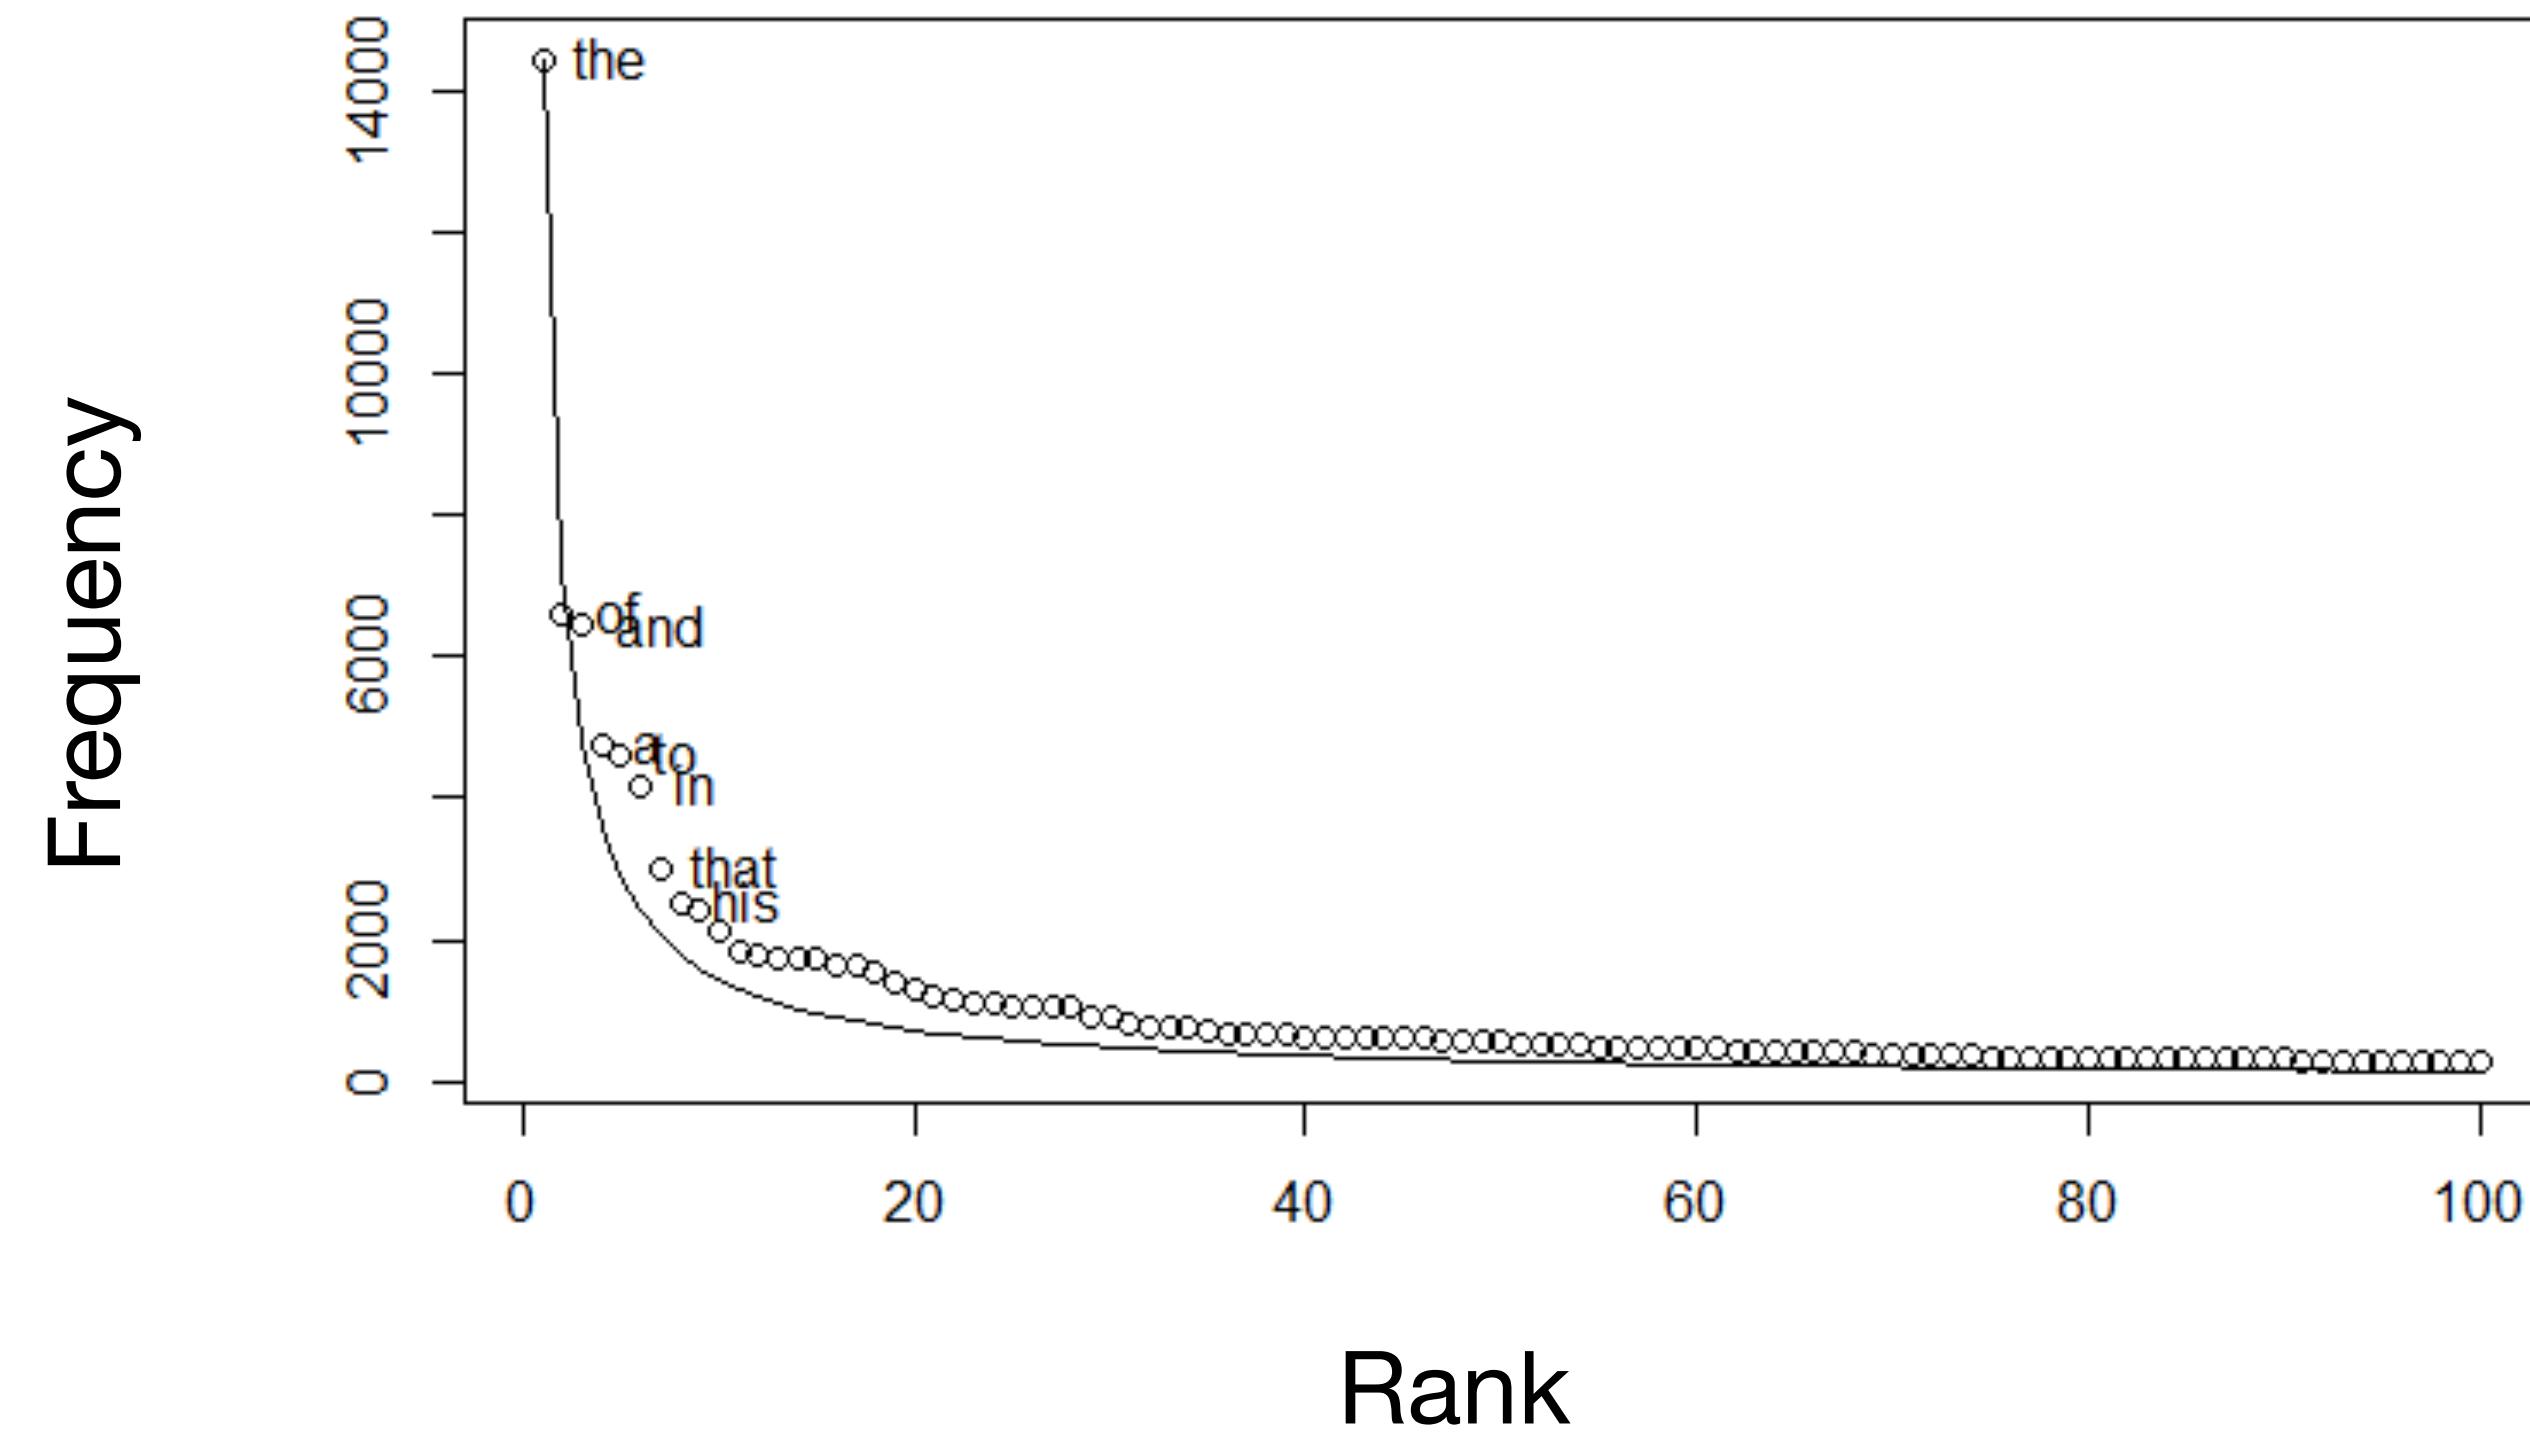
- Test corpus might have some that have zero probability under our model
 - Training set: *Google news*
 - Test set: *Shakespeare*
- $P(\text{affray} \mid \text{voice doth us}) = 0 \rightarrow P(\text{test corpus}) = 0$
- **Undefined perplexity**

Fun Fact!

Shakespeare as corpus

- $N=884,647$ tokens, $V=29,066$
- Shakespeare produced 300,000 bigram types out of $V^2= 844$ million possible bigrams.
 - So 99.96% of the possible bigrams were never seen (have zero entries in the table)

Sparsity in language



$$freq \propto \frac{1}{rank}$$

Zipf's Law

- Long tail of infrequent words
- Most finite-size corpora will have this problem.

Question

How might we fix this ?

Smoothing

- Handle sparsity by making sure all probabilities are non-zero in our model
 - **Additive**: Add a small amount to all probabilities
 - **Discounting**: Redistribute probability mass from observed n-grams to unobserved ones
 - **Back-off**: Use lower order n-grams if higher ones are too sparse
 - **Interpolation**: Use a combination of different granularities of n-grams

Smoothing intuition

When we have sparse statistics:

$P(w \mid \text{denied the})$

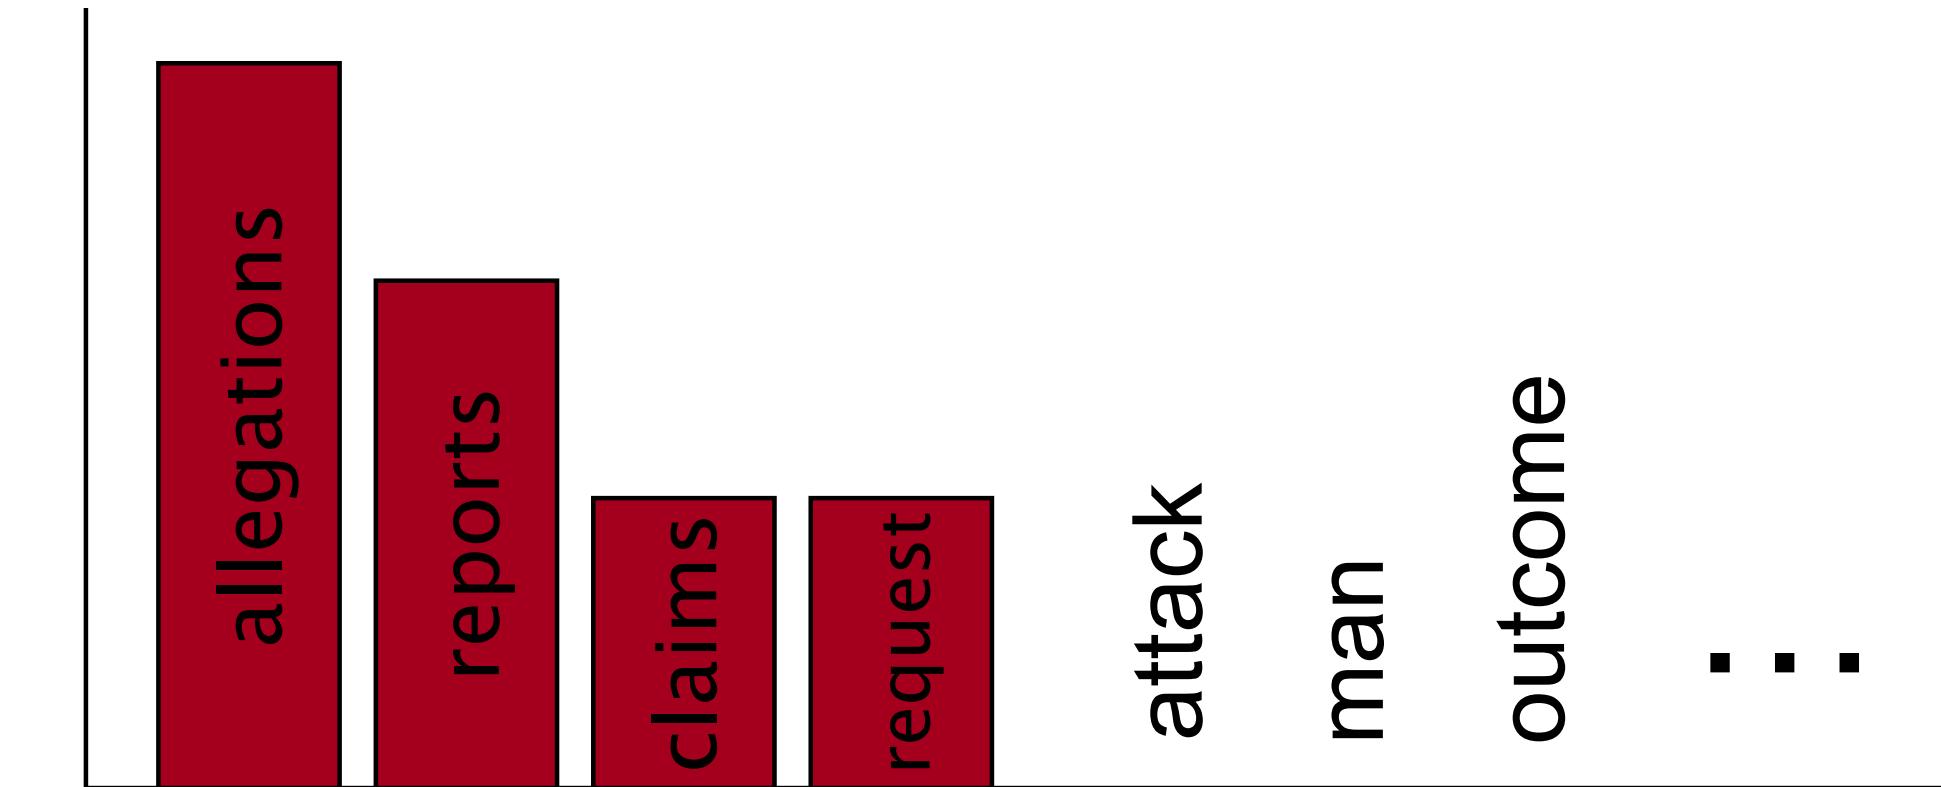
3 allegations

2 reports

1 claims

1 request

7 total



Steal probability mass to generalize better

$P(w \mid \text{denied the})$

2.5 allegations

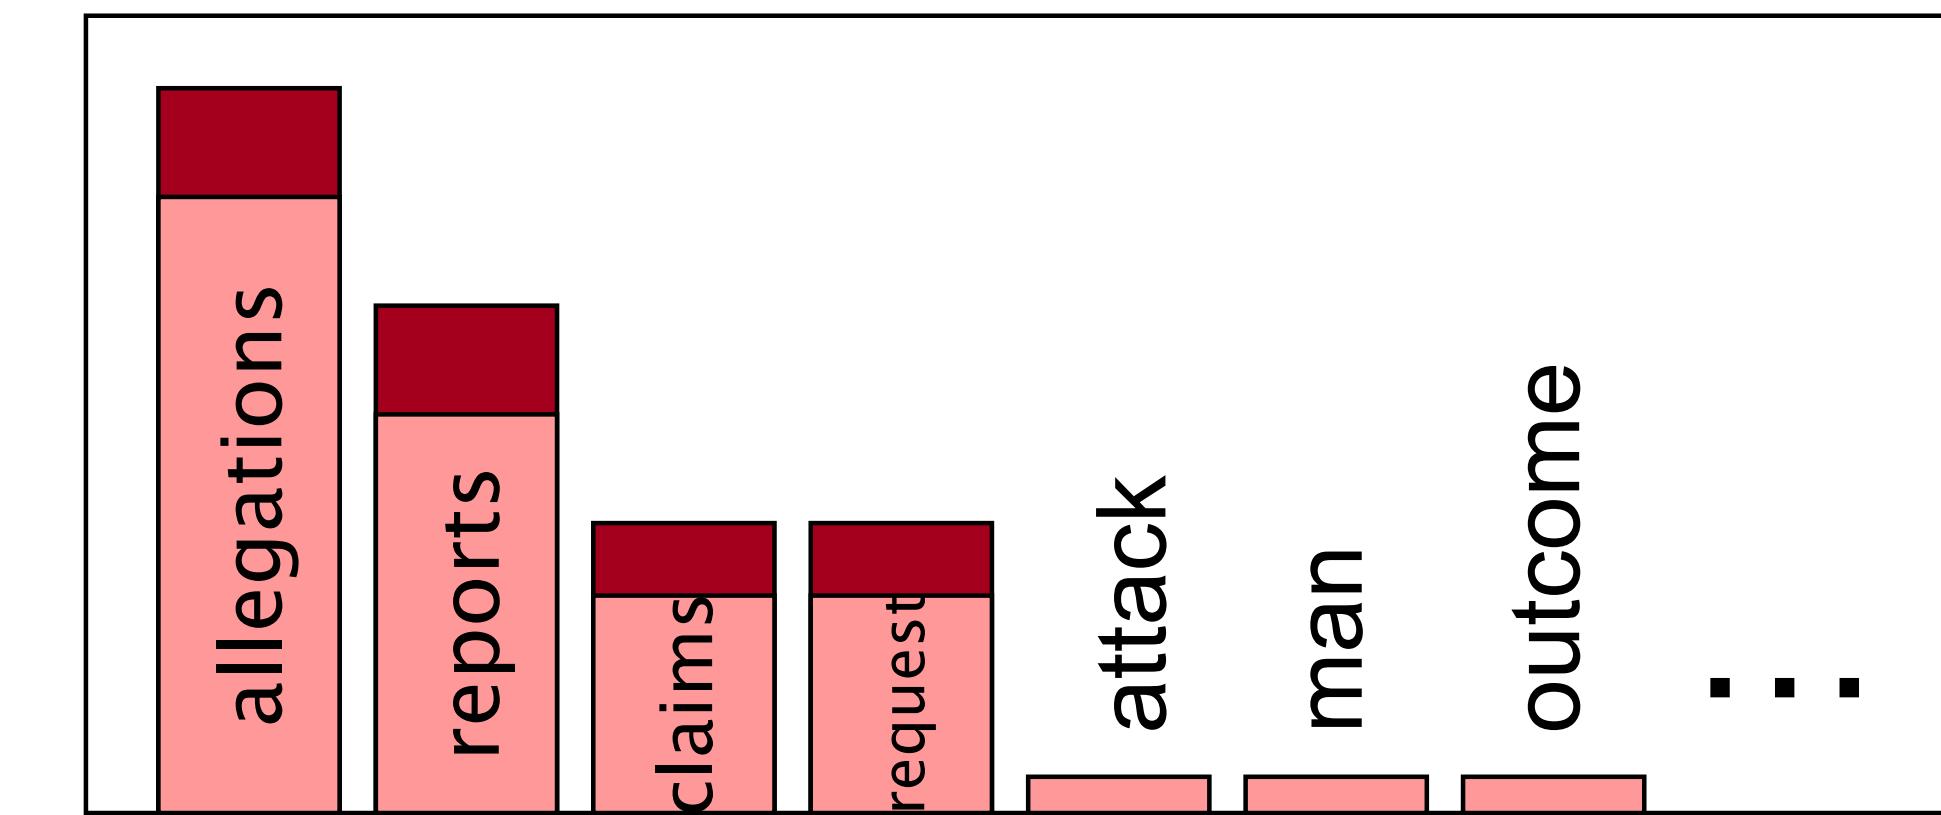
1.5 reports

0.5 claims

0.5 request

2 other

7 total



Laplace smoothing (add- α)

- **Simplest smoothing heuristic:** add α to all counts and renormalize!
- **Example:** max likelihood estimate for bigrams:

$$P(w_i|w_{i-1}) = \frac{C(w_{i-1}, w_i)}{C(w_{i-1})}$$

- After smoothing:

$$P(w_i|w_{i-1}) = \frac{C(w_{i-1}, w_i) + \alpha}{C(w_{i-1}) + \alpha|V|}$$

Raw bigram counts (Berkeley restaurant corpus)

- Out of 9222 sentences

| | i | want | to | eat | chinese | food | lunch | spend |
|---------|----|------|-----|-----|---------|------|-------|-------|
| i | 5 | 827 | 0 | 9 | 0 | 0 | 0 | 2 |
| want | 2 | 0 | 608 | 1 | 6 | 6 | 5 | 1 |
| to | 2 | 0 | 4 | 686 | 2 | 0 | 6 | 211 |
| eat | 0 | 0 | 2 | 0 | 16 | 2 | 42 | 0 |
| chinese | 1 | 0 | 0 | 0 | 0 | 82 | 1 | 0 |
| food | 15 | 0 | 15 | 0 | 1 | 4 | 0 | 0 |
| lunch | 2 | 0 | 0 | 0 | 0 | 1 | 0 | 0 |
| spend | 1 | 0 | 1 | 0 | 0 | 0 | 0 | 0 |

Smoothed bigram counts

| | i | want | to | eat | chinese | food | lunch | spend |
|---------|----|------|-----|-----|---------|------|-------|-------|
| i | 6 | 828 | 1 | 10 | 1 | 1 | 1 | 3 |
| want | 3 | 1 | 609 | 2 | 7 | 7 | 6 | 2 |
| to | 3 | 1 | 5 | 687 | 3 | 1 | 7 | 212 |
| eat | 1 | 1 | 3 | 1 | 17 | 3 | 43 | 1 |
| chinese | 2 | 1 | 1 | 1 | 1 | 83 | 2 | 1 |
| food | 16 | 1 | 16 | 1 | 2 | 5 | 1 | 1 |
| lunch | 3 | 1 | 1 | 1 | 1 | 2 | 1 | 1 |
| spend | 2 | 1 | 2 | 1 | 1 | 1 | 1 | 1 |

Add 1 to all the entries in the matrix

Smoothed bigram probabilities

$$P^*(w_i | w_{i-1}) = \frac{C(w_{i-1}w_i) + 1}{C(w_{i-1}) + V}$$

| | i | want | to | eat | chinese | food | lunch | spend |
|---------|----------------|----------------|----------------|----------------|----------------|----------------|----------------|----------------|
| i | 0.0015 | 0.21 | 0.00025 | 0.0025 | 0.00025 | 0.00025 | 0.00025 | 0.00075 |
| want | 0.0013 | 0.00042 | 0.26 | 0.00084 | 0.0029 | 0.0029 | 0.0025 | 0.00084 |
| to | 0.00078 | 0.00026 | 0.0013 | 0.18 | 0.00078 | 0.00026 | 0.0018 | 0.055 |
| eat | 0.00046 | 0.00046 | 0.0014 | 0.00046 | 0.0078 | 0.0014 | 0.02 | 0.00046 |
| chinese | 0.0012 | 0.00062 | 0.00062 | 0.00062 | 0.00062 | 0.052 | 0.0012 | 0.00062 |
| food | 0.0063 | 0.00039 | 0.0063 | 0.00039 | 0.00079 | 0.002 | 0.00039 | 0.00039 |
| lunch | 0.0017 | 0.00056 | 0.00056 | 0.00056 | 0.00056 | 0.0011 | 0.00056 | 0.00056 |
| spend | 0.0012 | 0.00058 | 0.0012 | 0.00058 | 0.00058 | 0.00058 | 0.00058 | 0.00058 |

Discounting

| Bigram count in training | Bigram count in heldout set |
|--------------------------|-----------------------------|
| 0 | .0000270 |
| 1 | 0.448 |
| 2 | 1.25 |
| 3 | 2.24 |
| 4 | 3.23 |
| 5 | 4.21 |
| 6 | 5.23 |
| 7 | 6.21 |
| 8 | 7.21 |
| 9 | 8.26 |

- Determine some “mass” to remove from probability estimates
- Redistribute mass among unseen n-grams
- Just choose an absolute value to discount (usually <1)

$$P_{\text{abs_discount}}(w_i | w_{i-1}) = \frac{c(w_{i-1}, w_i) - d}{c(w_{i-1})} \quad \text{if } c(w_{i-1}, w_i) > 0$$

Unigram probabilities

$$\alpha(w_{i-1}) \frac{P(w_i)}{\sum_{w'} P(w')} \quad \text{for all } w' \text{ s.t. } c(w_{i-1}, w') = 0 \quad \text{if } c(w_{i-1}, w_i) = 0$$

40

Back-off

- Use n-gram if enough evidence, else back off to (n-1)-gram

n here is length of *n*-gram, not length of sequence (as in previous slides)

$$P_{bo}(w_i \mid w_{i-n+1} \cdots w_{i-1}) = \begin{cases} d_{w_{i-n+1} \cdots w_i} \frac{C(w_{i-n+1} \cdots w_{i-1} w_i)}{C(w_{i-n+1} \cdots w_{i-1})} & \text{if } C(w_{i-n+1} \cdots w_i) > k \\ \alpha_{w_{i-n+1} \cdots w_{i-1}} P_{bo}(w_i \mid w_{i-n+2} \cdots w_{i-1}) & \text{otherwise} \end{cases}$$

(Katz back-off)

- d = amount of discounting
- α = back-off weight

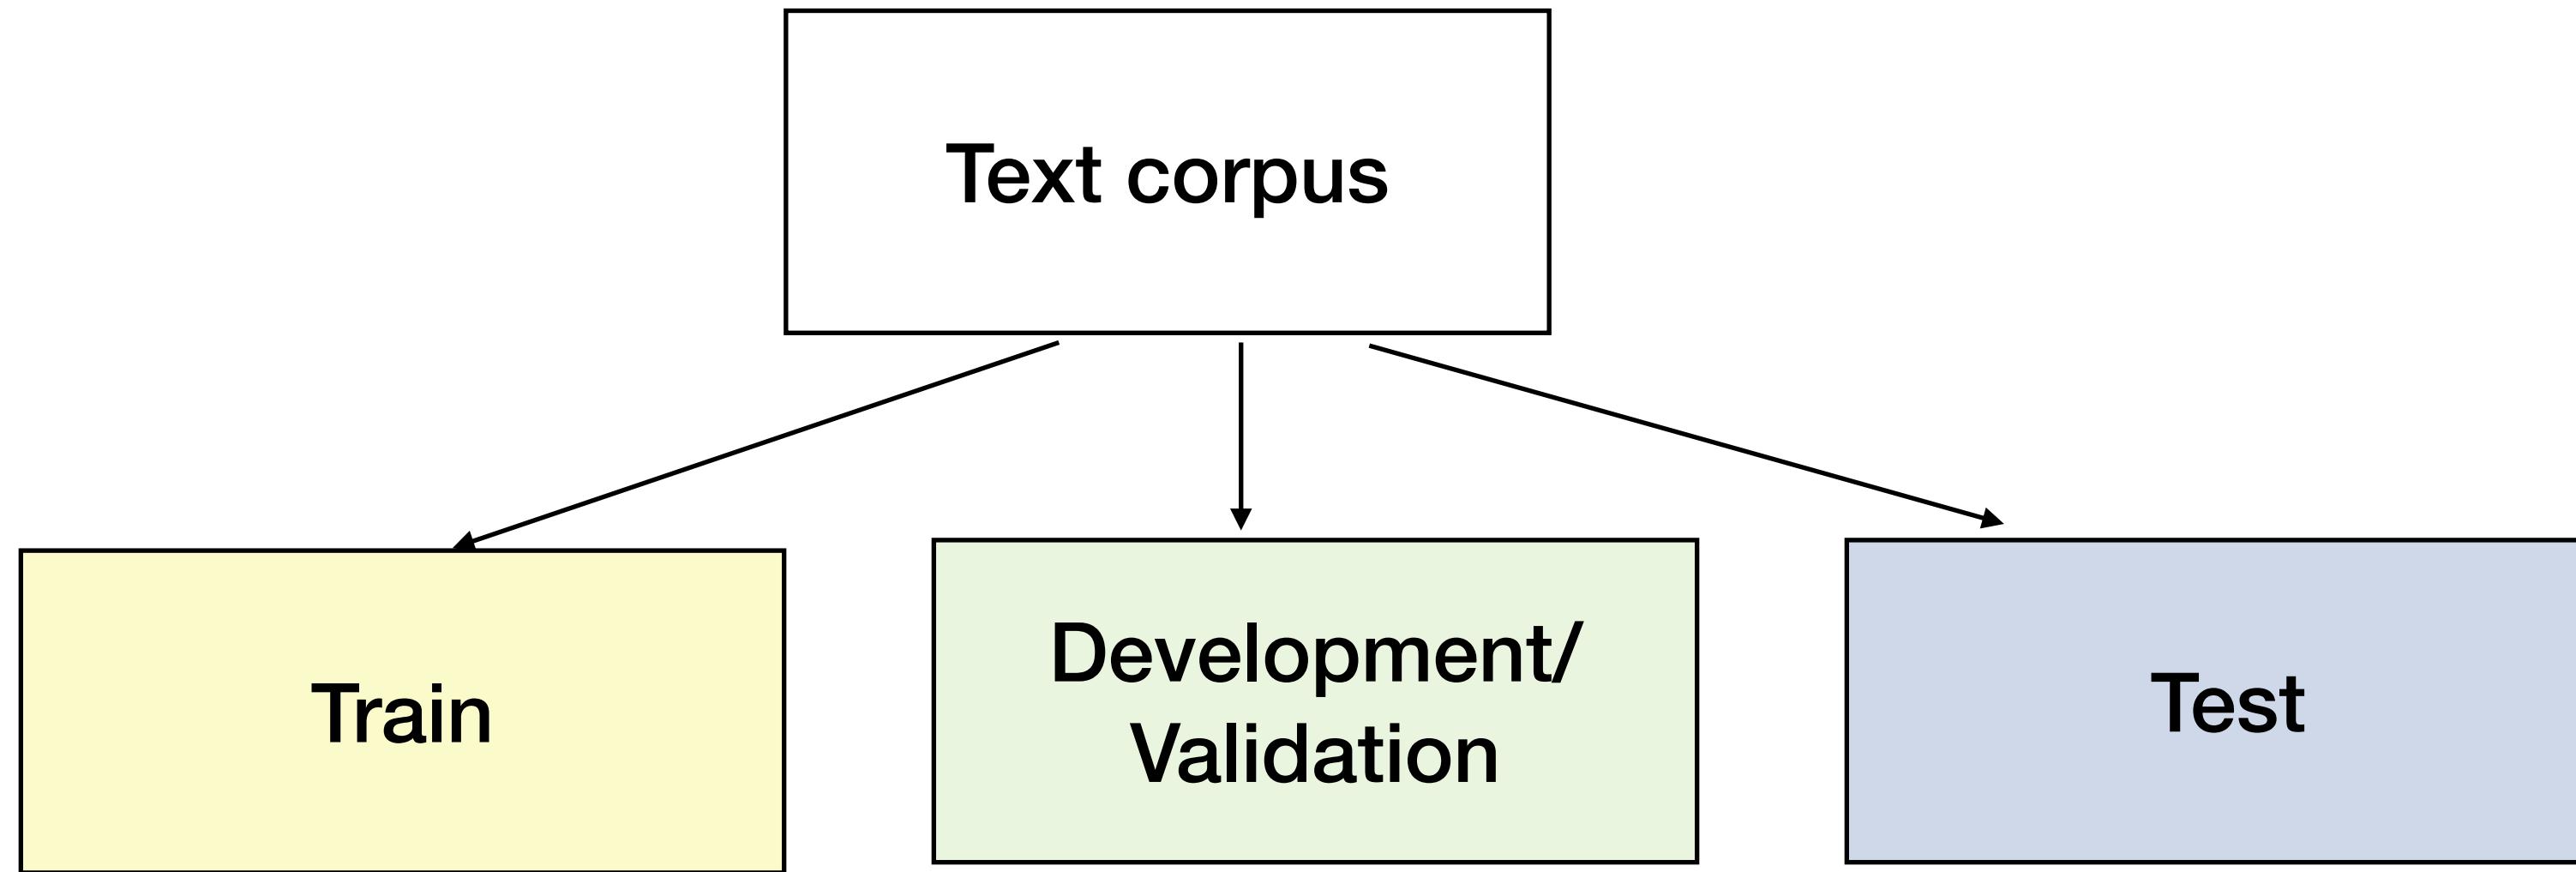
Linear Interpolation

$$\hat{P}(w_i|w_{i-1}, w_{i-2}) = \lambda_1 P(w_i|w_{i-1}, w_{i-2}) \quad \text{Trigram}$$
$$+ \lambda_2 P(w_i|w_{i-1}) \quad \text{Bigram}$$
$$+ \lambda_3 P(w_i) \quad \text{Unigram}$$

$$\sum_j \lambda_j = 1$$

- Use a combination of models to estimate probability
- Strong empirical performance

Choosing lambdas



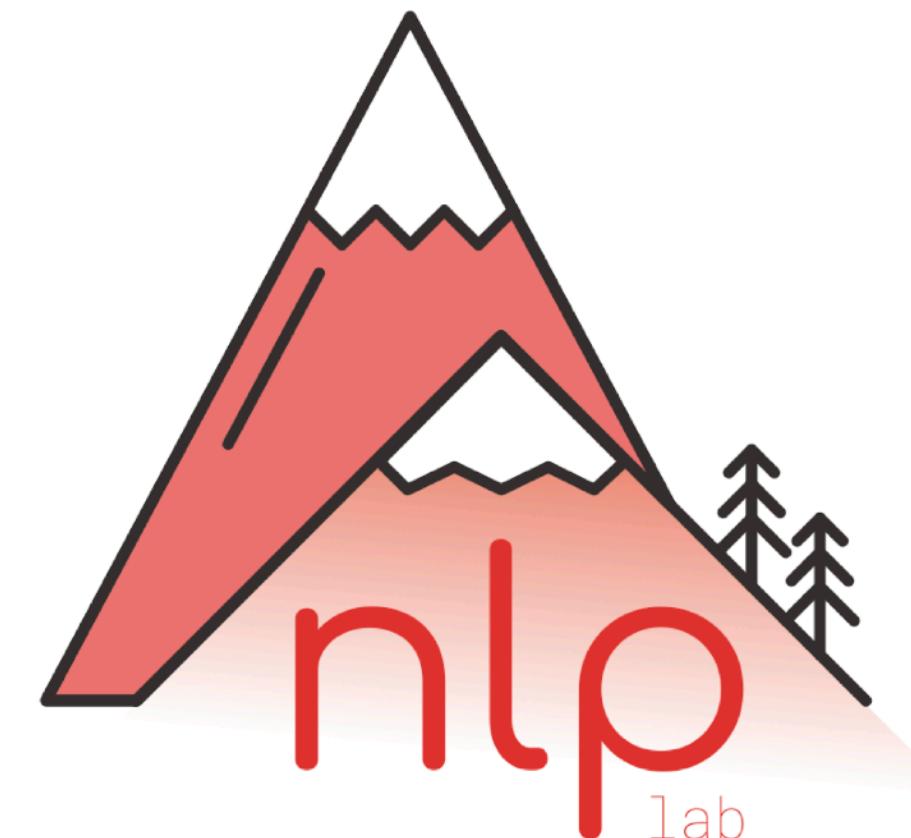
- First, estimate n-gram prob. on training set
- Then, estimate lambdas (hyperparameters) to maximize probability on the held-out development/validation set
- Use best model from above to evaluate on test set

Recap

- n -gram language models are simple, effective methods for estimating the probability of sequences
- **Problem #1:** Number of modelable sequences grows exponentially with n
- **Problem #2:** Smoothing heuristics must be used to estimate the probability of sequences unseen during training

Language Models: Fixed-context Neural Models

Antoine Bosselut



Today's Outline

- **Part 1:** Language models introduction, count-based language models, evaluating language models, smoothing
- **Part 2:** Fixed-context Language Models
- **Additional Slides:** Neural networks recap

Question

What's another issue with n-gram language models?

Problem

With n-gram language models, there is no notion of similarity unless sequence overlaps!

We treat all words / prefixes as independent of one another!

students opened their __

pupils opened their __

scholars opened their __

undergraduates opened their __

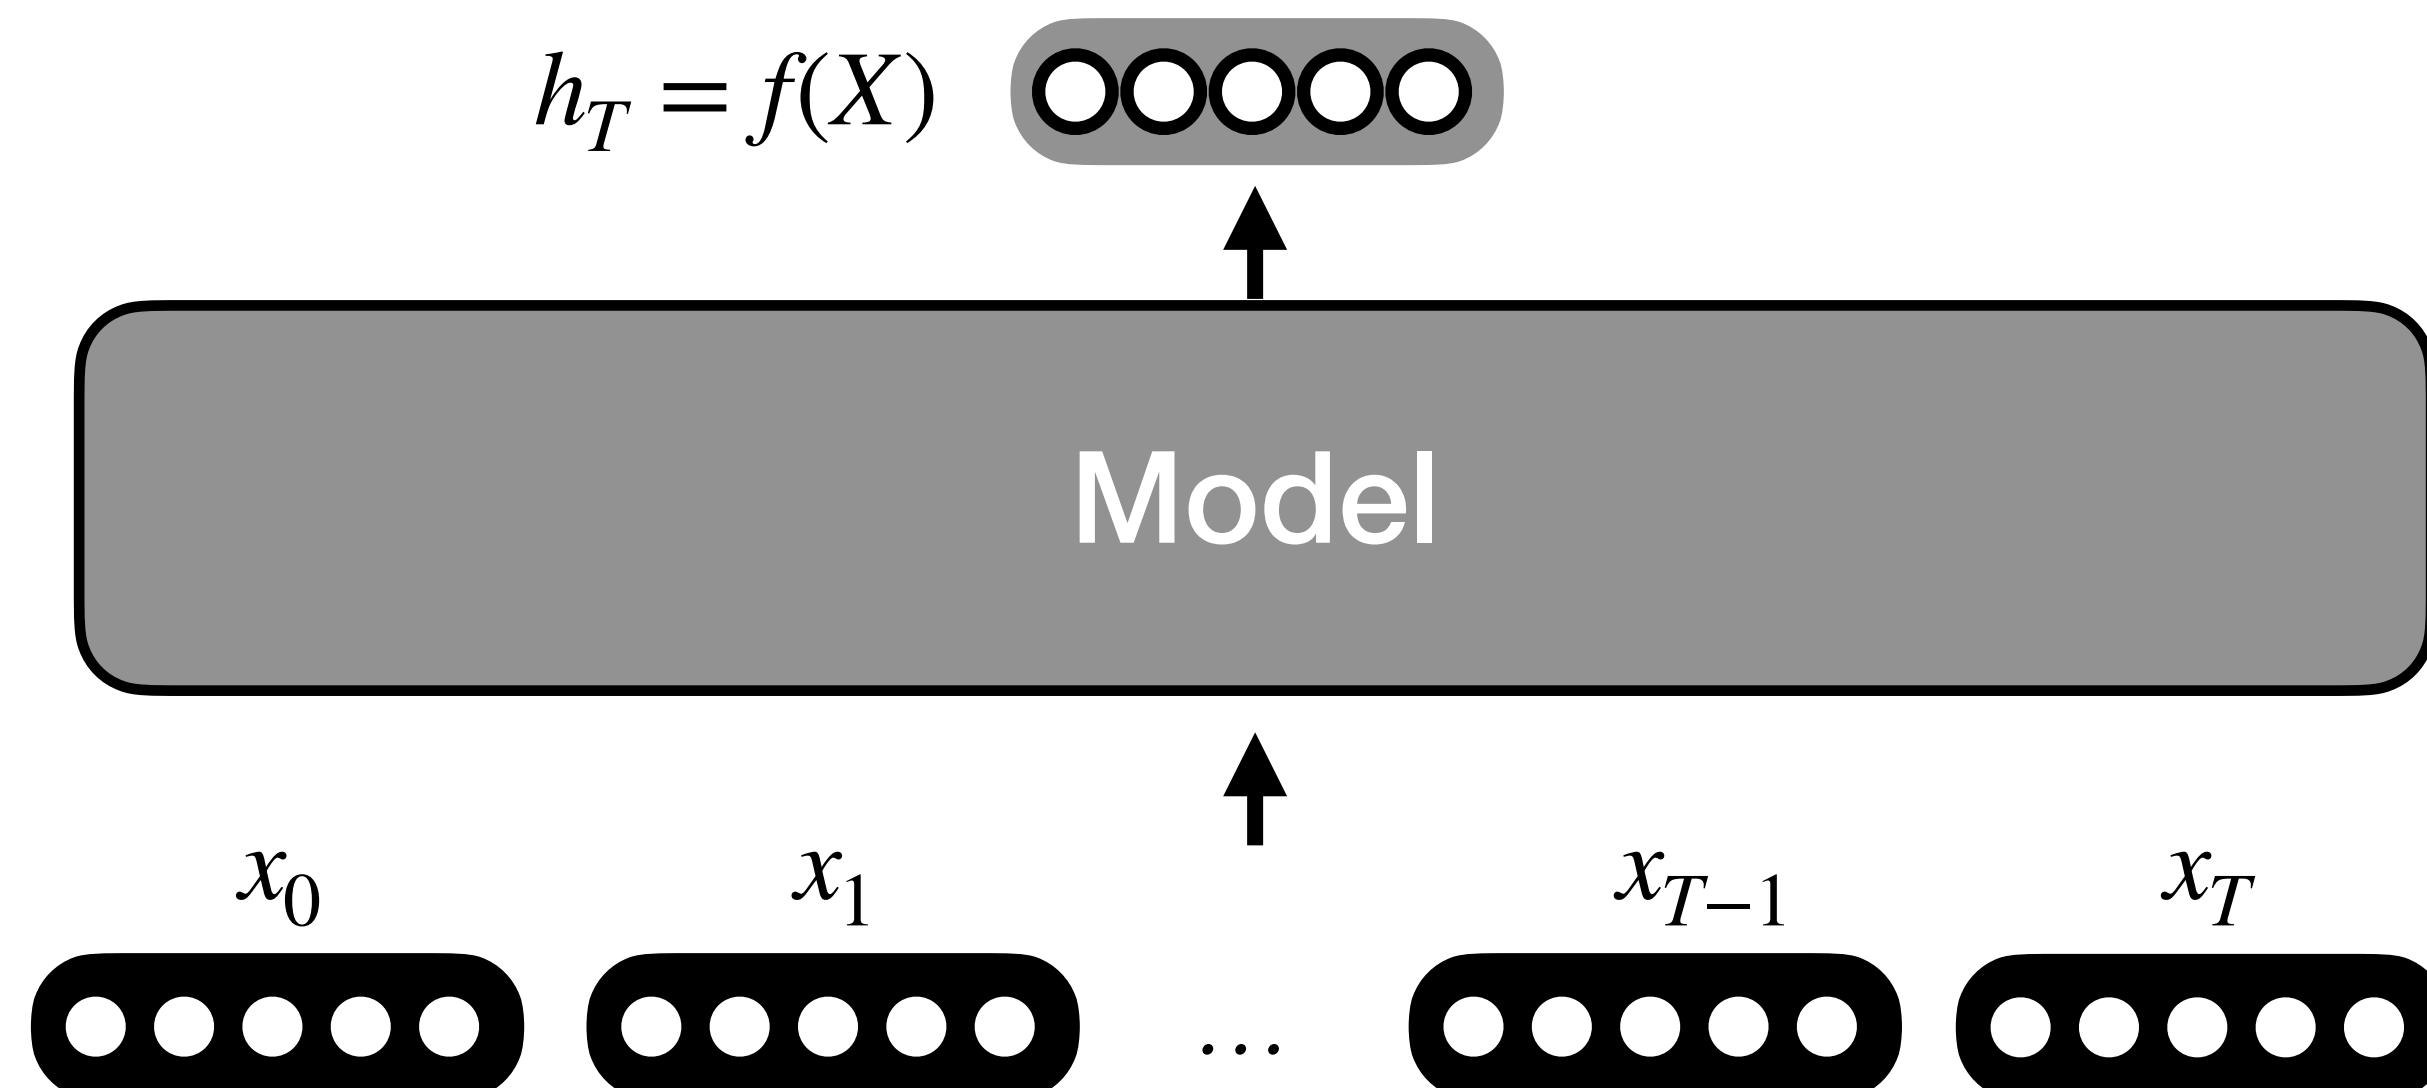
students turned the pages of their __

students attentively perused their __

Ideally, we should share information across these same prefixes!

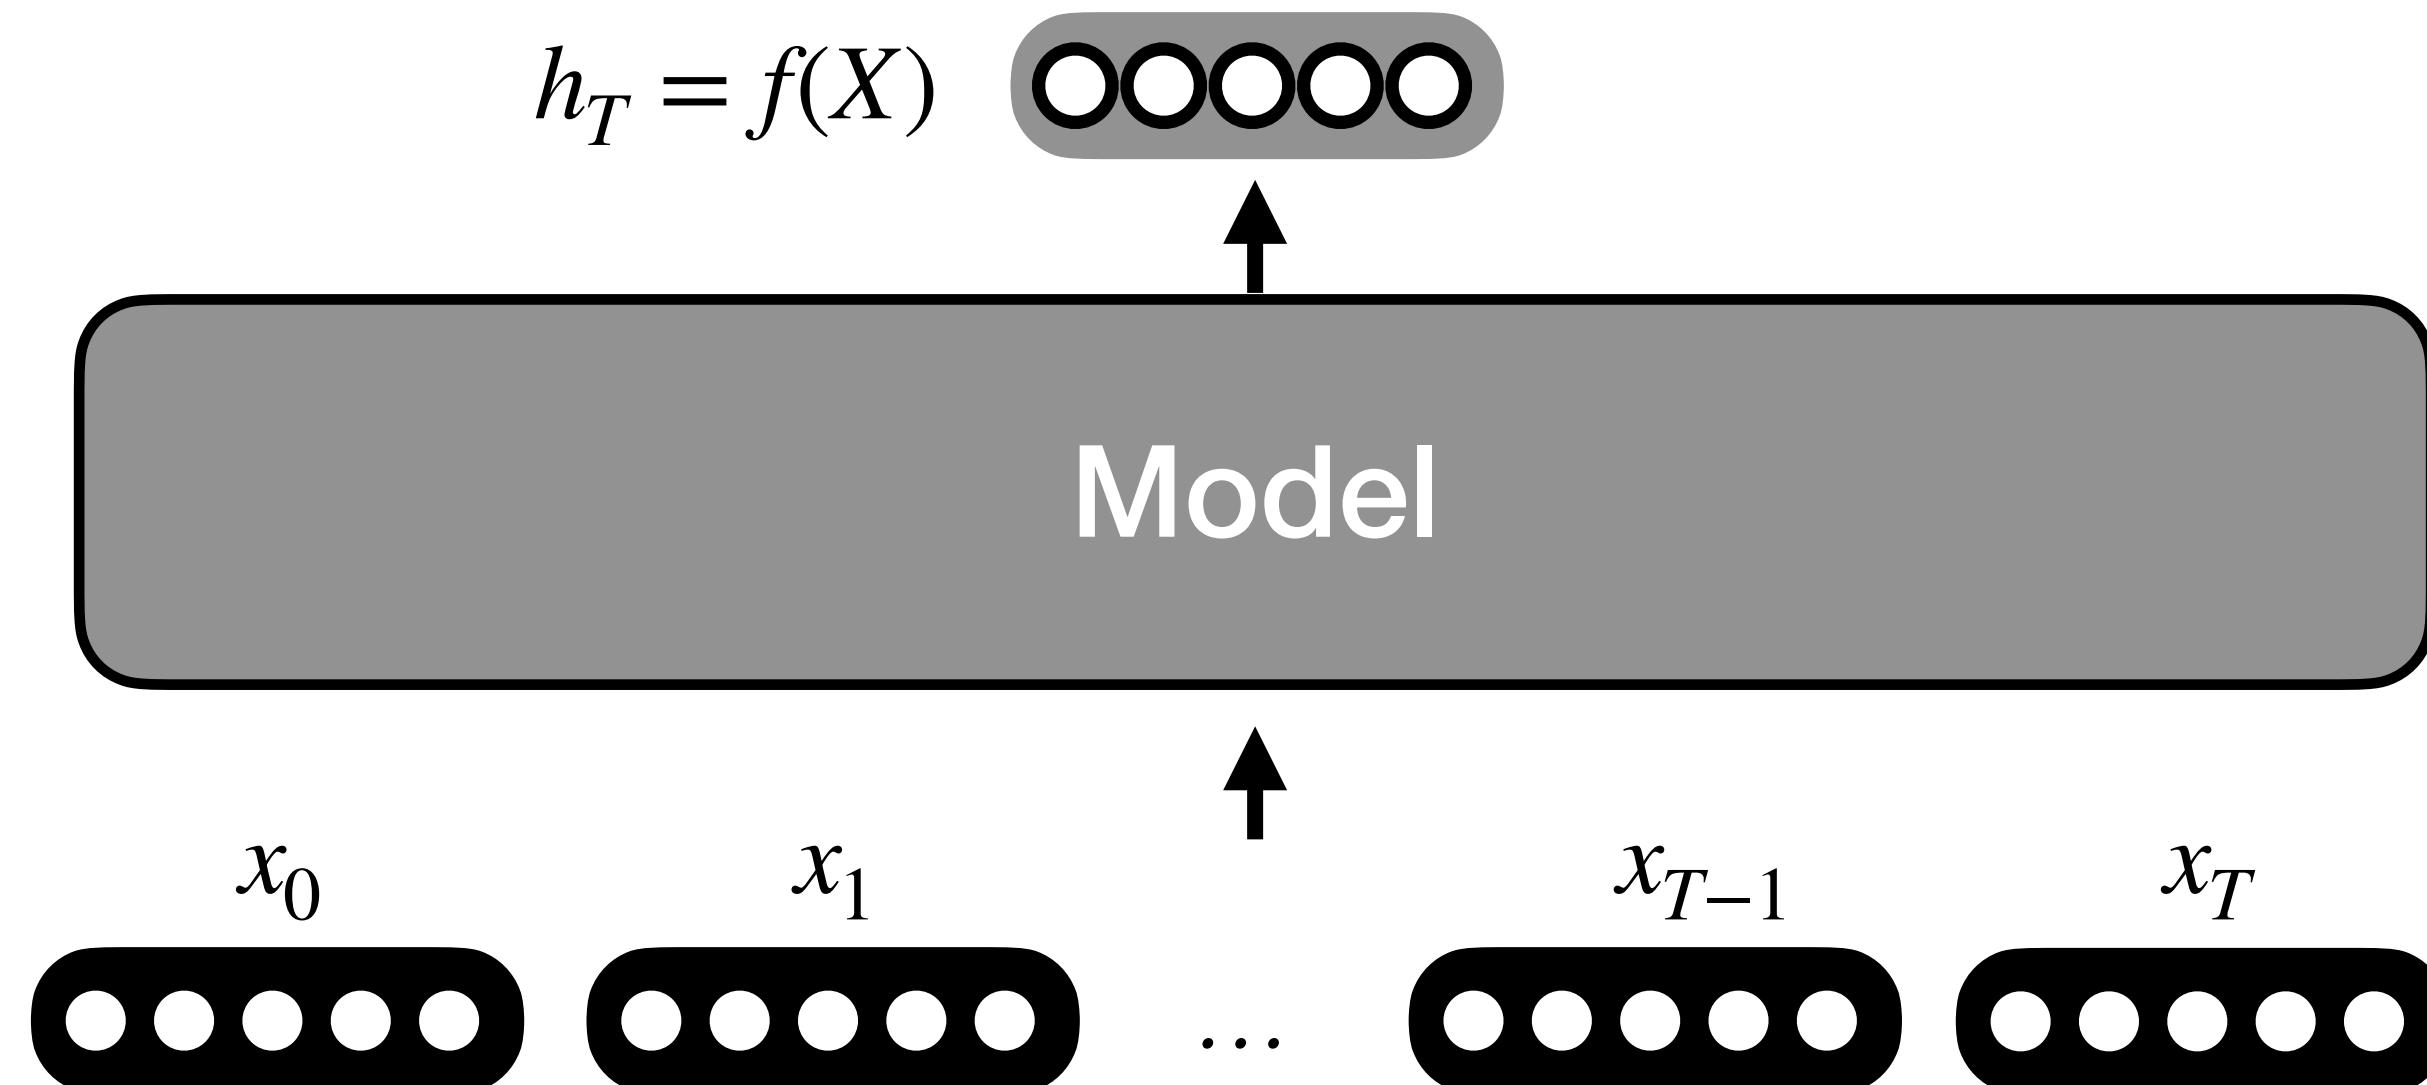
Recall

- neural networks **compose** word embeddings into vectors for natural language phrases, sentences, and documents



Recall

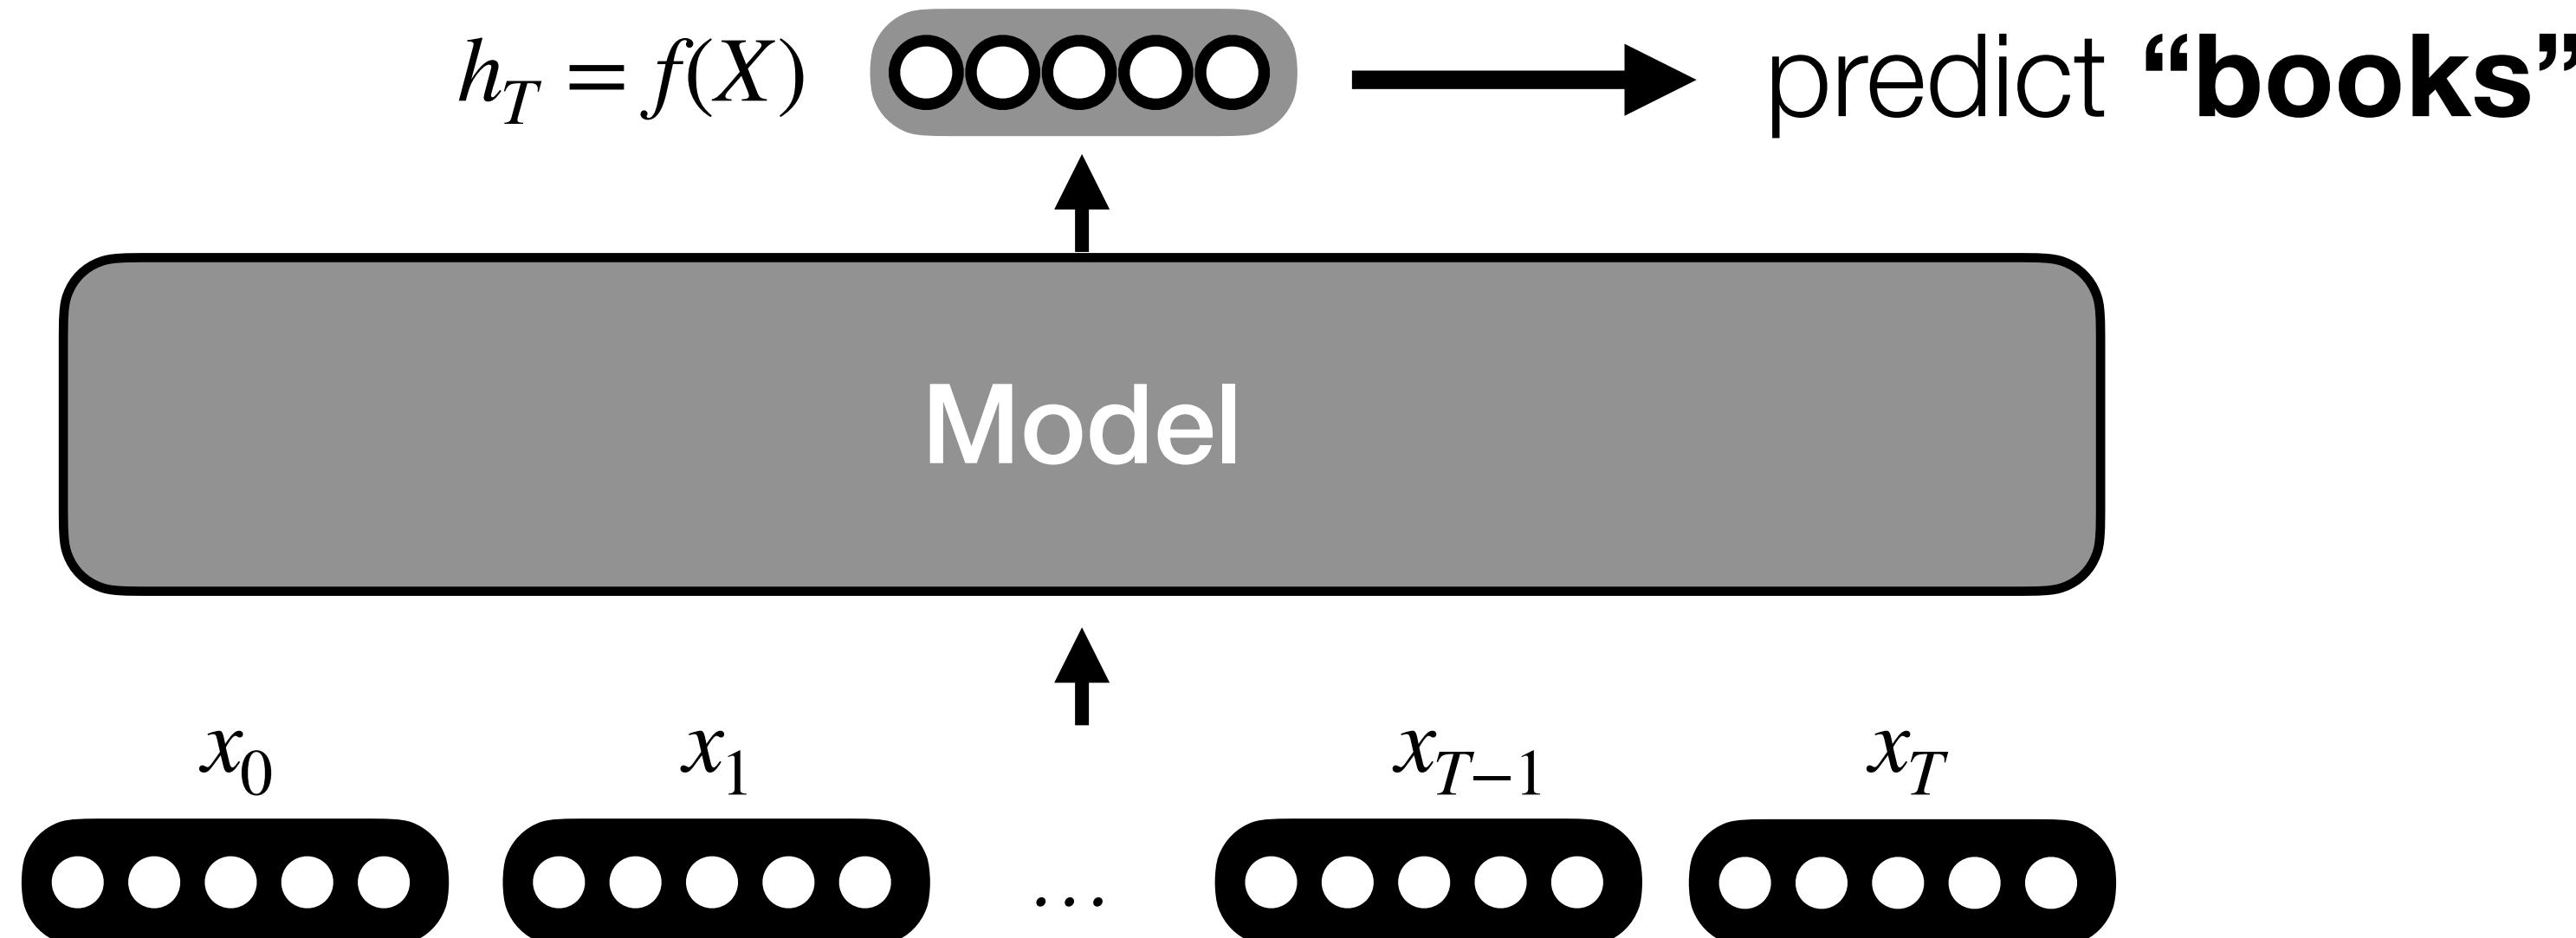
- neural networks **compose** word embeddings into vectors for natural language phrases, sentences, and documents
- Predict the next word from composed prefix representation



How does this happen?

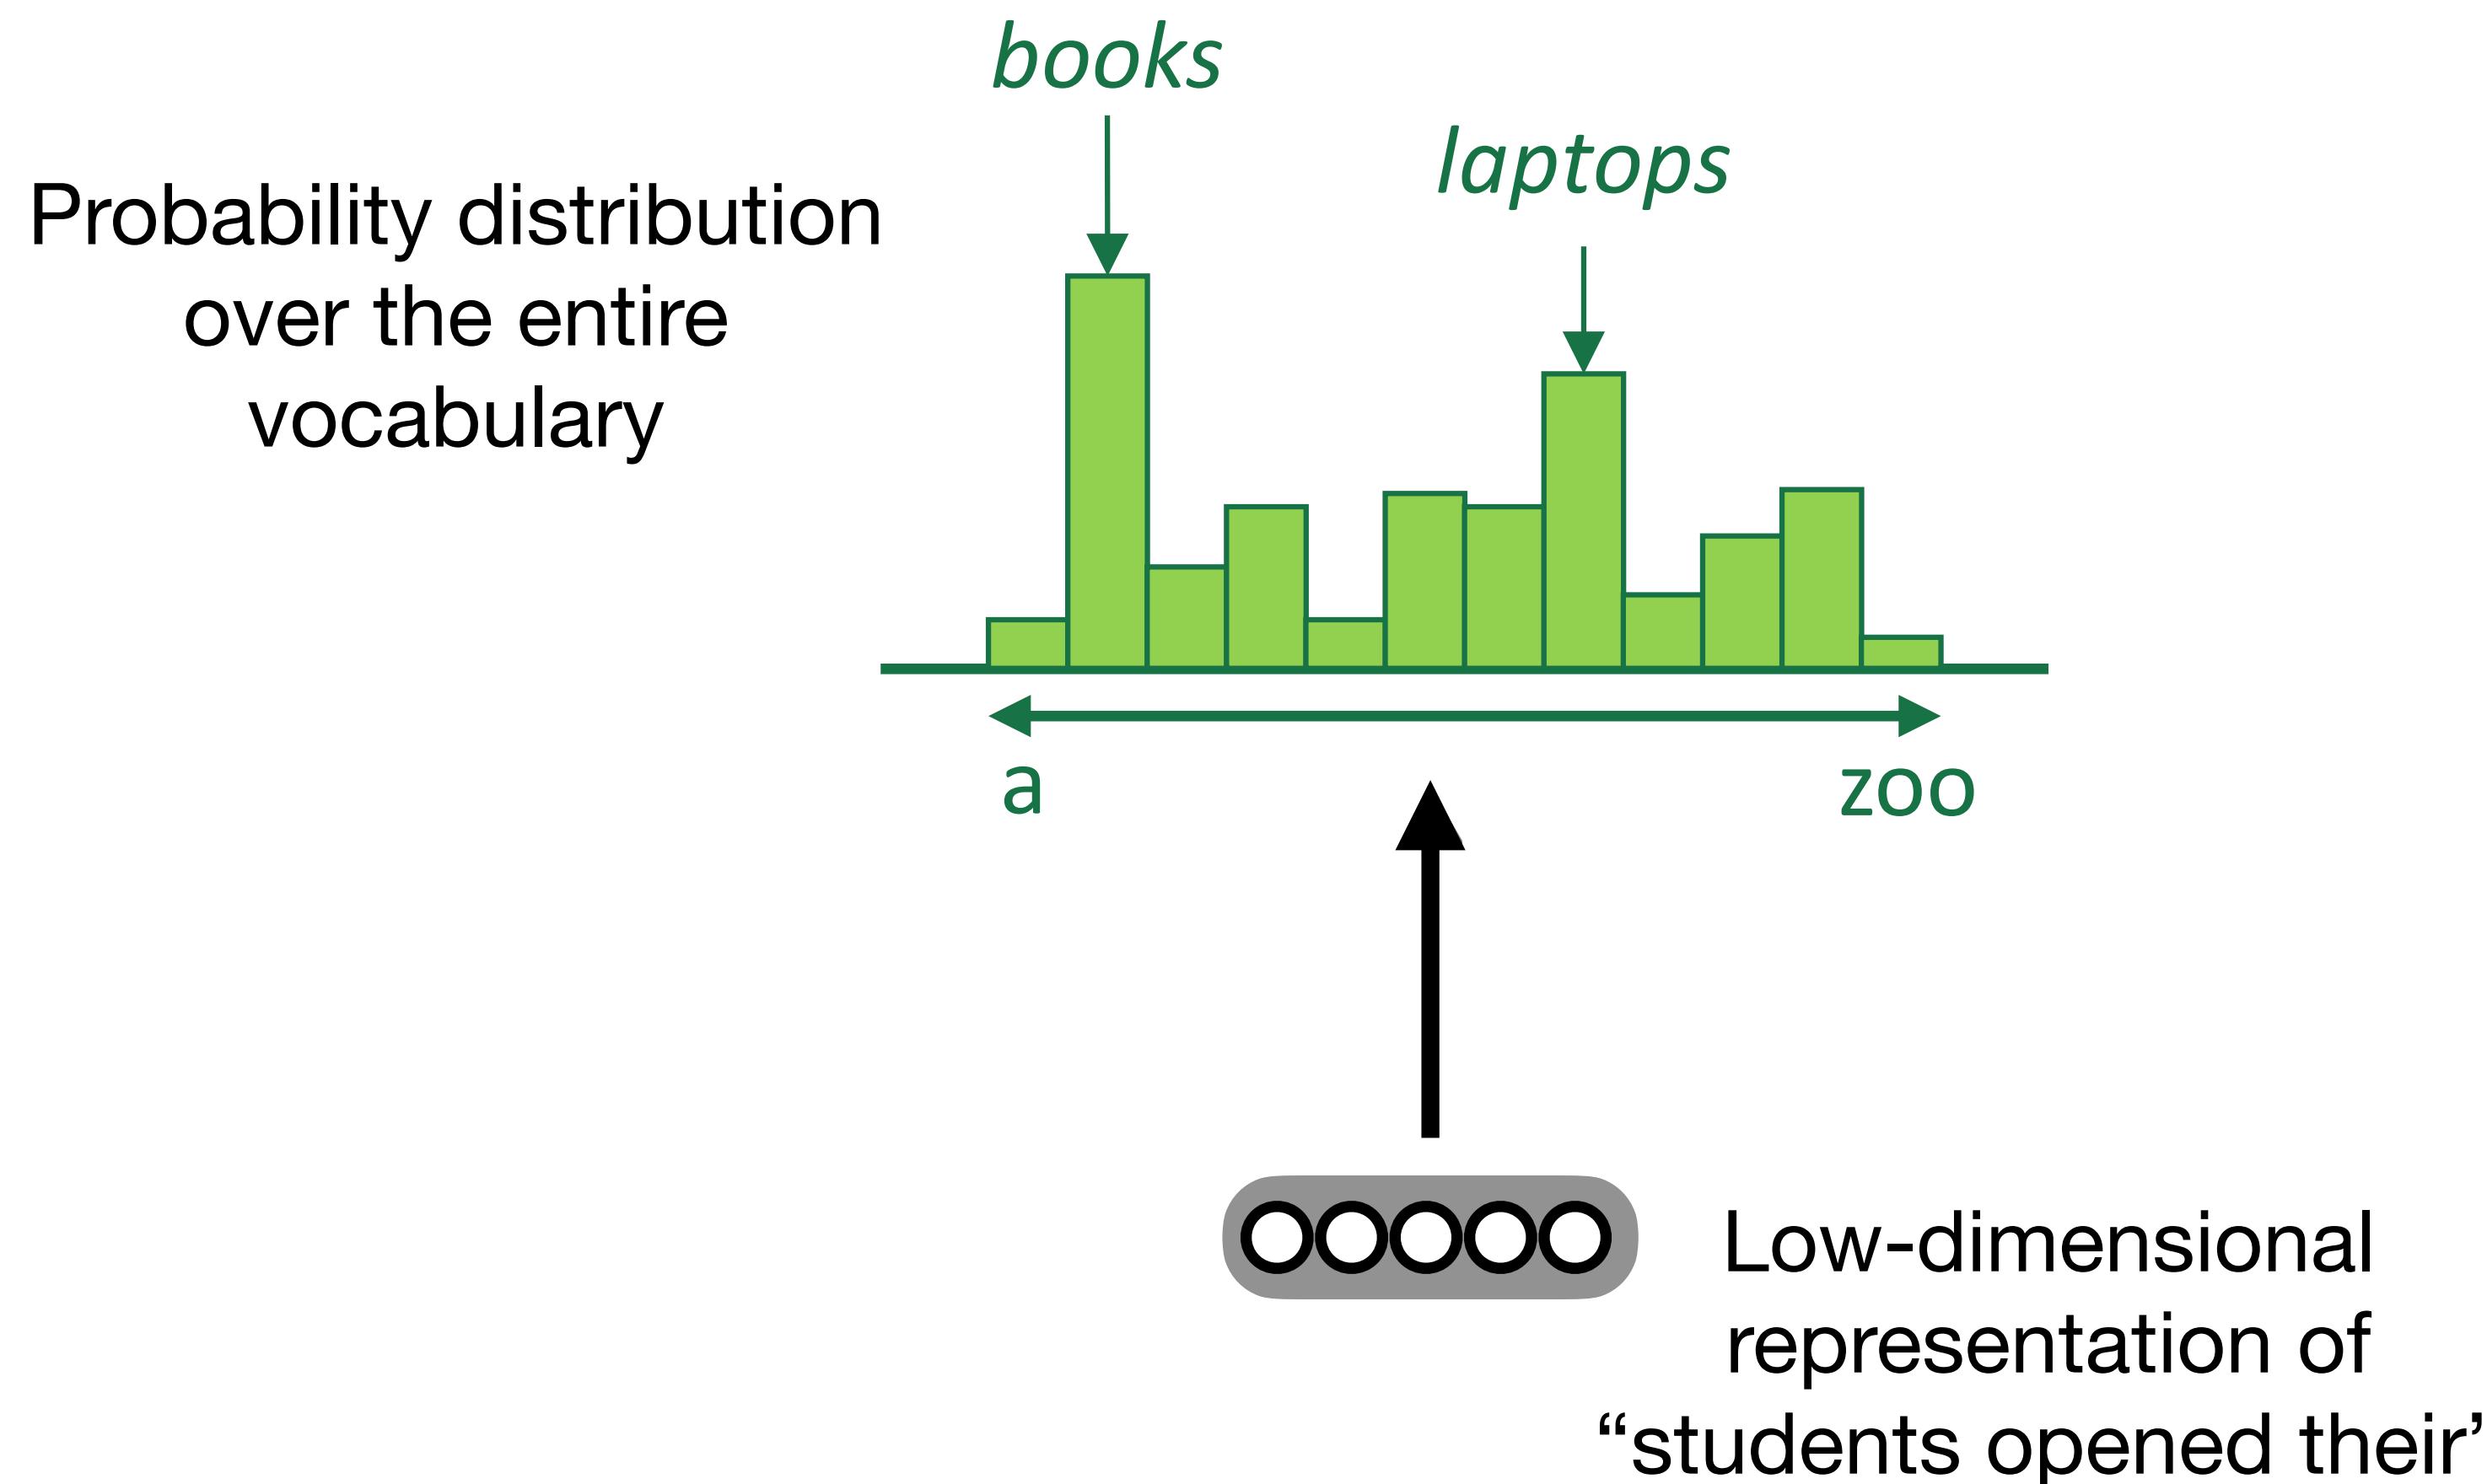
Vocab Projection + Softmax layers:

Converts a vector representation
into a probability distribution over
the entire vocabulary



Vocabulary Space Projection

$P(w_i | \text{vector for "students opened their"})$



Recap

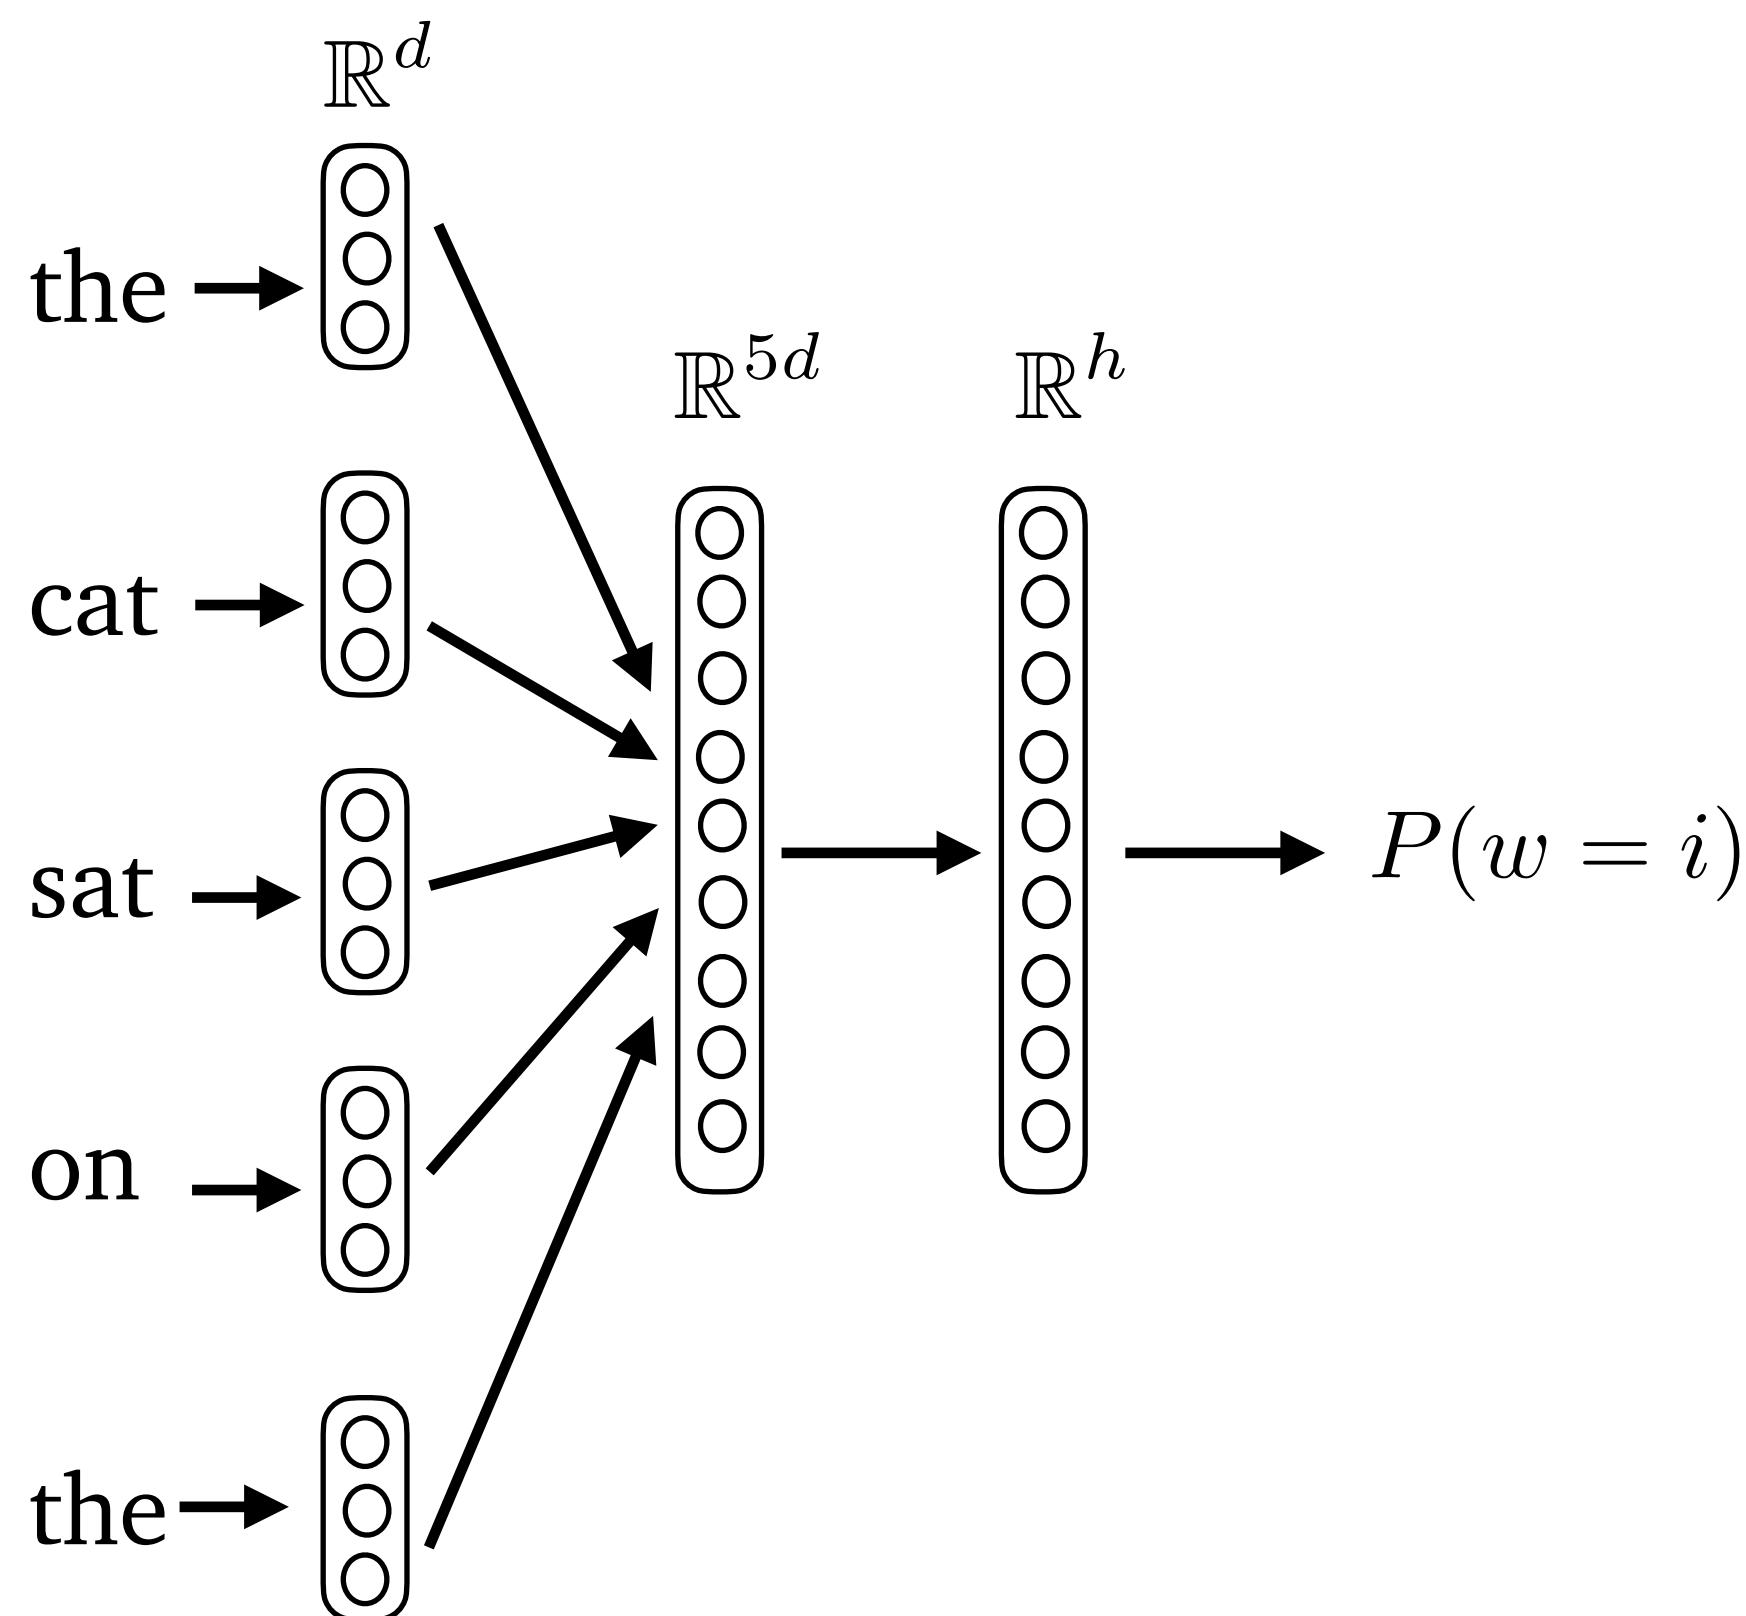
- Given a d -dimensional vector representation \mathbf{x} of a prefix, we do the following to predict the next word:
 1. Project it to a V -dimensional vector using a matrix-vector product (a.k.a. a “linear layer”, or a “feedforward layer”), where V is the size of the vocabulary
 2. Apply the softmax function to transform the resulting vector into a probability distribution

Question

**How can we use these principles to come up with
a better language model than the n-gram one?**

Fixed-context Neural Language Models

- $P(\text{mat} \mid \text{the cat sat on the}) = ?$



- Input layer ($n = 5$):

$$\mathbf{x} = [\mathbf{e}_{\text{the}}; \mathbf{e}_{\text{cat}}; \mathbf{e}_{\text{sat}}; \mathbf{e}_{\text{on}}; \mathbf{e}_{\text{the}}] \in \mathbb{R}^{dn}$$

- Hidden layer

$$\mathbf{h} = \tanh(\mathbf{W}\mathbf{x} + \mathbf{b}) \in \mathbb{R}^h$$

- Output layer (softmax)

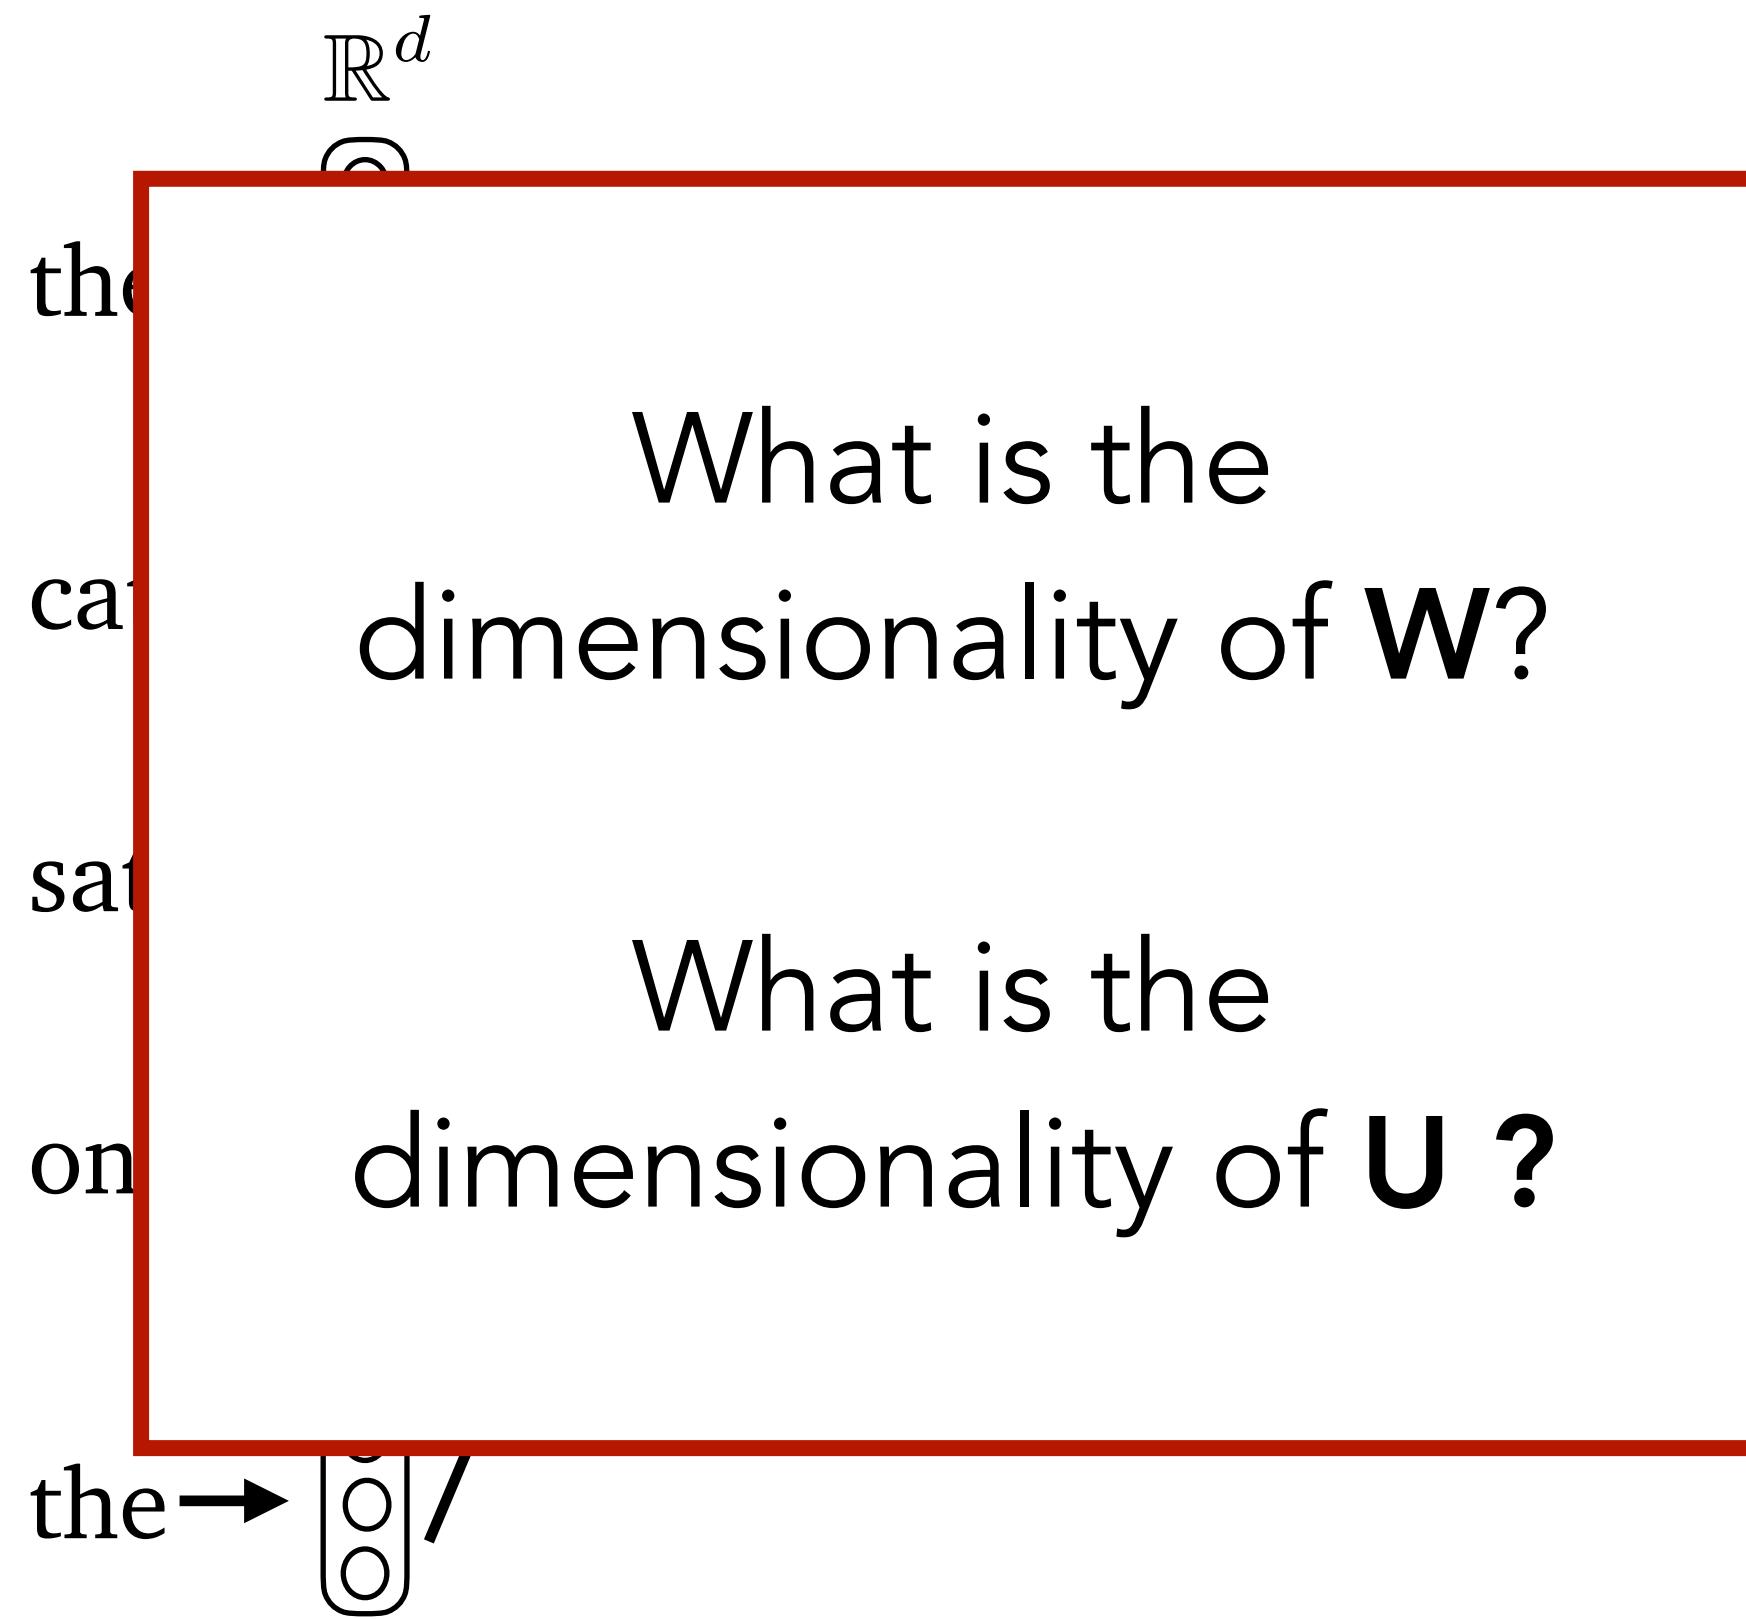
$$\mathbf{z} = \mathbf{U}\mathbf{h} \in \mathbb{R}^{|V|}$$

$$P(w = i \mid \text{the cat sat on the})$$

$$= \text{softmax}_i(\mathbf{z}) = \frac{e^{z_i}}{\sum_k e^{z_k}}$$

Fixed-context Neural Language Models

- $P(\text{mat} \mid \text{the cat sat on the}) = ?$



- Input layer ($n = 5$):

$$\mathbf{x} = [\mathbf{e}_{\text{the}}; \mathbf{e}_{\text{cat}}; \mathbf{e}_{\text{sat}}; \mathbf{e}_{\text{on}}; \mathbf{e}_{\text{the}}] \in \mathbb{R}^{dn}$$

- Hidden layer

$$\mathbf{h} = \tanh(\mathbf{W}\mathbf{x} + \mathbf{b}) \in \mathbb{R}^h$$

- Output layer (softmax)

$$\mathbf{z} = \mathbf{U}\mathbf{h} \in \mathbb{R}^{|V|}$$

$$P(w = i \mid \text{the cat sat on the})$$

$$= \text{softmax}_i(\mathbf{z}) = \frac{e^{z_i}}{\sum_k e^{z_k}}$$

Question

What algorithm did we see last week do fixed context language models resemble the most?

Question

What algorithm did we see last week do fixed context language models resemble the most?

CBOW!

Instead of predicting the middle word we predict the next word

Advantages vs. Disadvantages

- No more sparsity problem
 - All sequences can be estimated with non-zero probability ! **Why?**
softmax never predicts 0 prob
- Model size is much smaller!
 - Depends on number of weights in model, not number of sequences!
- Fixed windows are still too small to encode **long-range dependencies**
- Enlarging the window size makes the weight matrix **W** larger (more computationally expensive!)
- Weights in **W** aren't shared across embeddings in the window (computationally inefficient!)

Sequences are not of consistent length

The **cat** chased the **mouse**

The starving **cat** chased the **mouse**

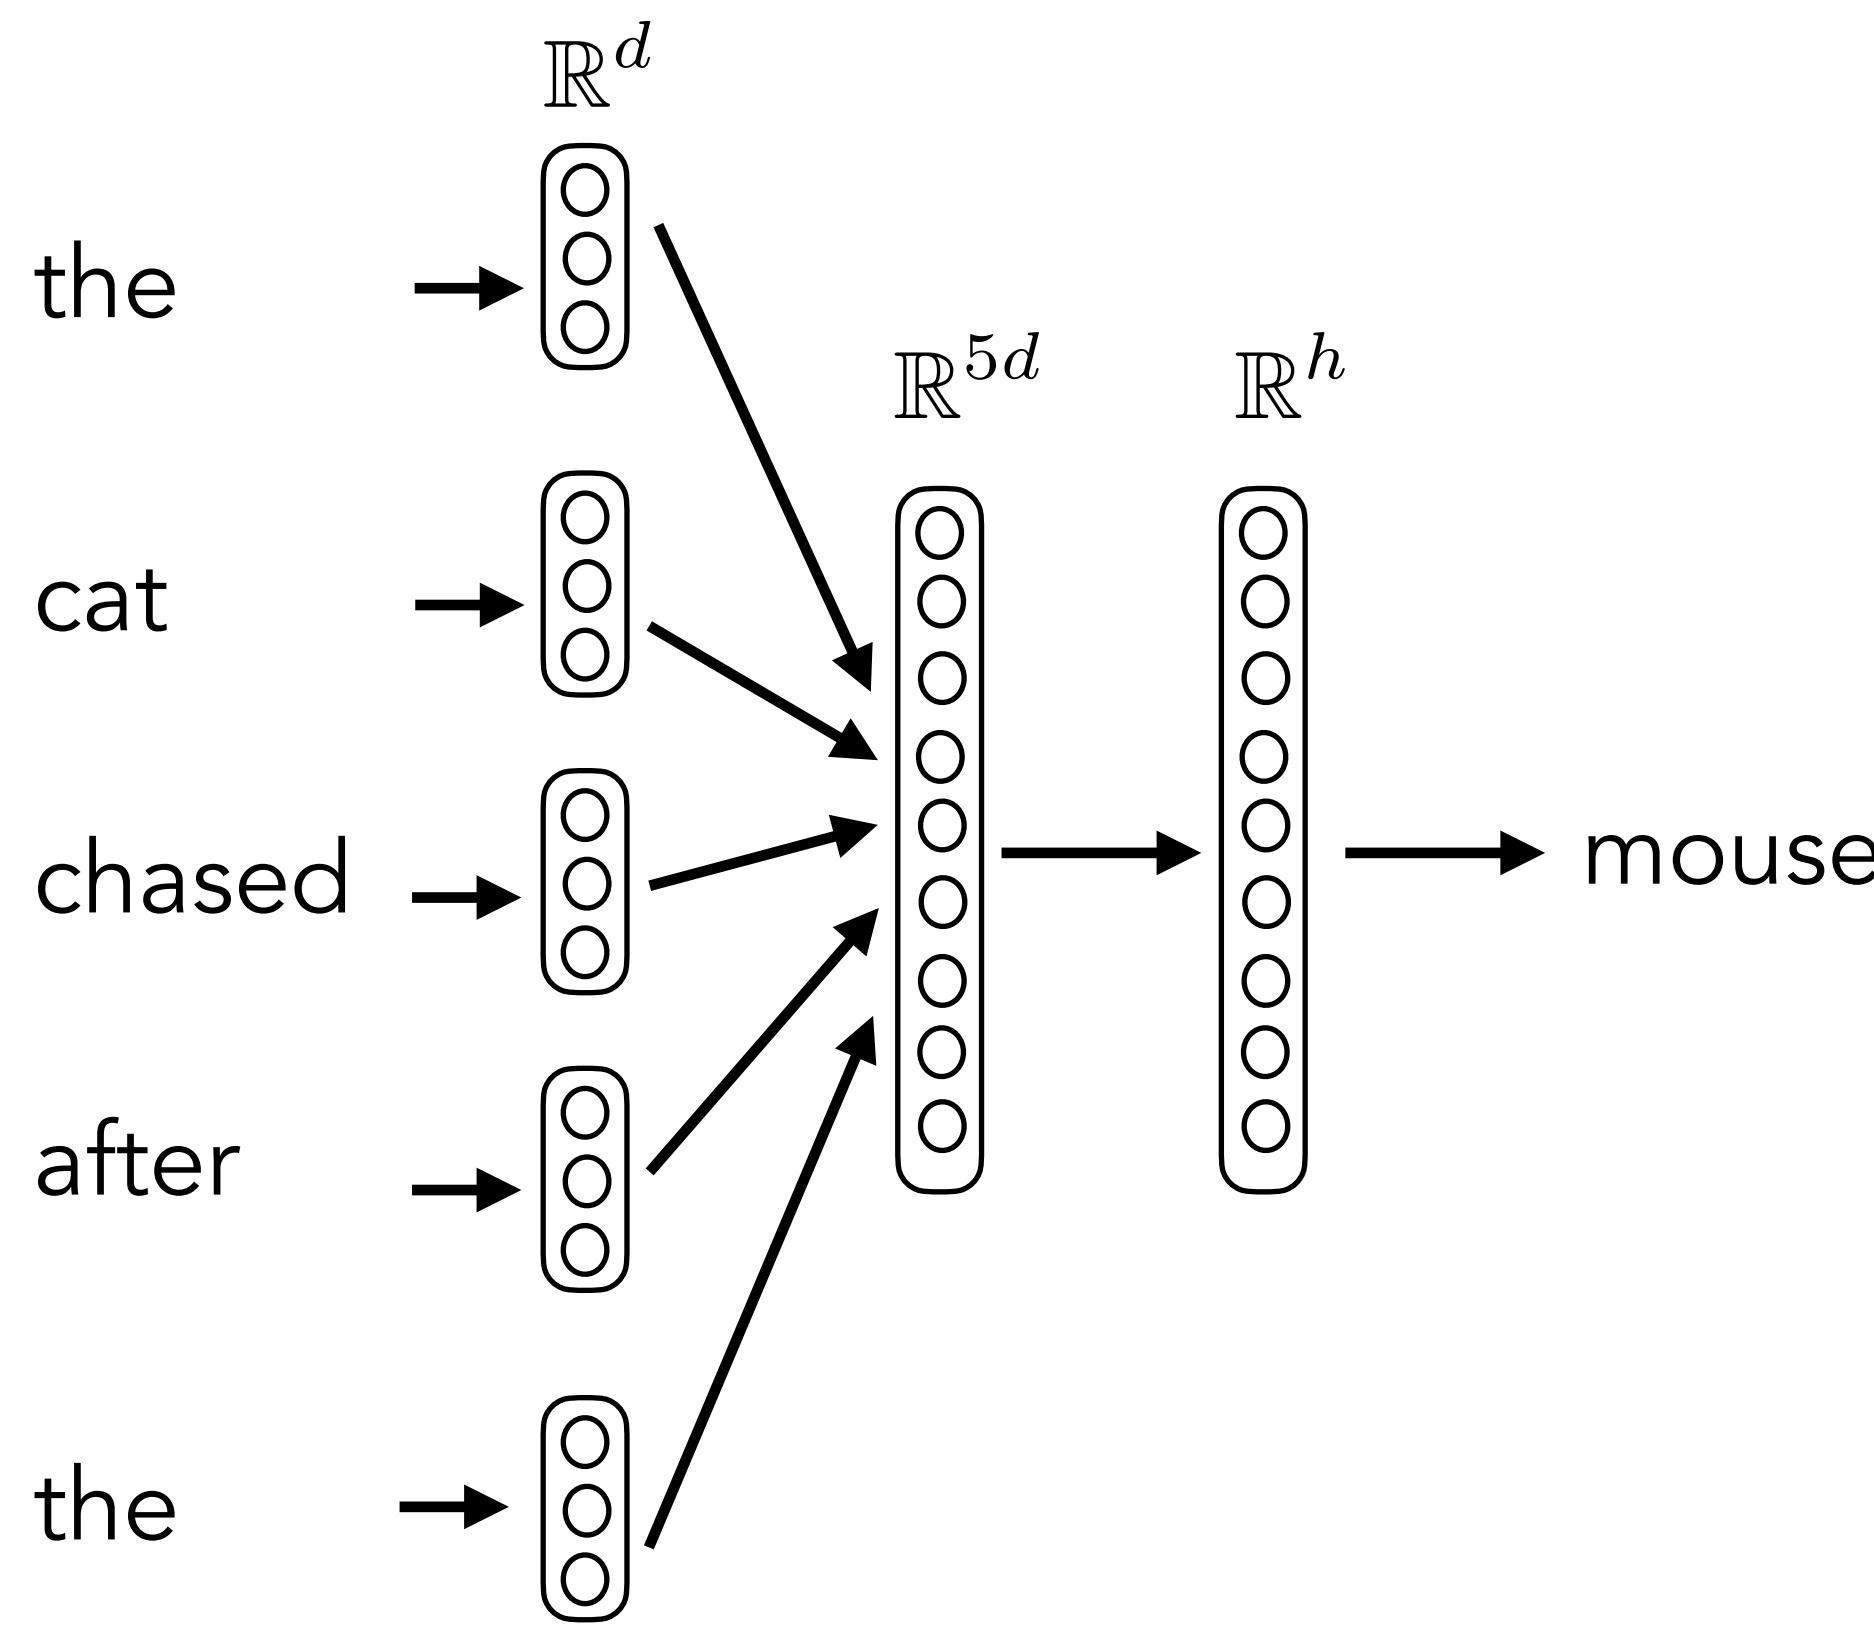
The starving **cat** fanatically chased the **mouse**

The starving **cat** fanatically chased the elusive **mouse**

The starving **cat**, who had not eaten in six days, fanatically chased the elusive **mouse**



Fixed-context Neural Language Models



$P(\text{mouse} \mid \text{the cat chase the})$



$P(\text{mouse} \mid \text{the starving cat chased the})$



$P(\text{mouse} \mid \text{starving cat chased after the})$



$P(\text{mouse} \mid \text{cat fanatically chased after the})$

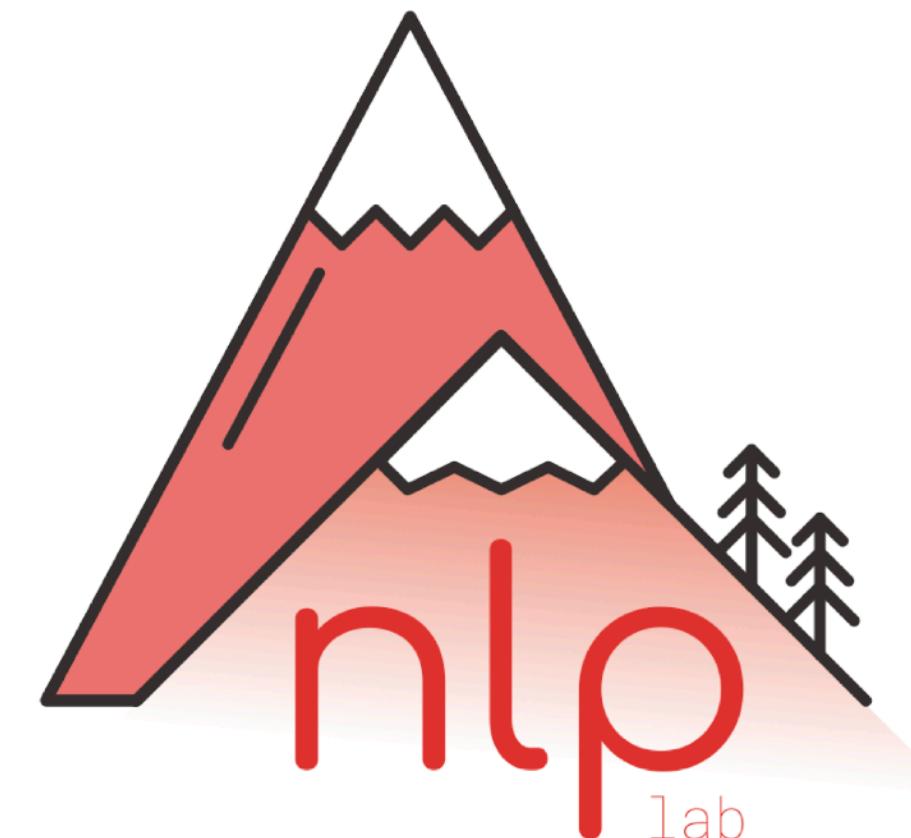


$P(\text{mouse} \mid \text{fanatically chased after the elusive})$



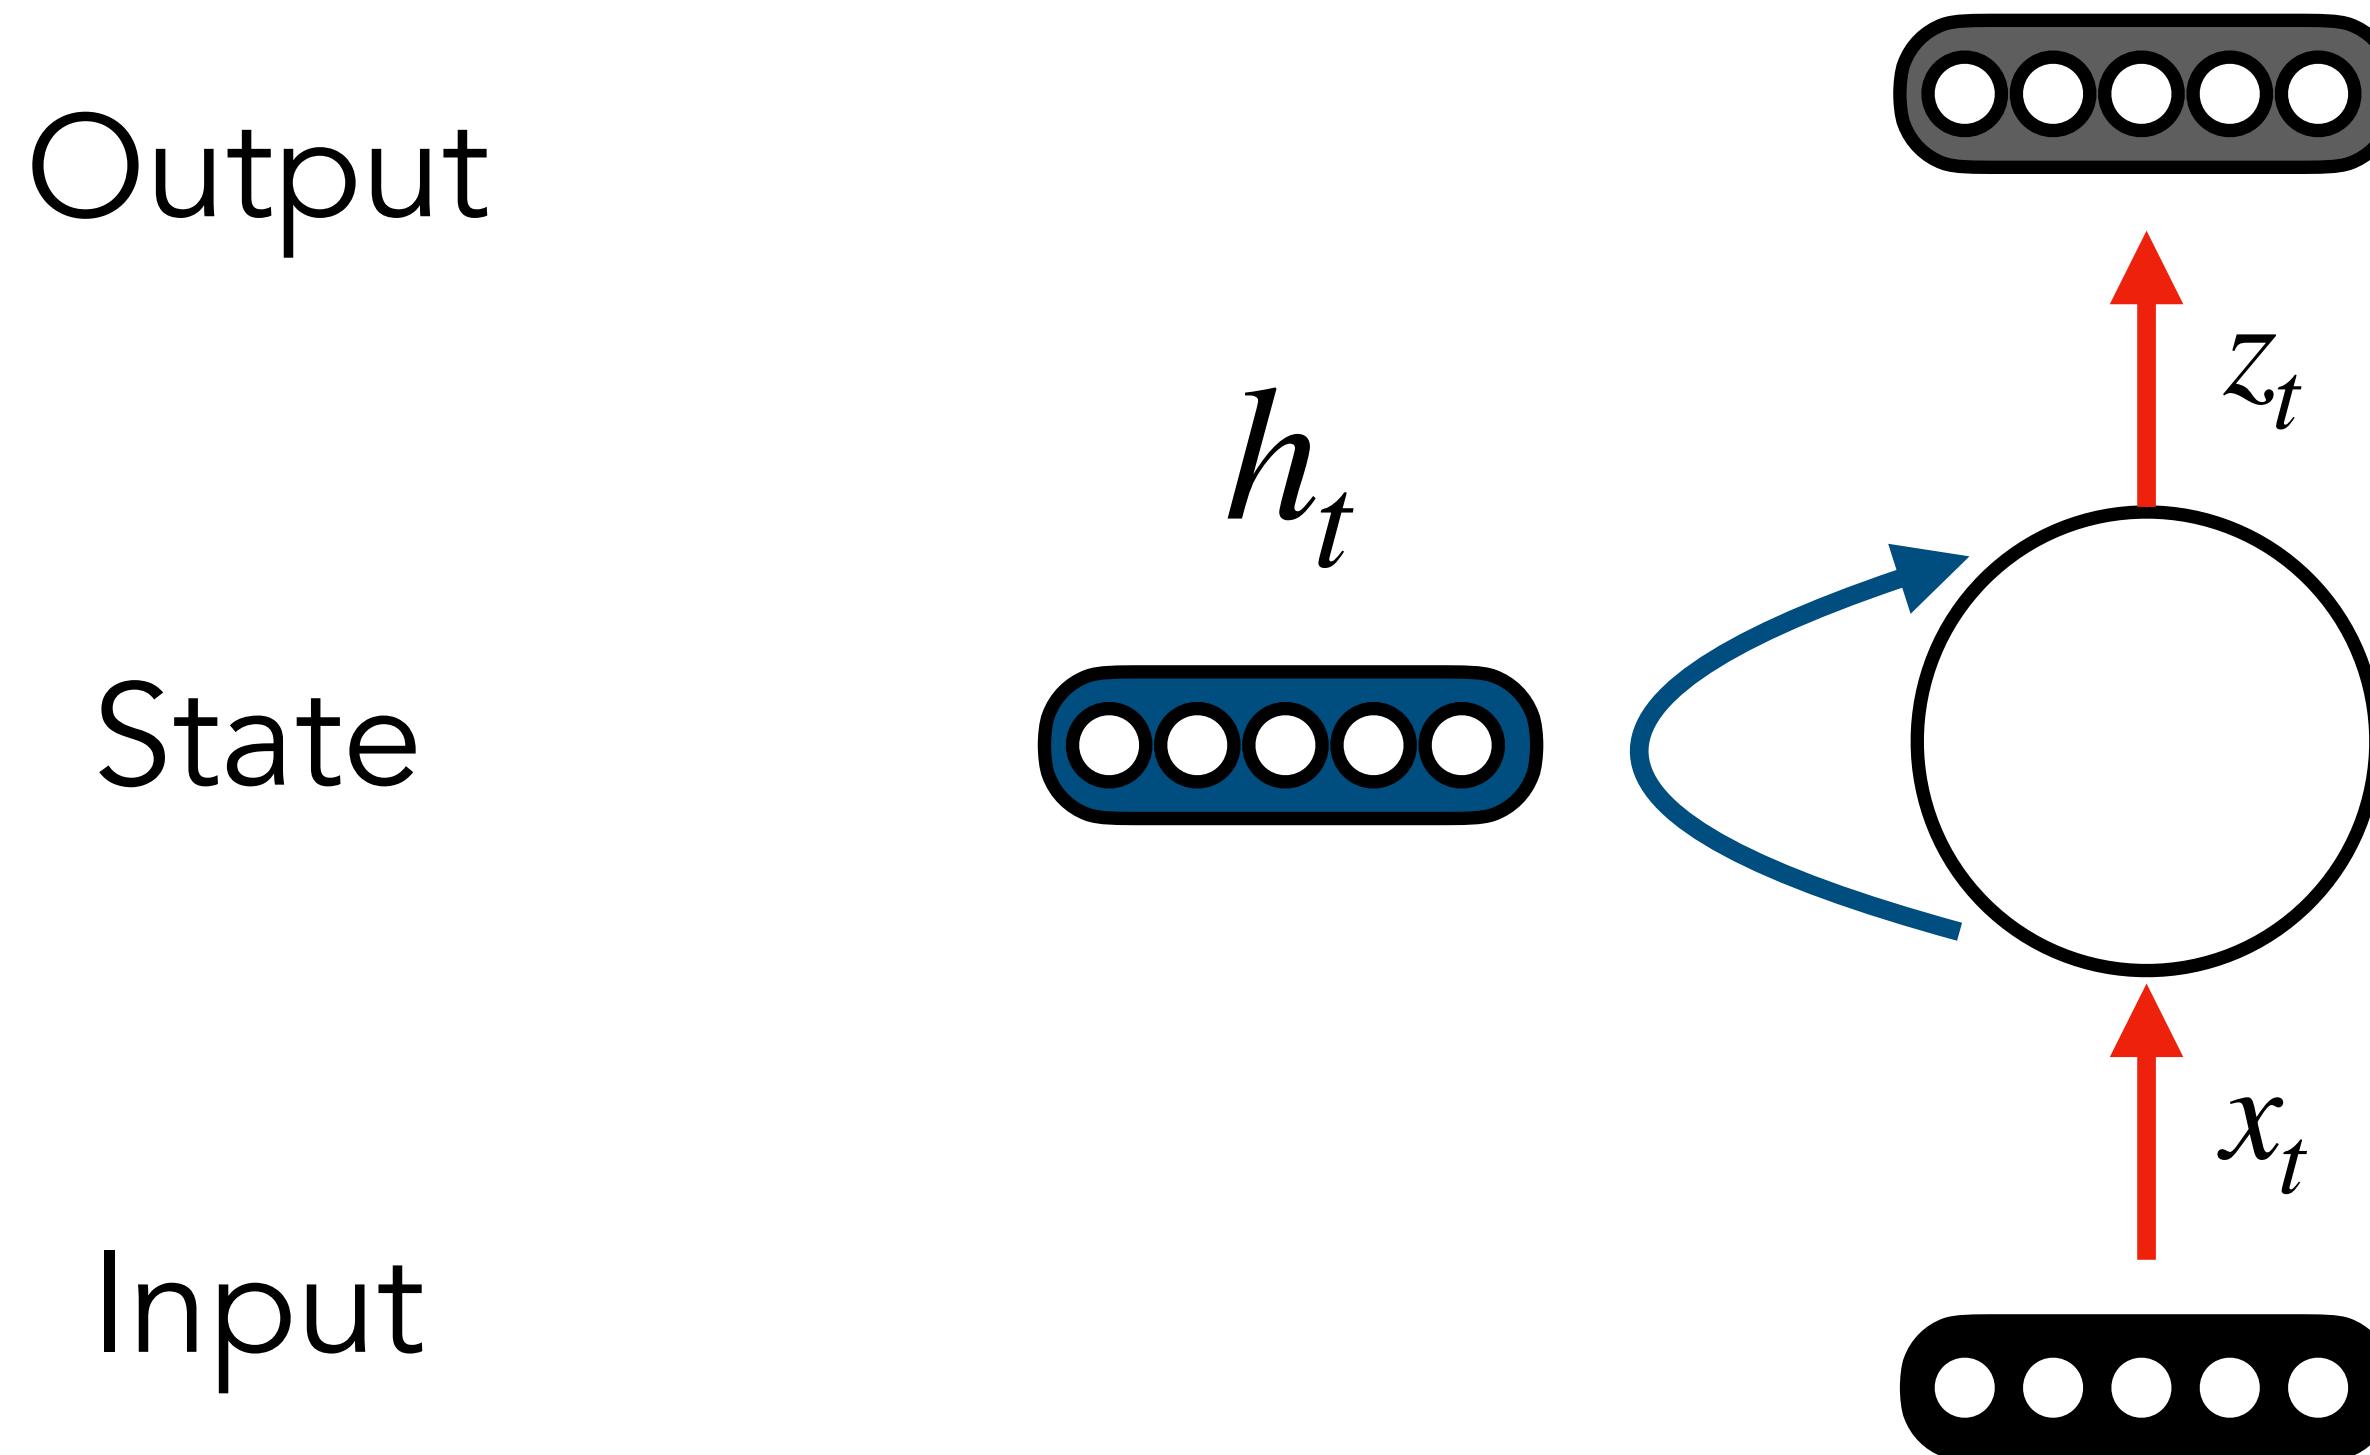
Language Models: Recurrent Neural Networks

Antoine Bosselut



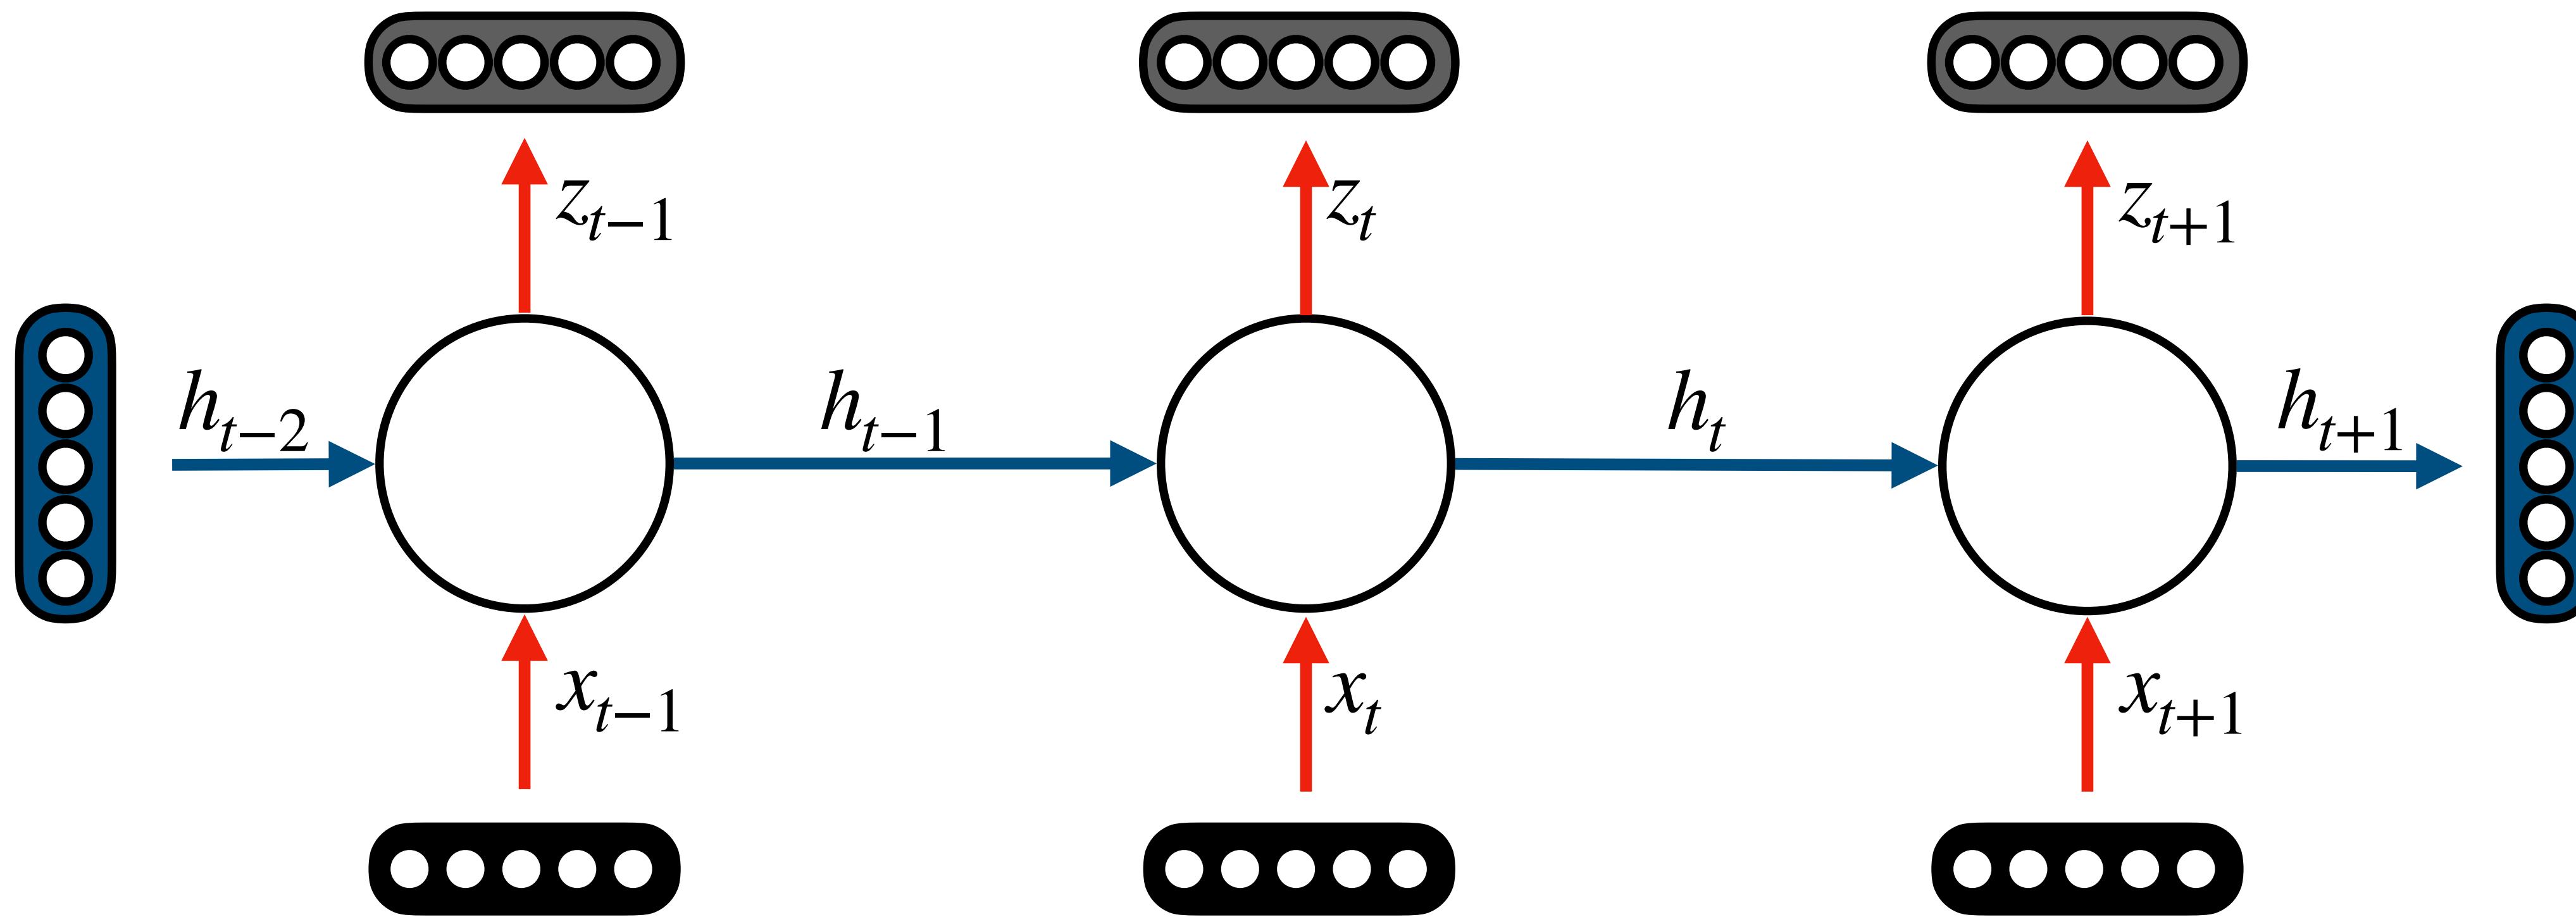
Recurrent Neural Networks

- **Solution:** Recurrent neural networks — NNs with feedback loops



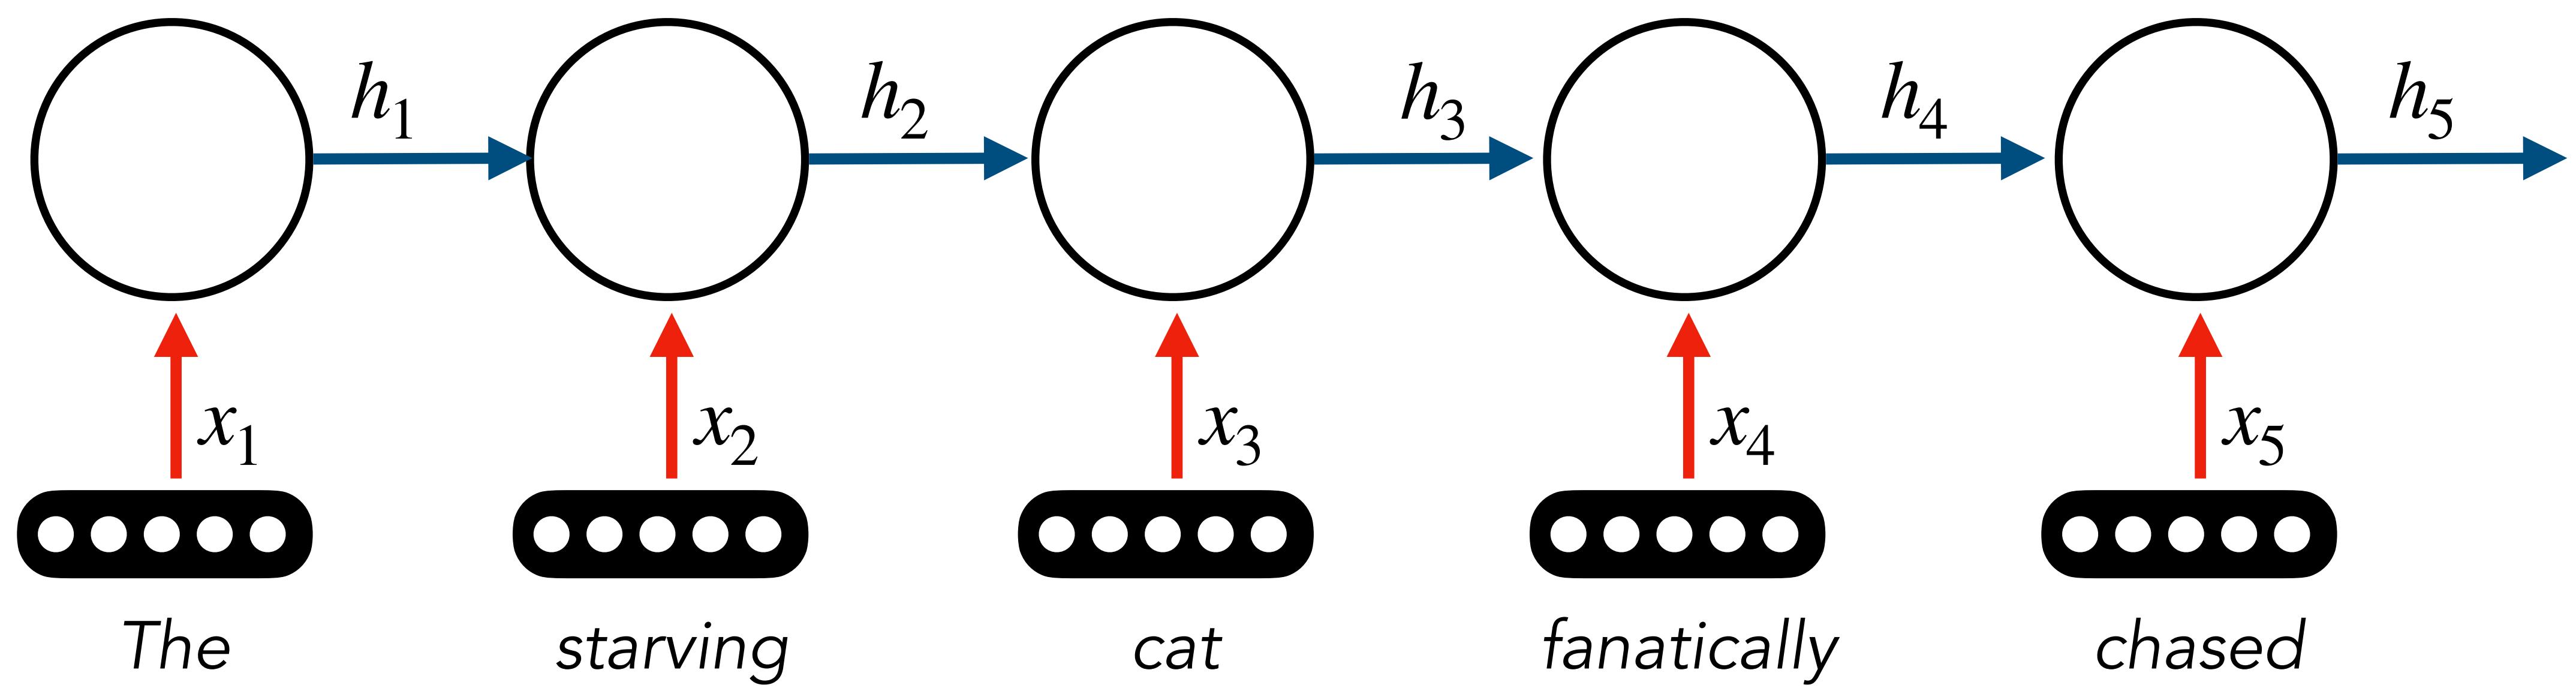
Unrolling the RNN

Unrolling the RNN across all time steps gives full computation graph

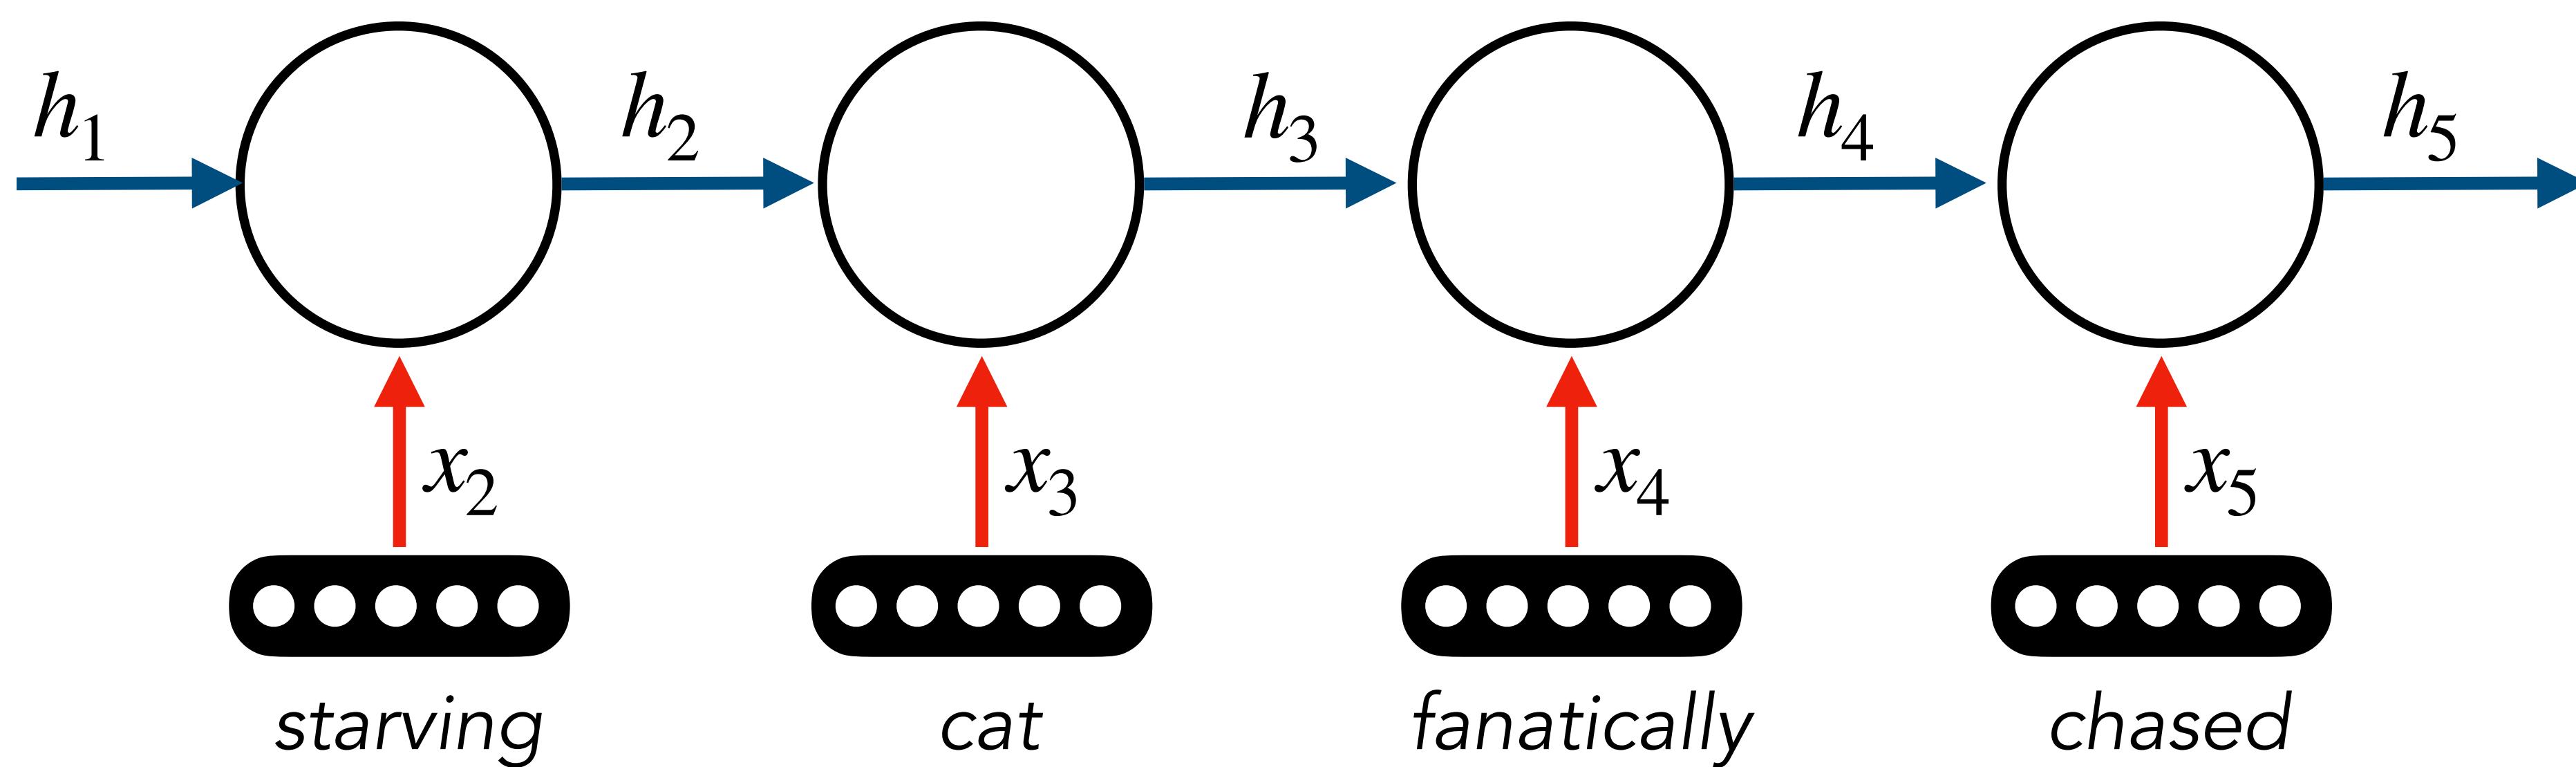


Allows for learning from entire sequence history, regardless of length

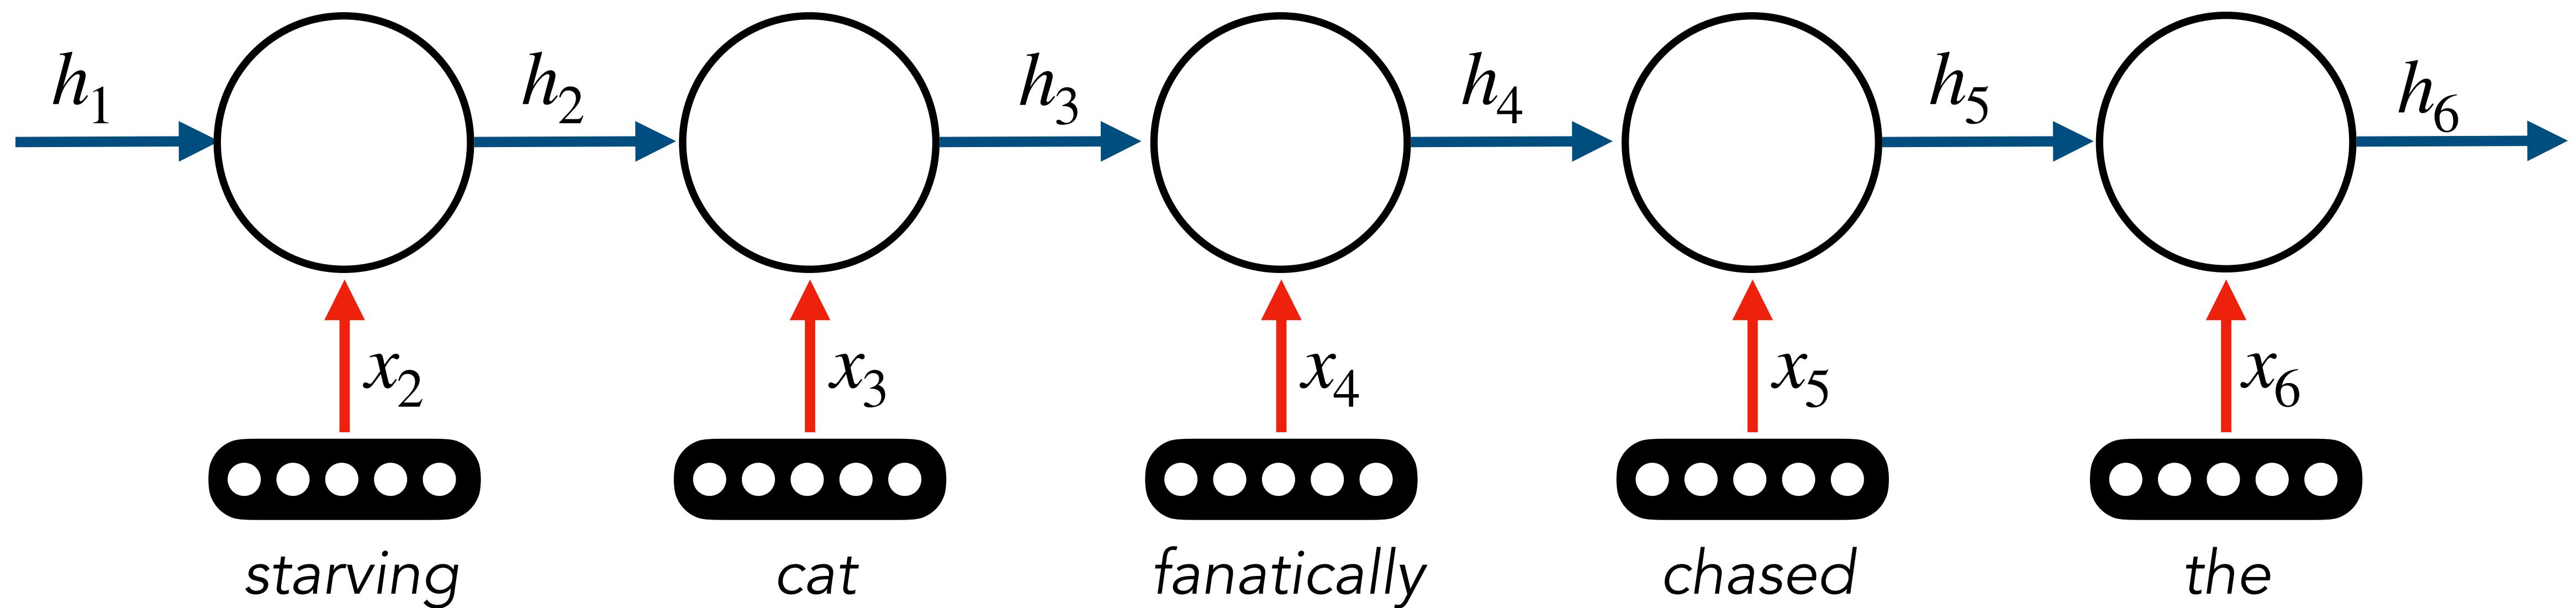
Unrolling the RNN



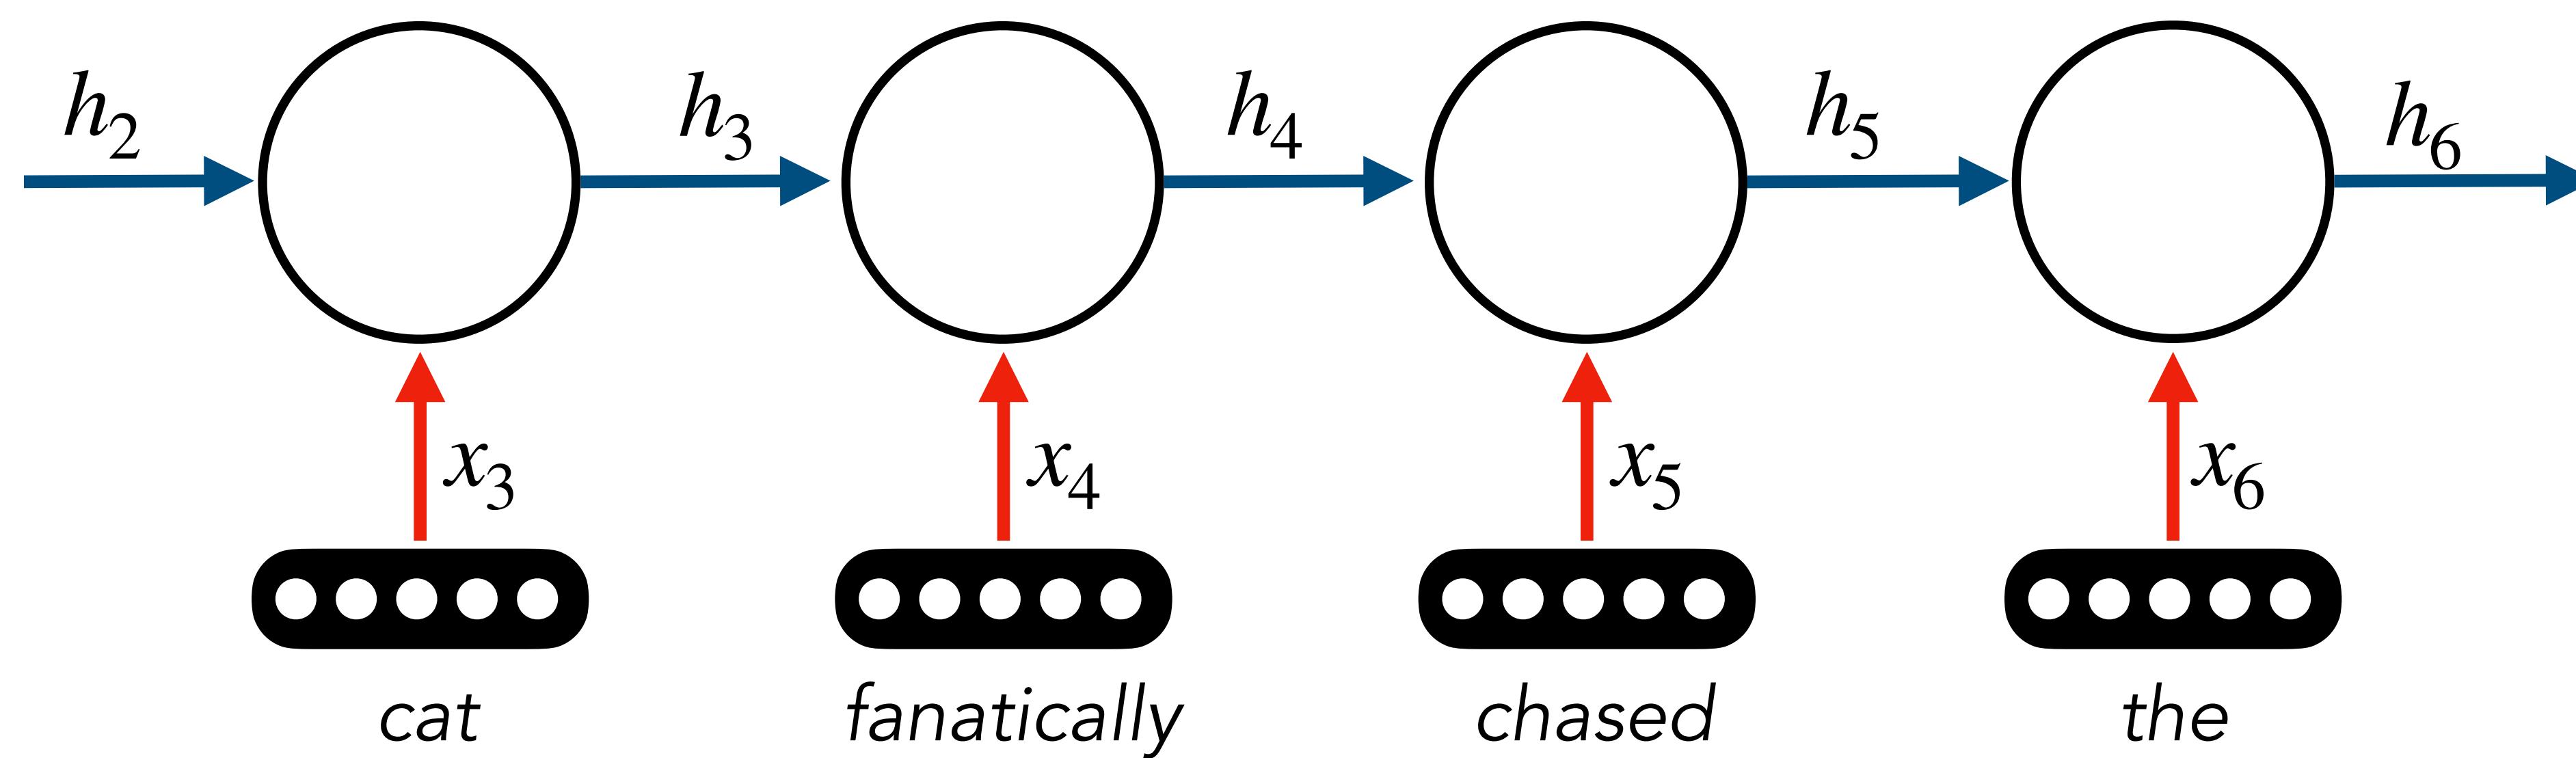
Unrolling the RNN



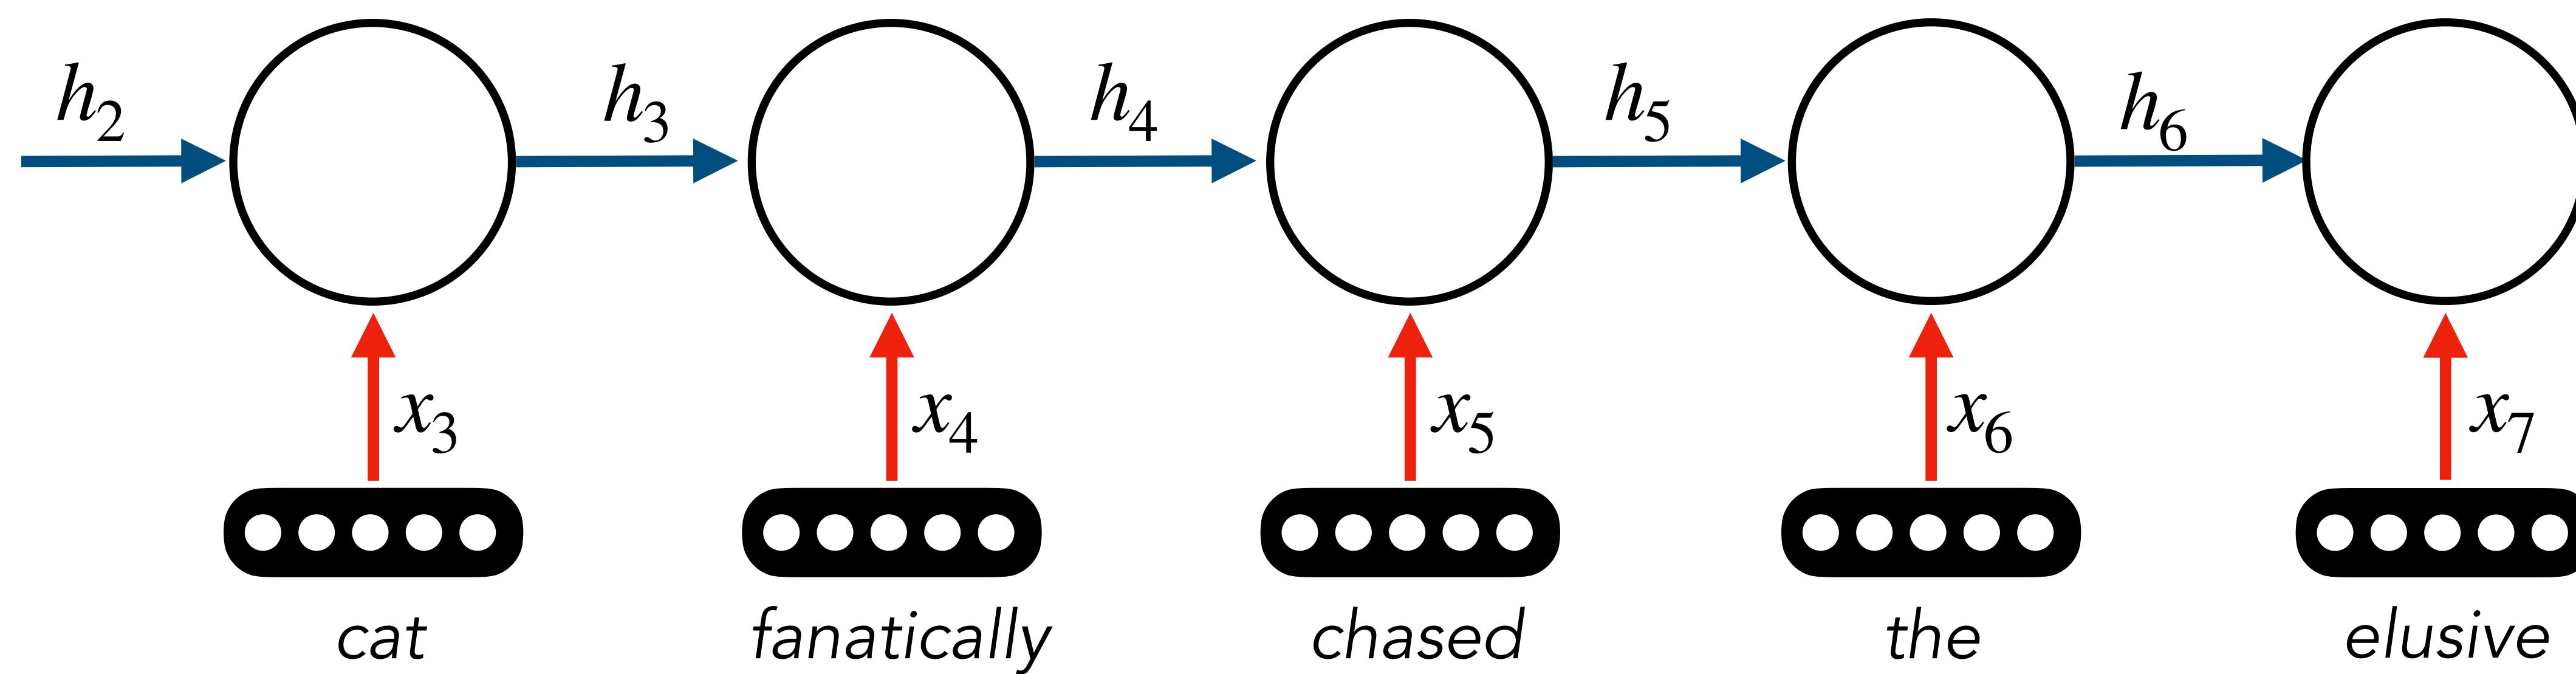
Unrolling the RNN



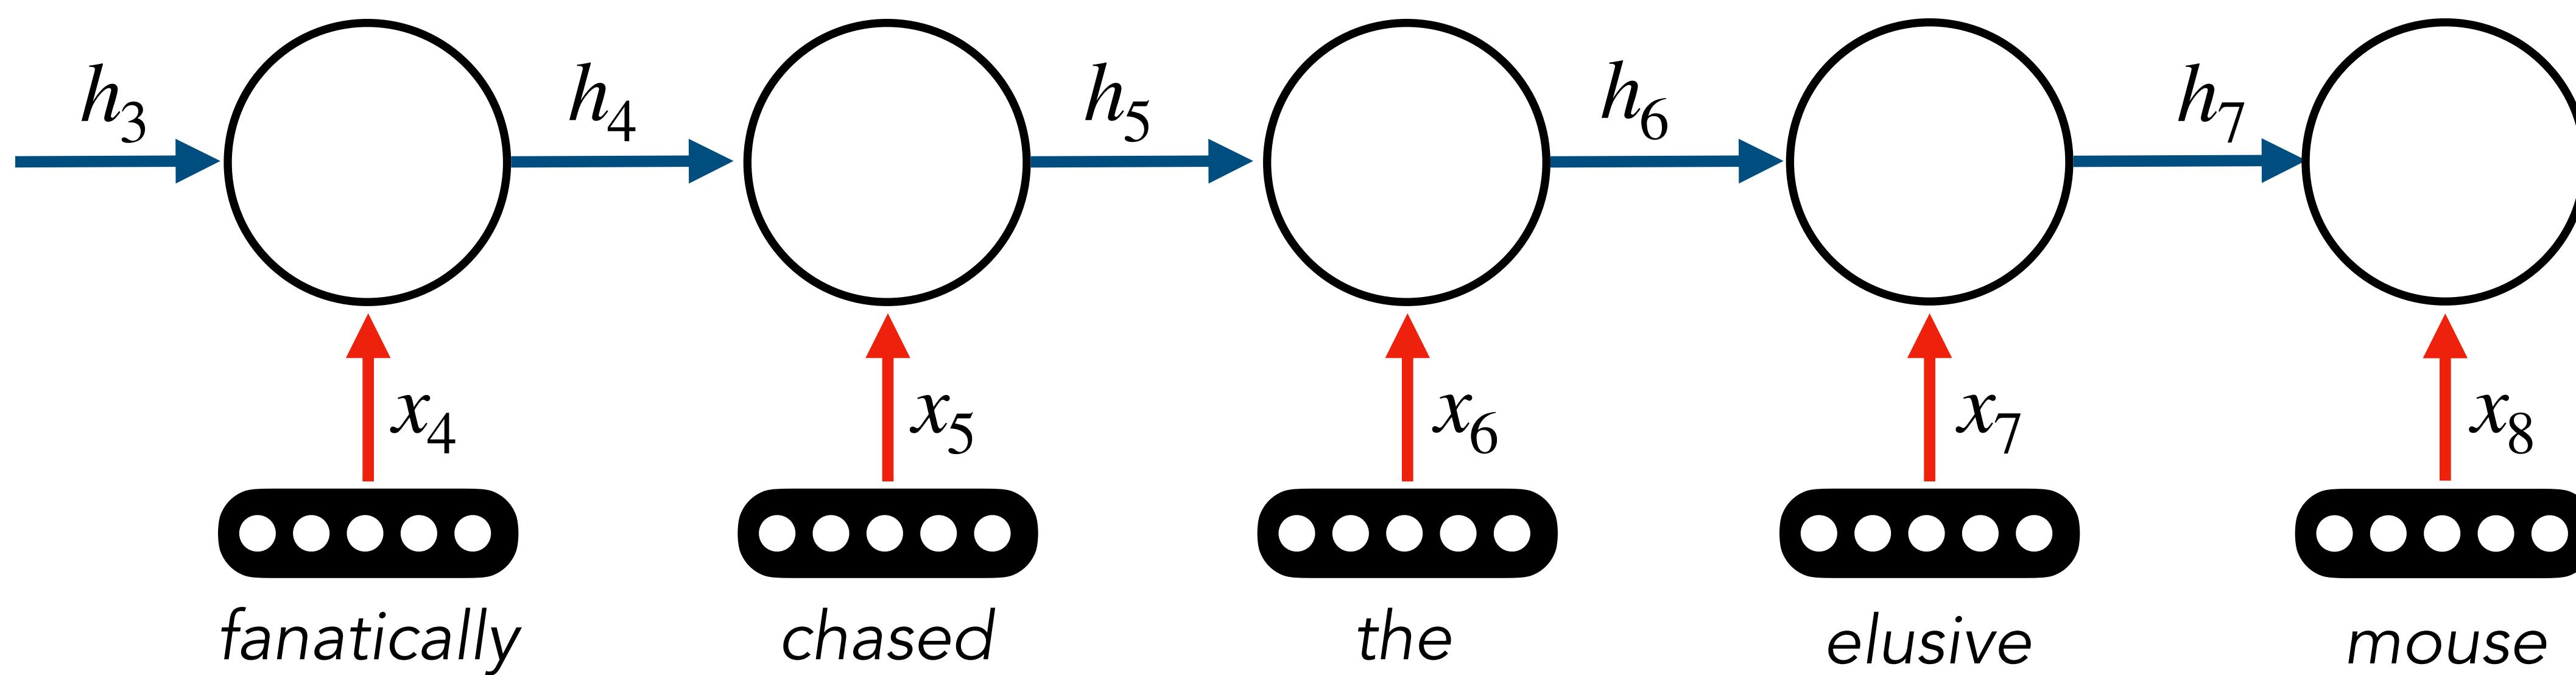
Unrolling the RNN



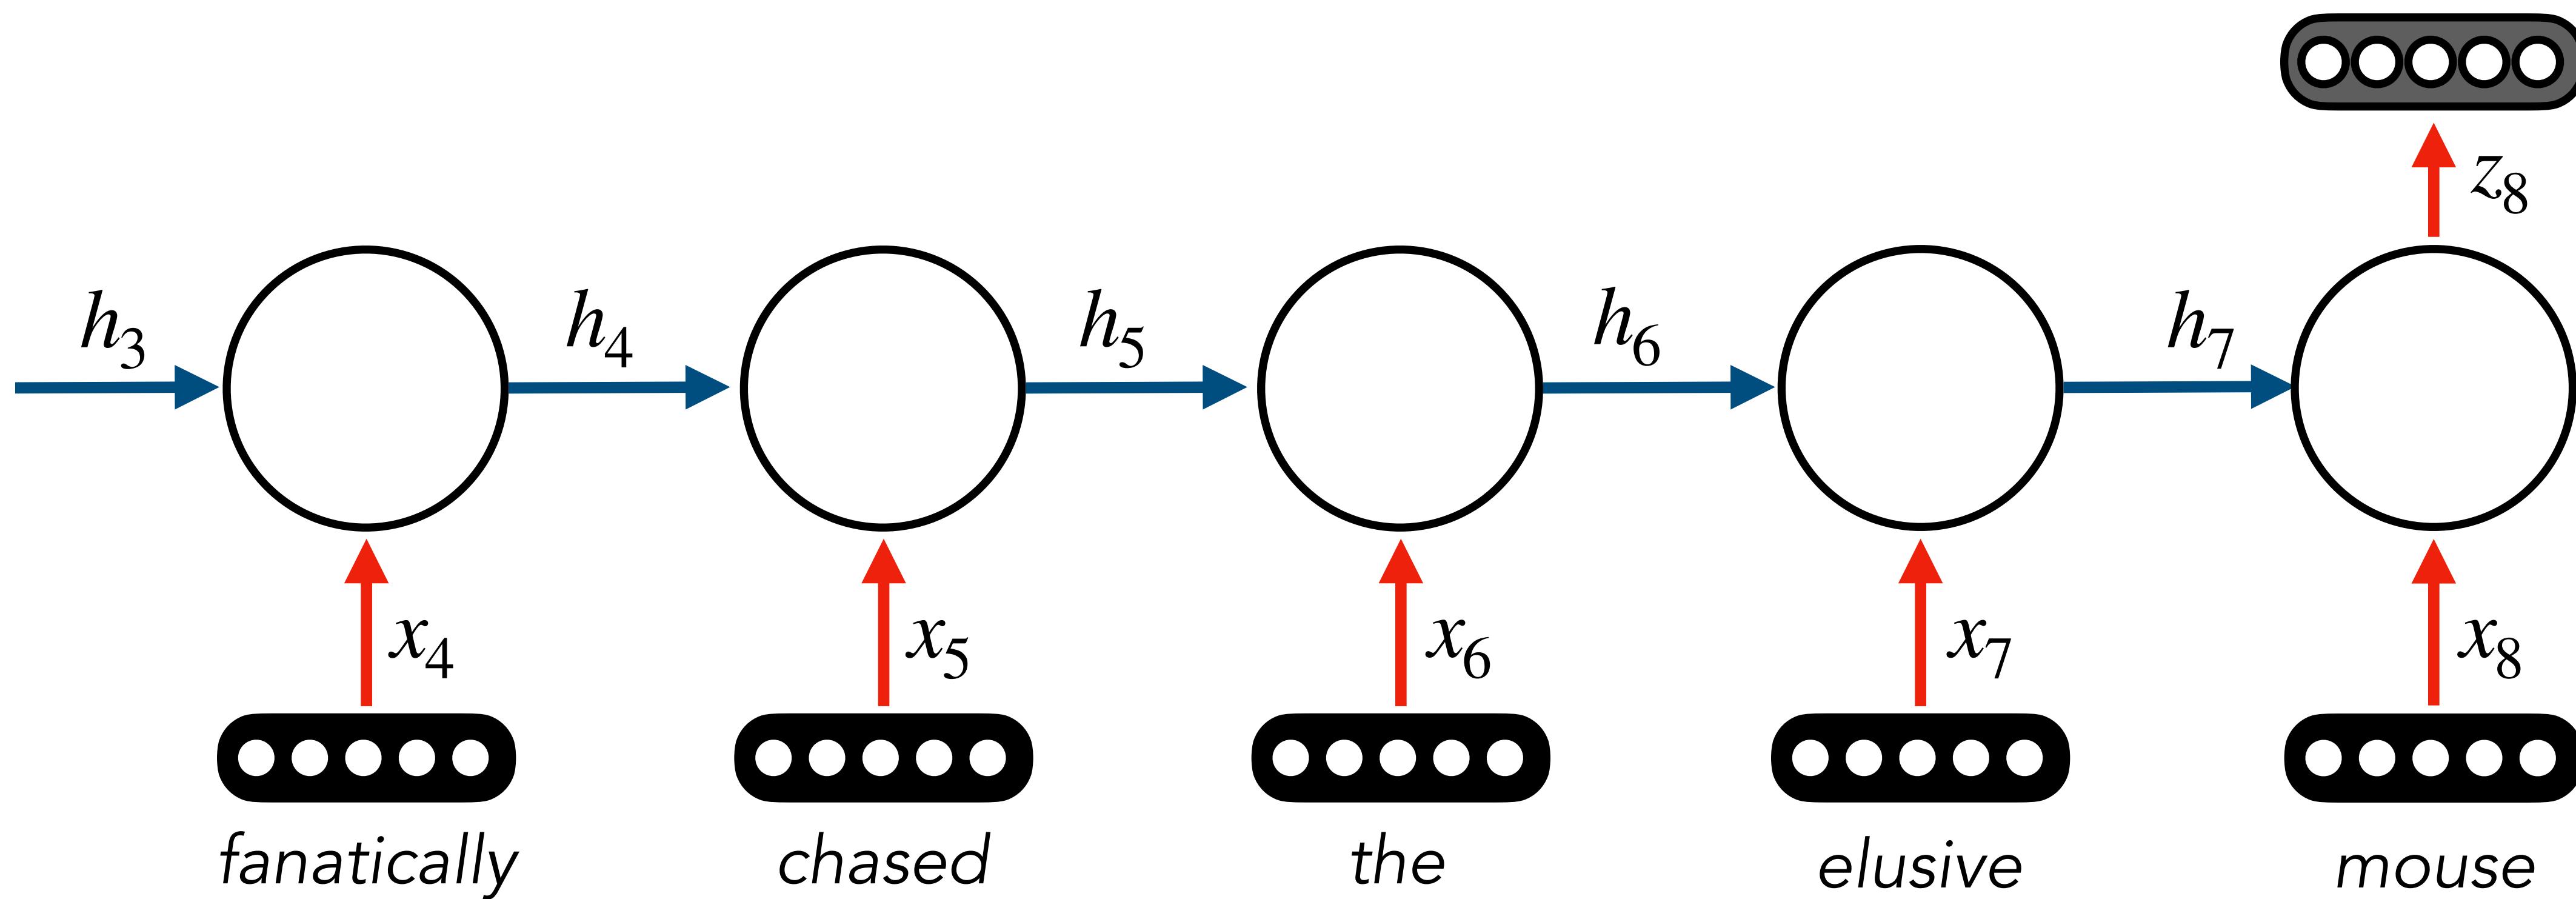
Unrolling the RNN



Unrolling the RNN

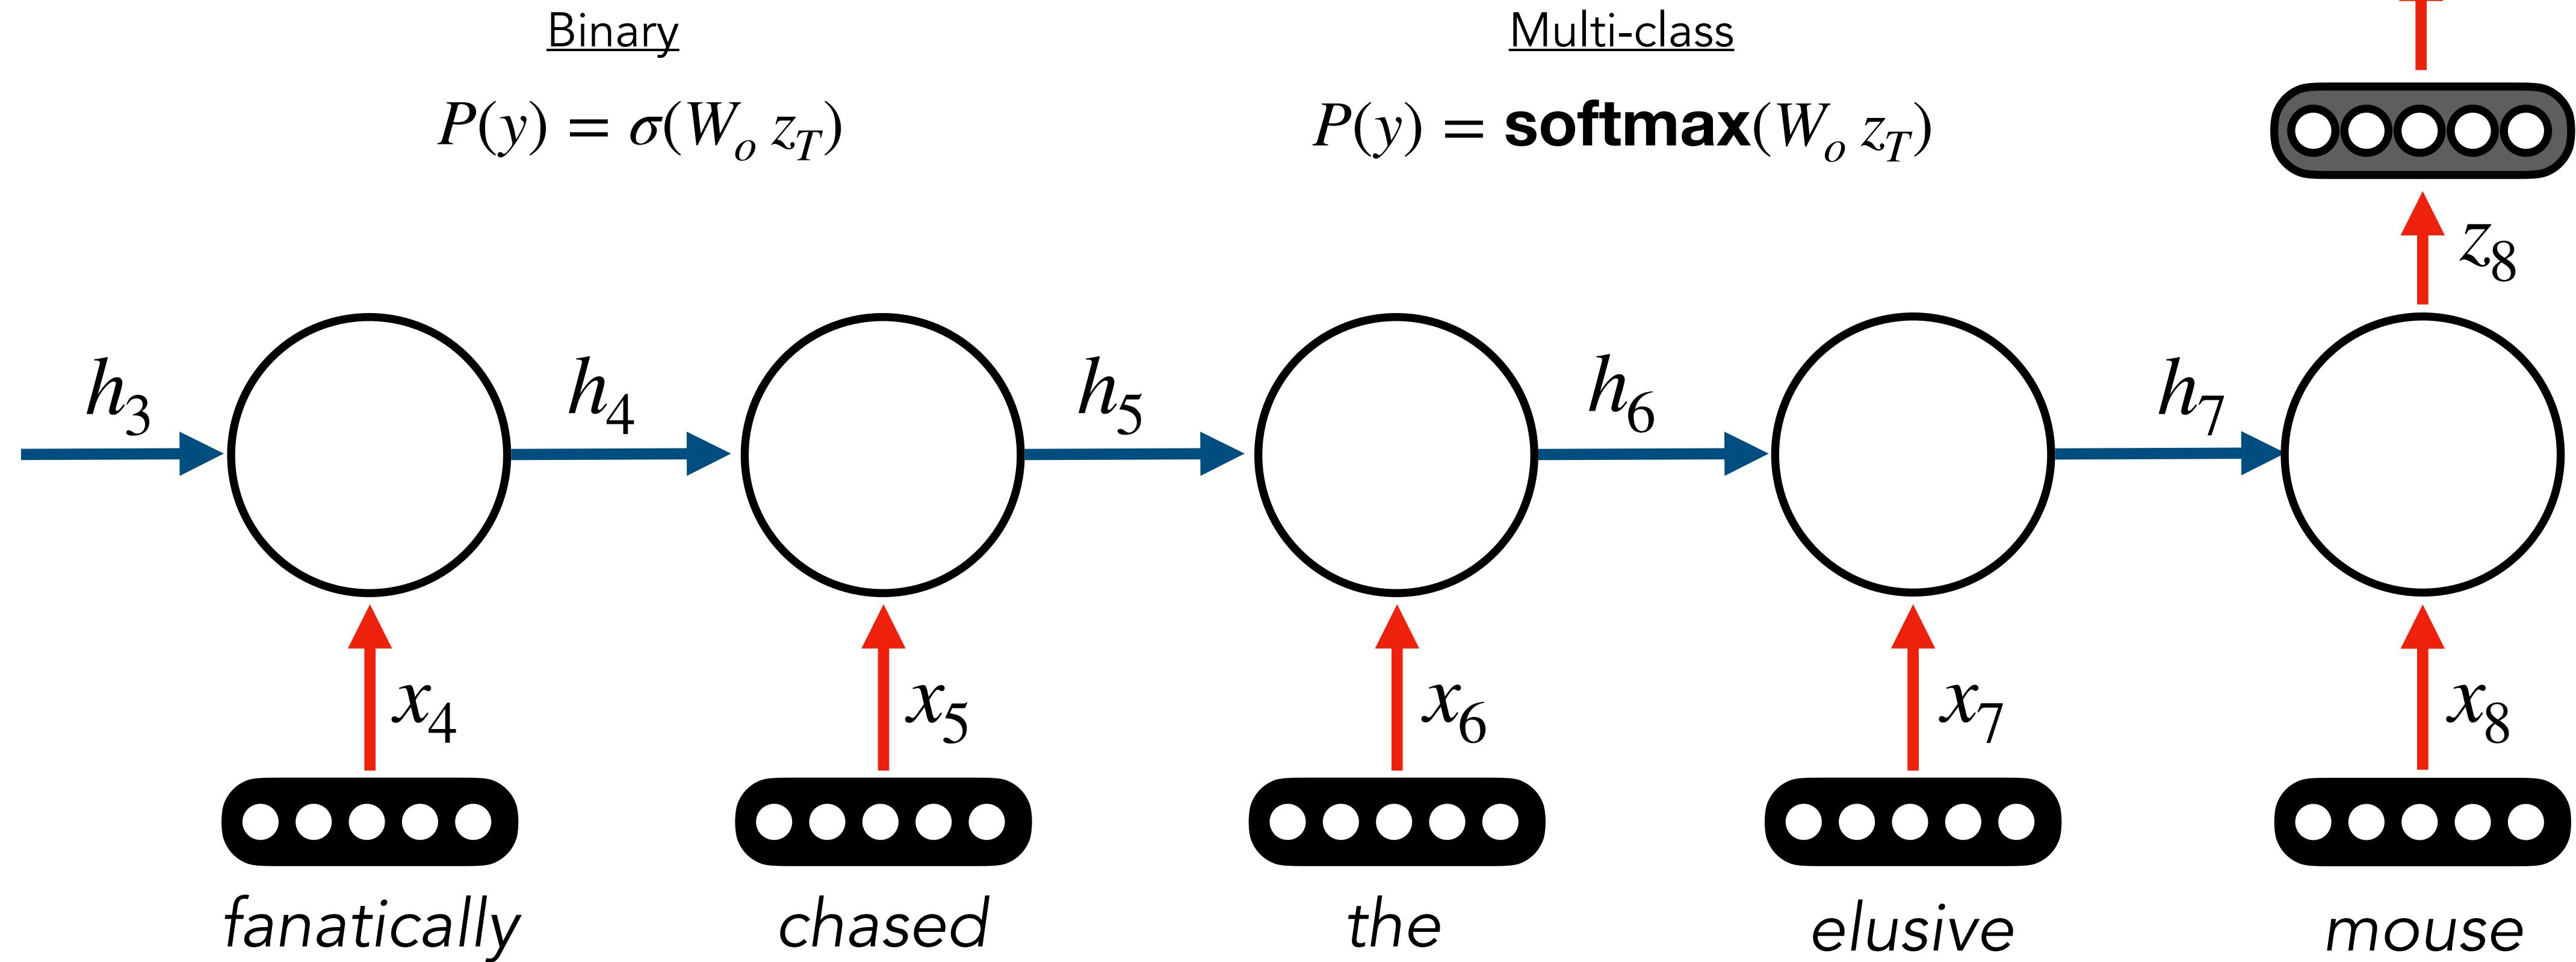


Unrolling the RNN



Classification

- Classifier is just an output projection followed by a softmax!

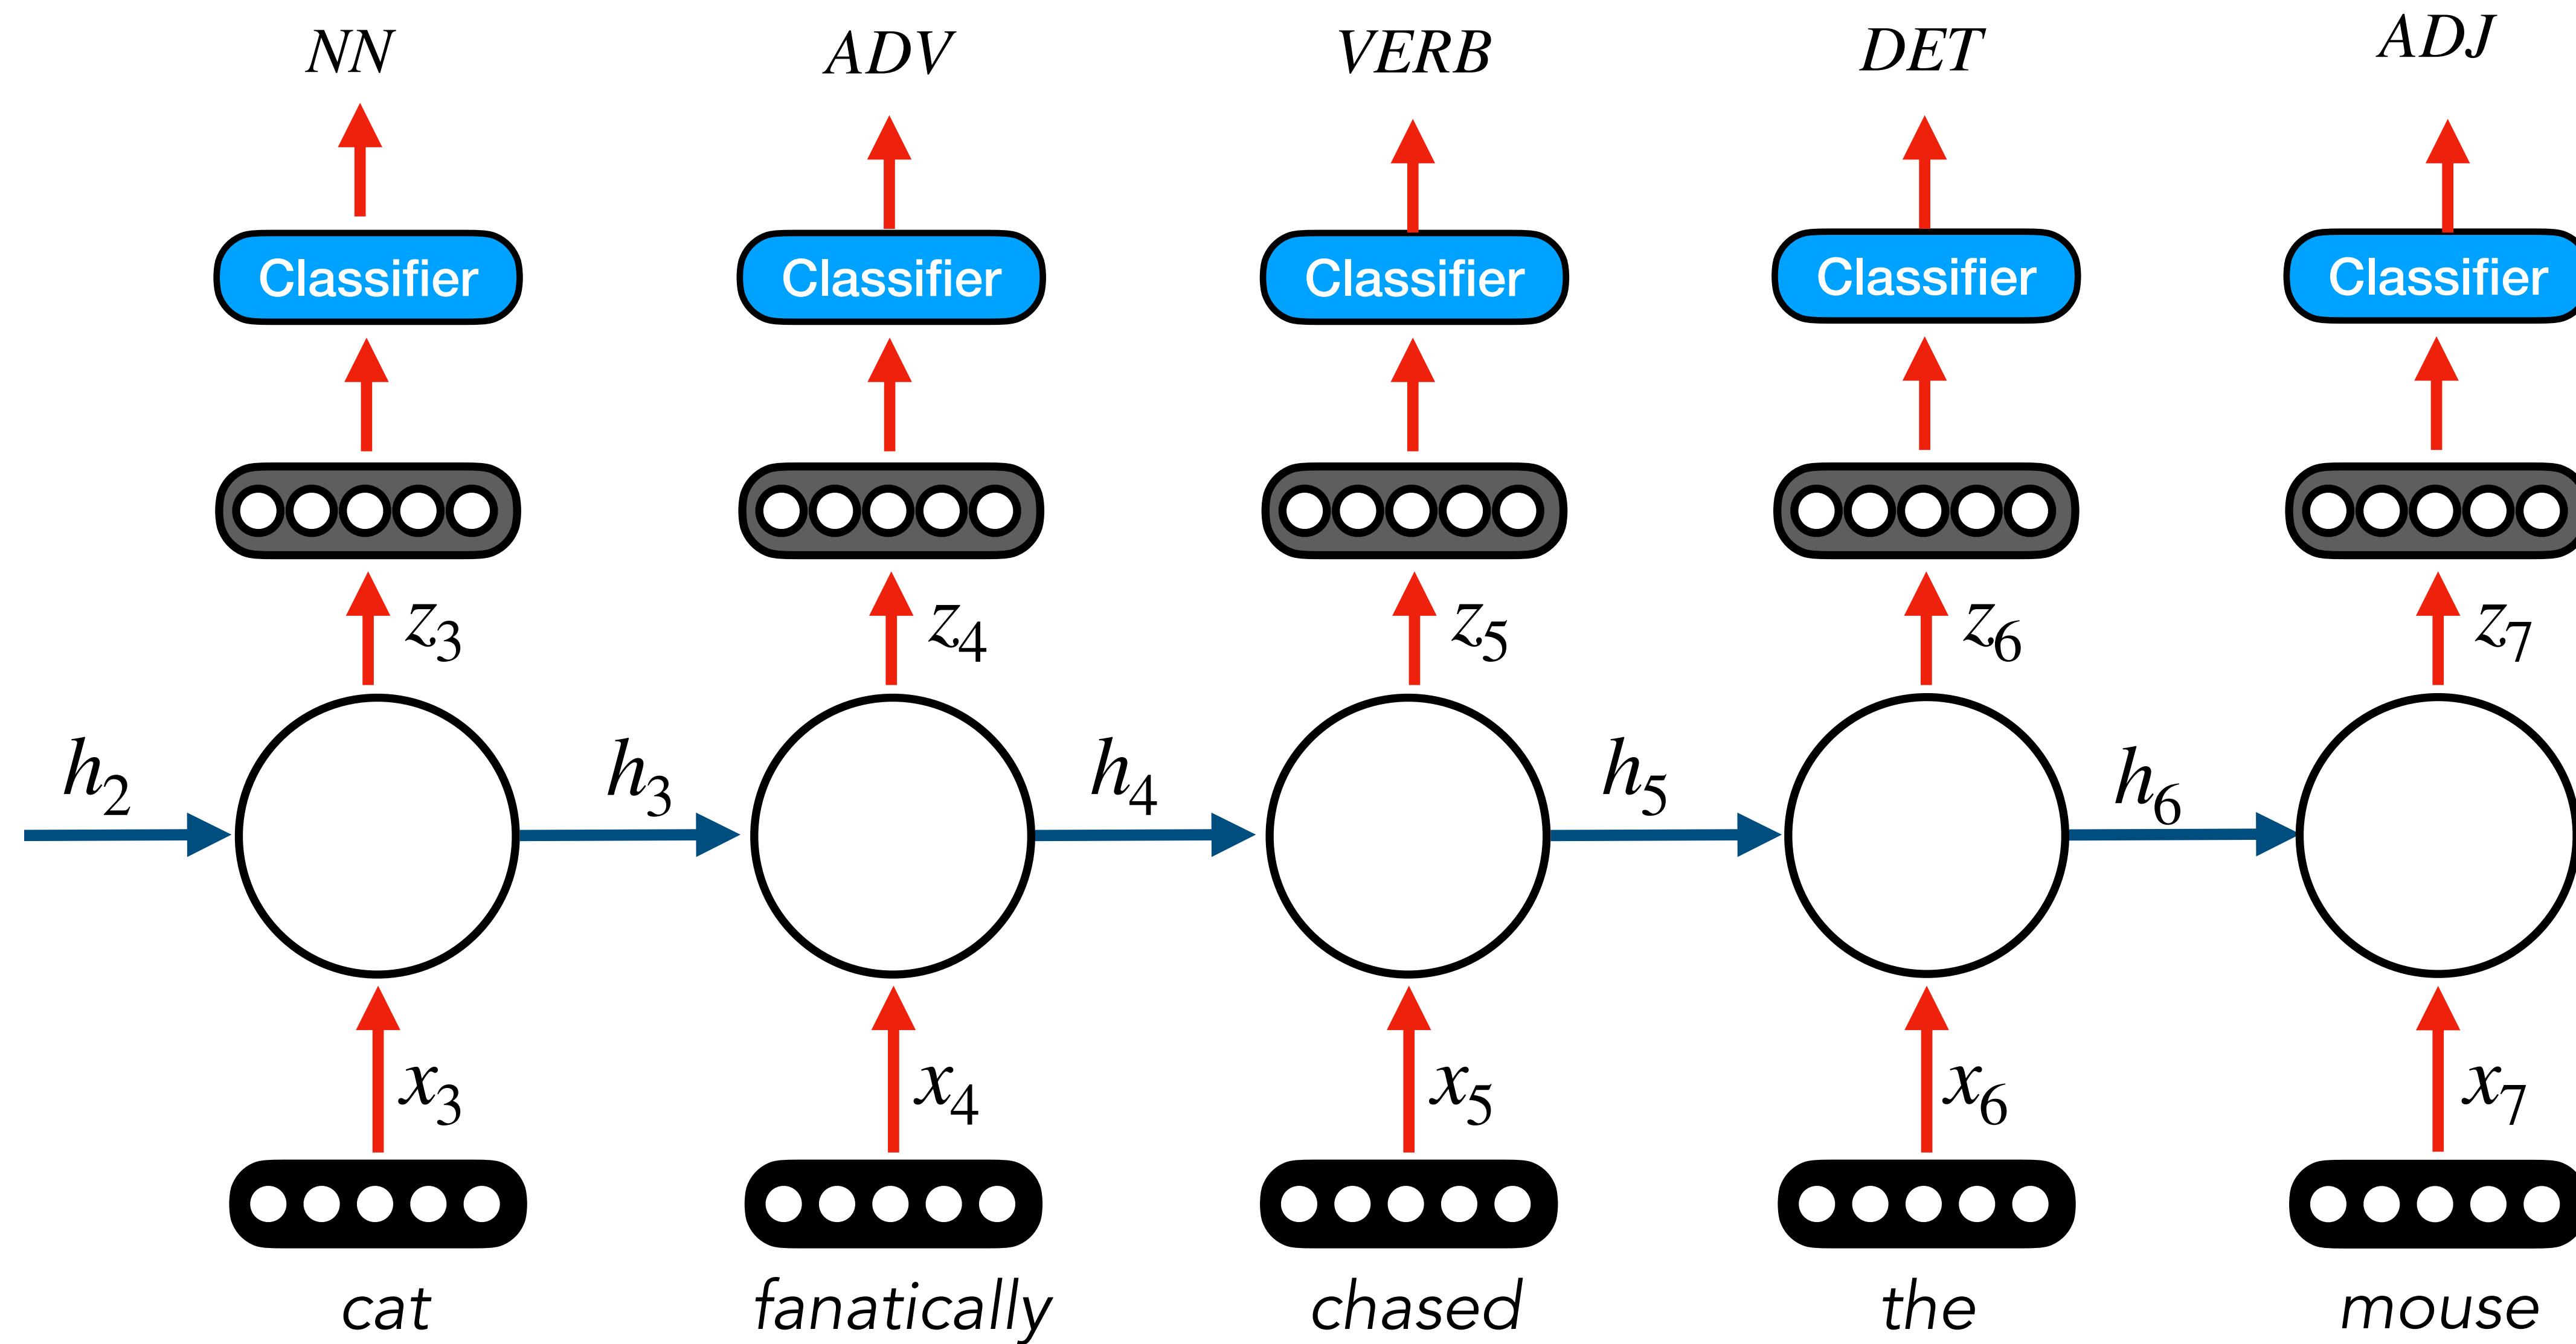


Question

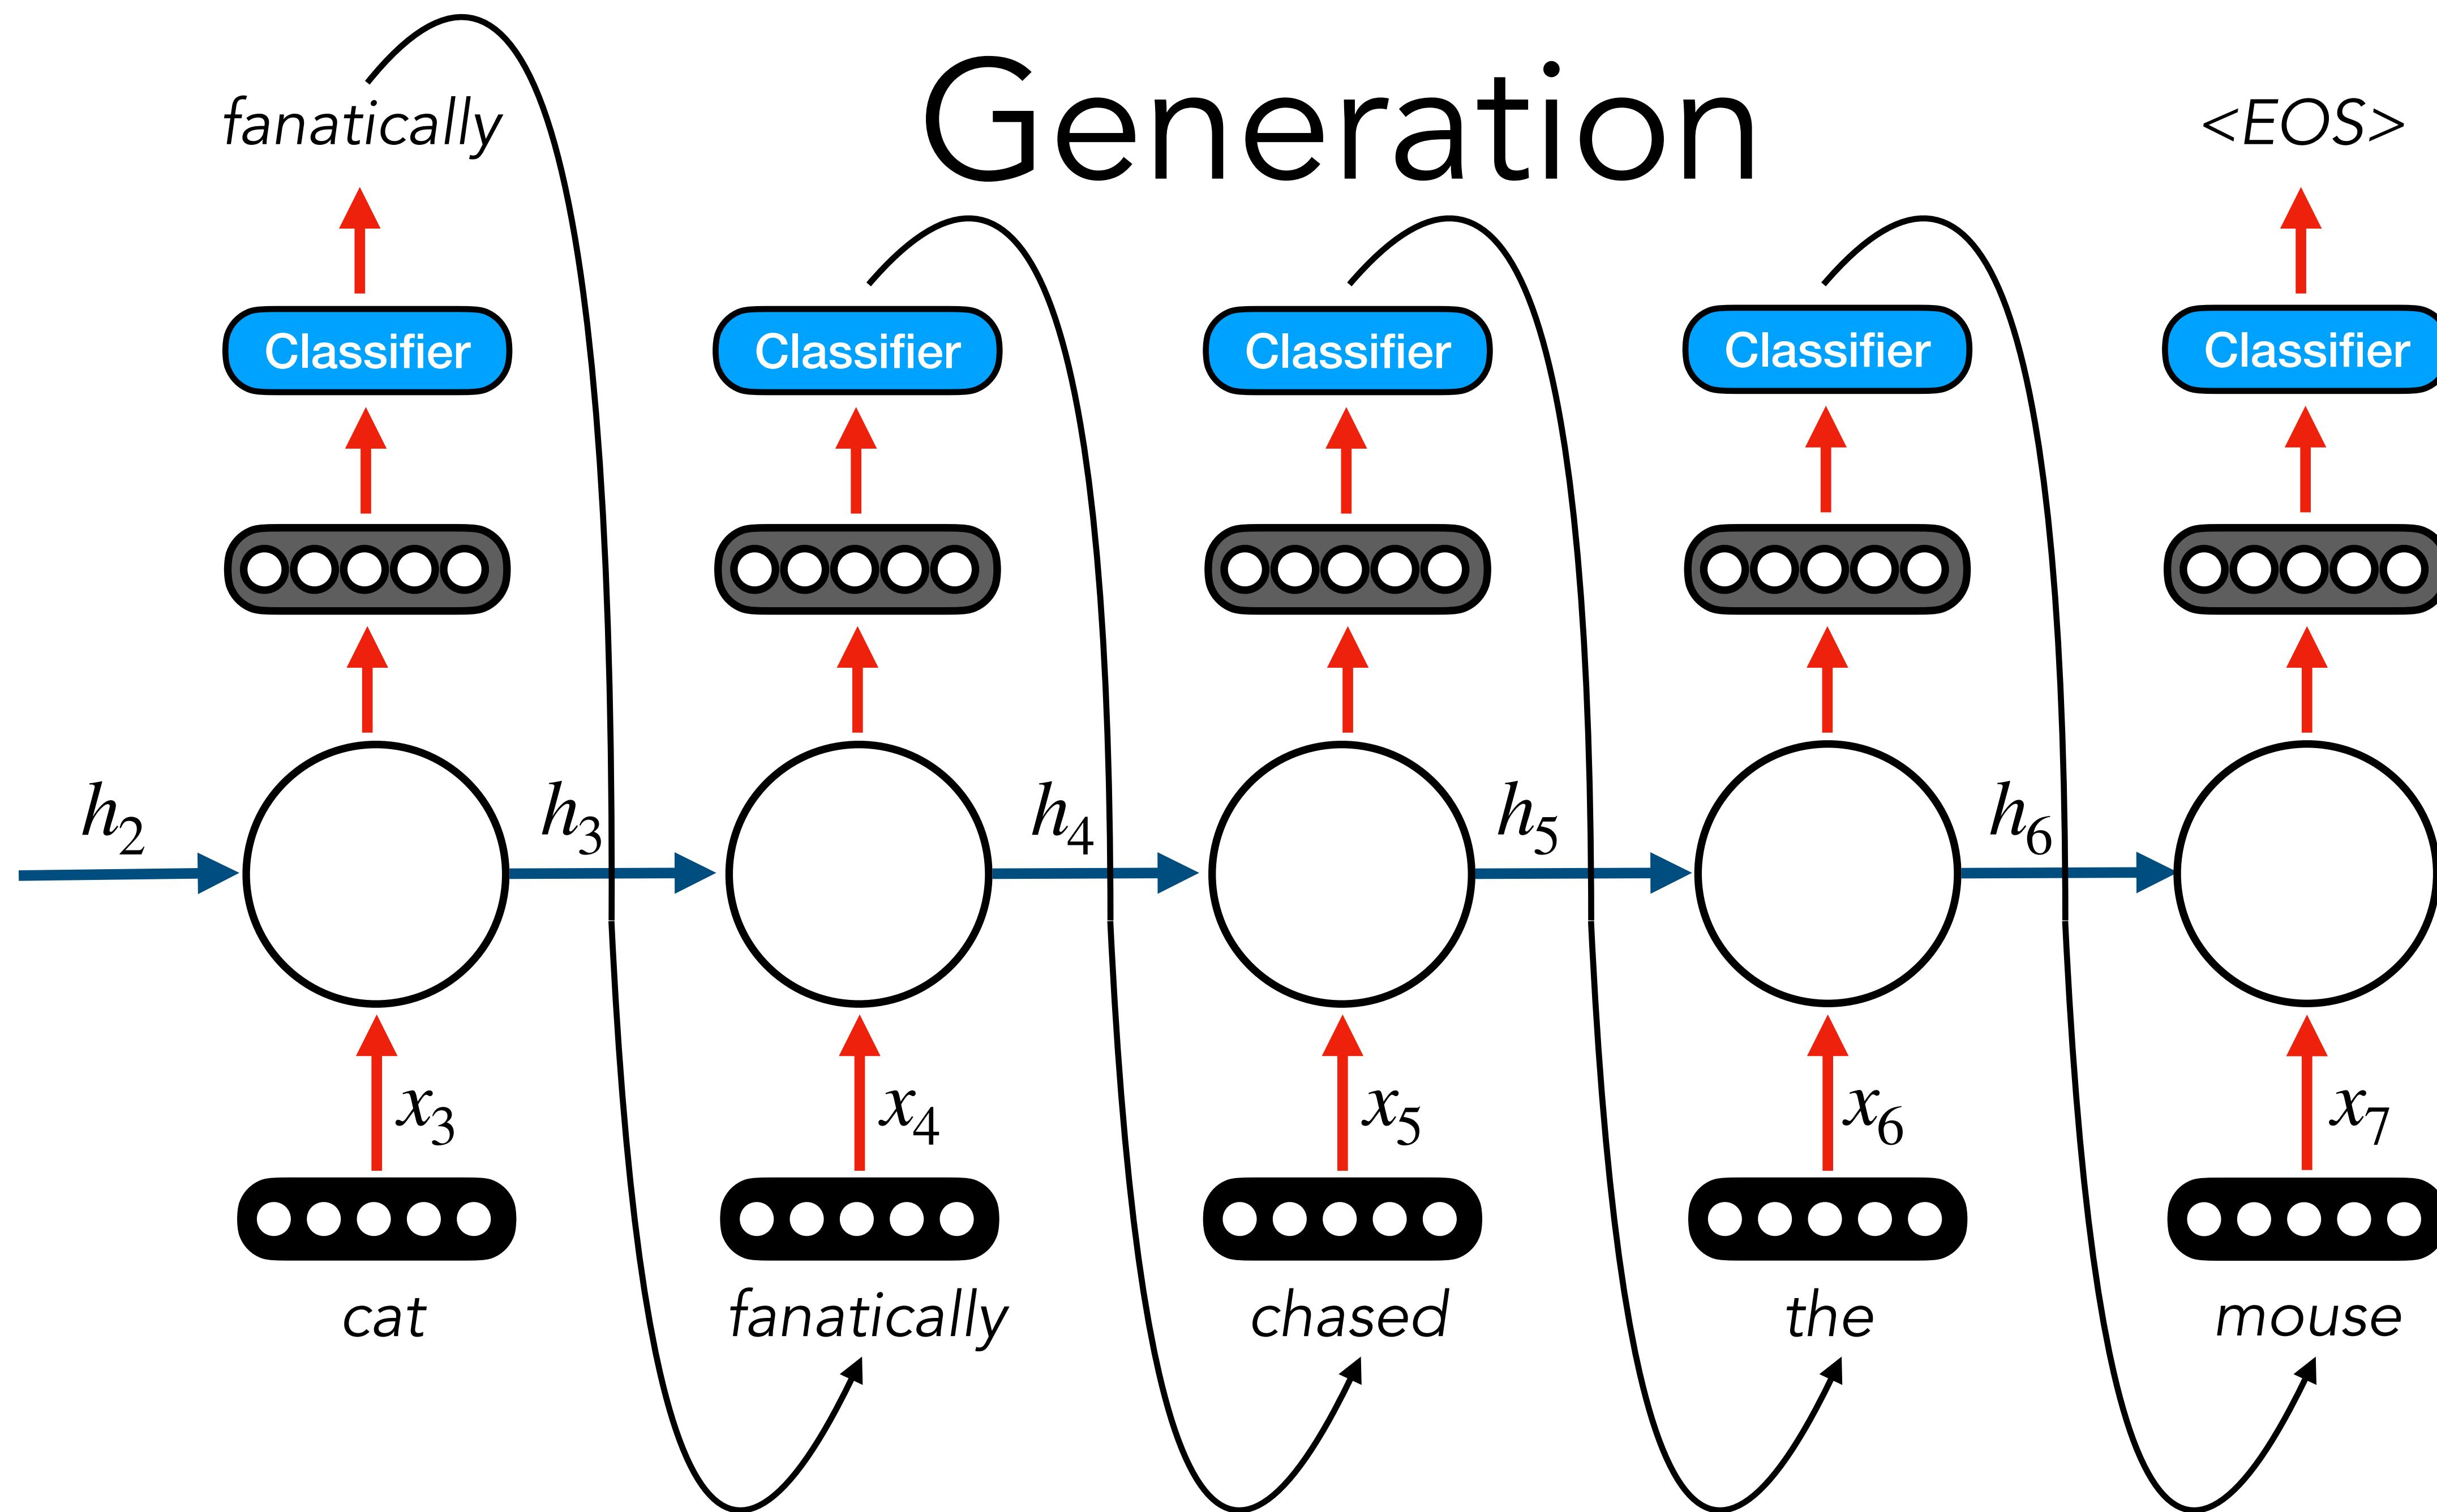
Why would you use the output of the last recurrent unit as the one to predict a label?

It's the only one that has seen the whole sequence!

Sequence Labeling



Generation



Example

PANDARUS:

Alas, I think he shall be come approached and the day
When little strain would be attain'd into being never fed,
And who is but a chain and subjects of his death,
I should not sleep.

Second Senator:

They are away this miseries, produced upon my soul,
Breaking and strongly should be buried, when I perish
The earth and thoughts of many states.

DUKE VINCENTIO:

Well, your wit is in the care of side and that.

Second Lord:

They would be ruled after this chamber, and
my fair nues begun out of the fact, to be conveyed,
Whose noble souls I'll have the heart of the wars.

Clown:

Come, sir, I will make did behold your worship.

VIOLA:

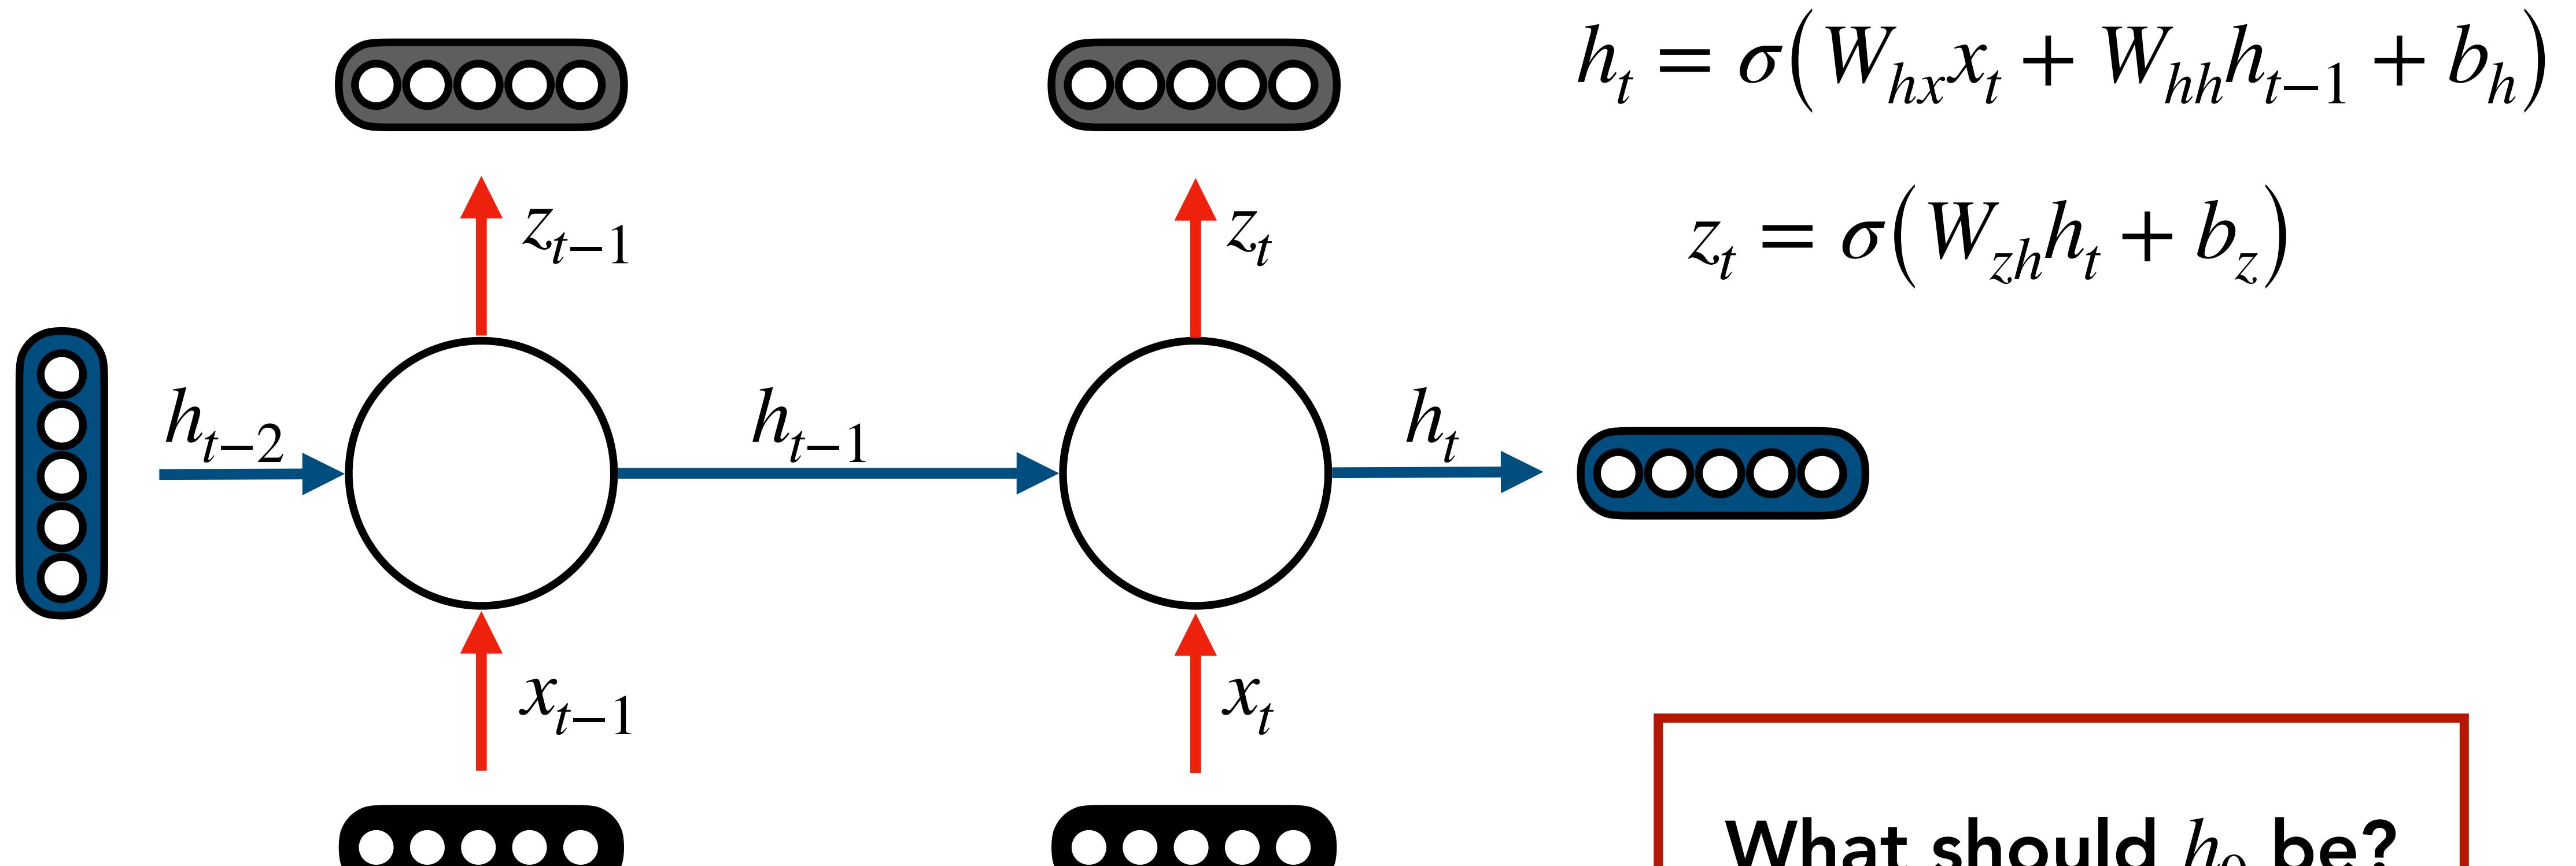
I'll drink it.

**Not bad for generating
Shakespeare!**

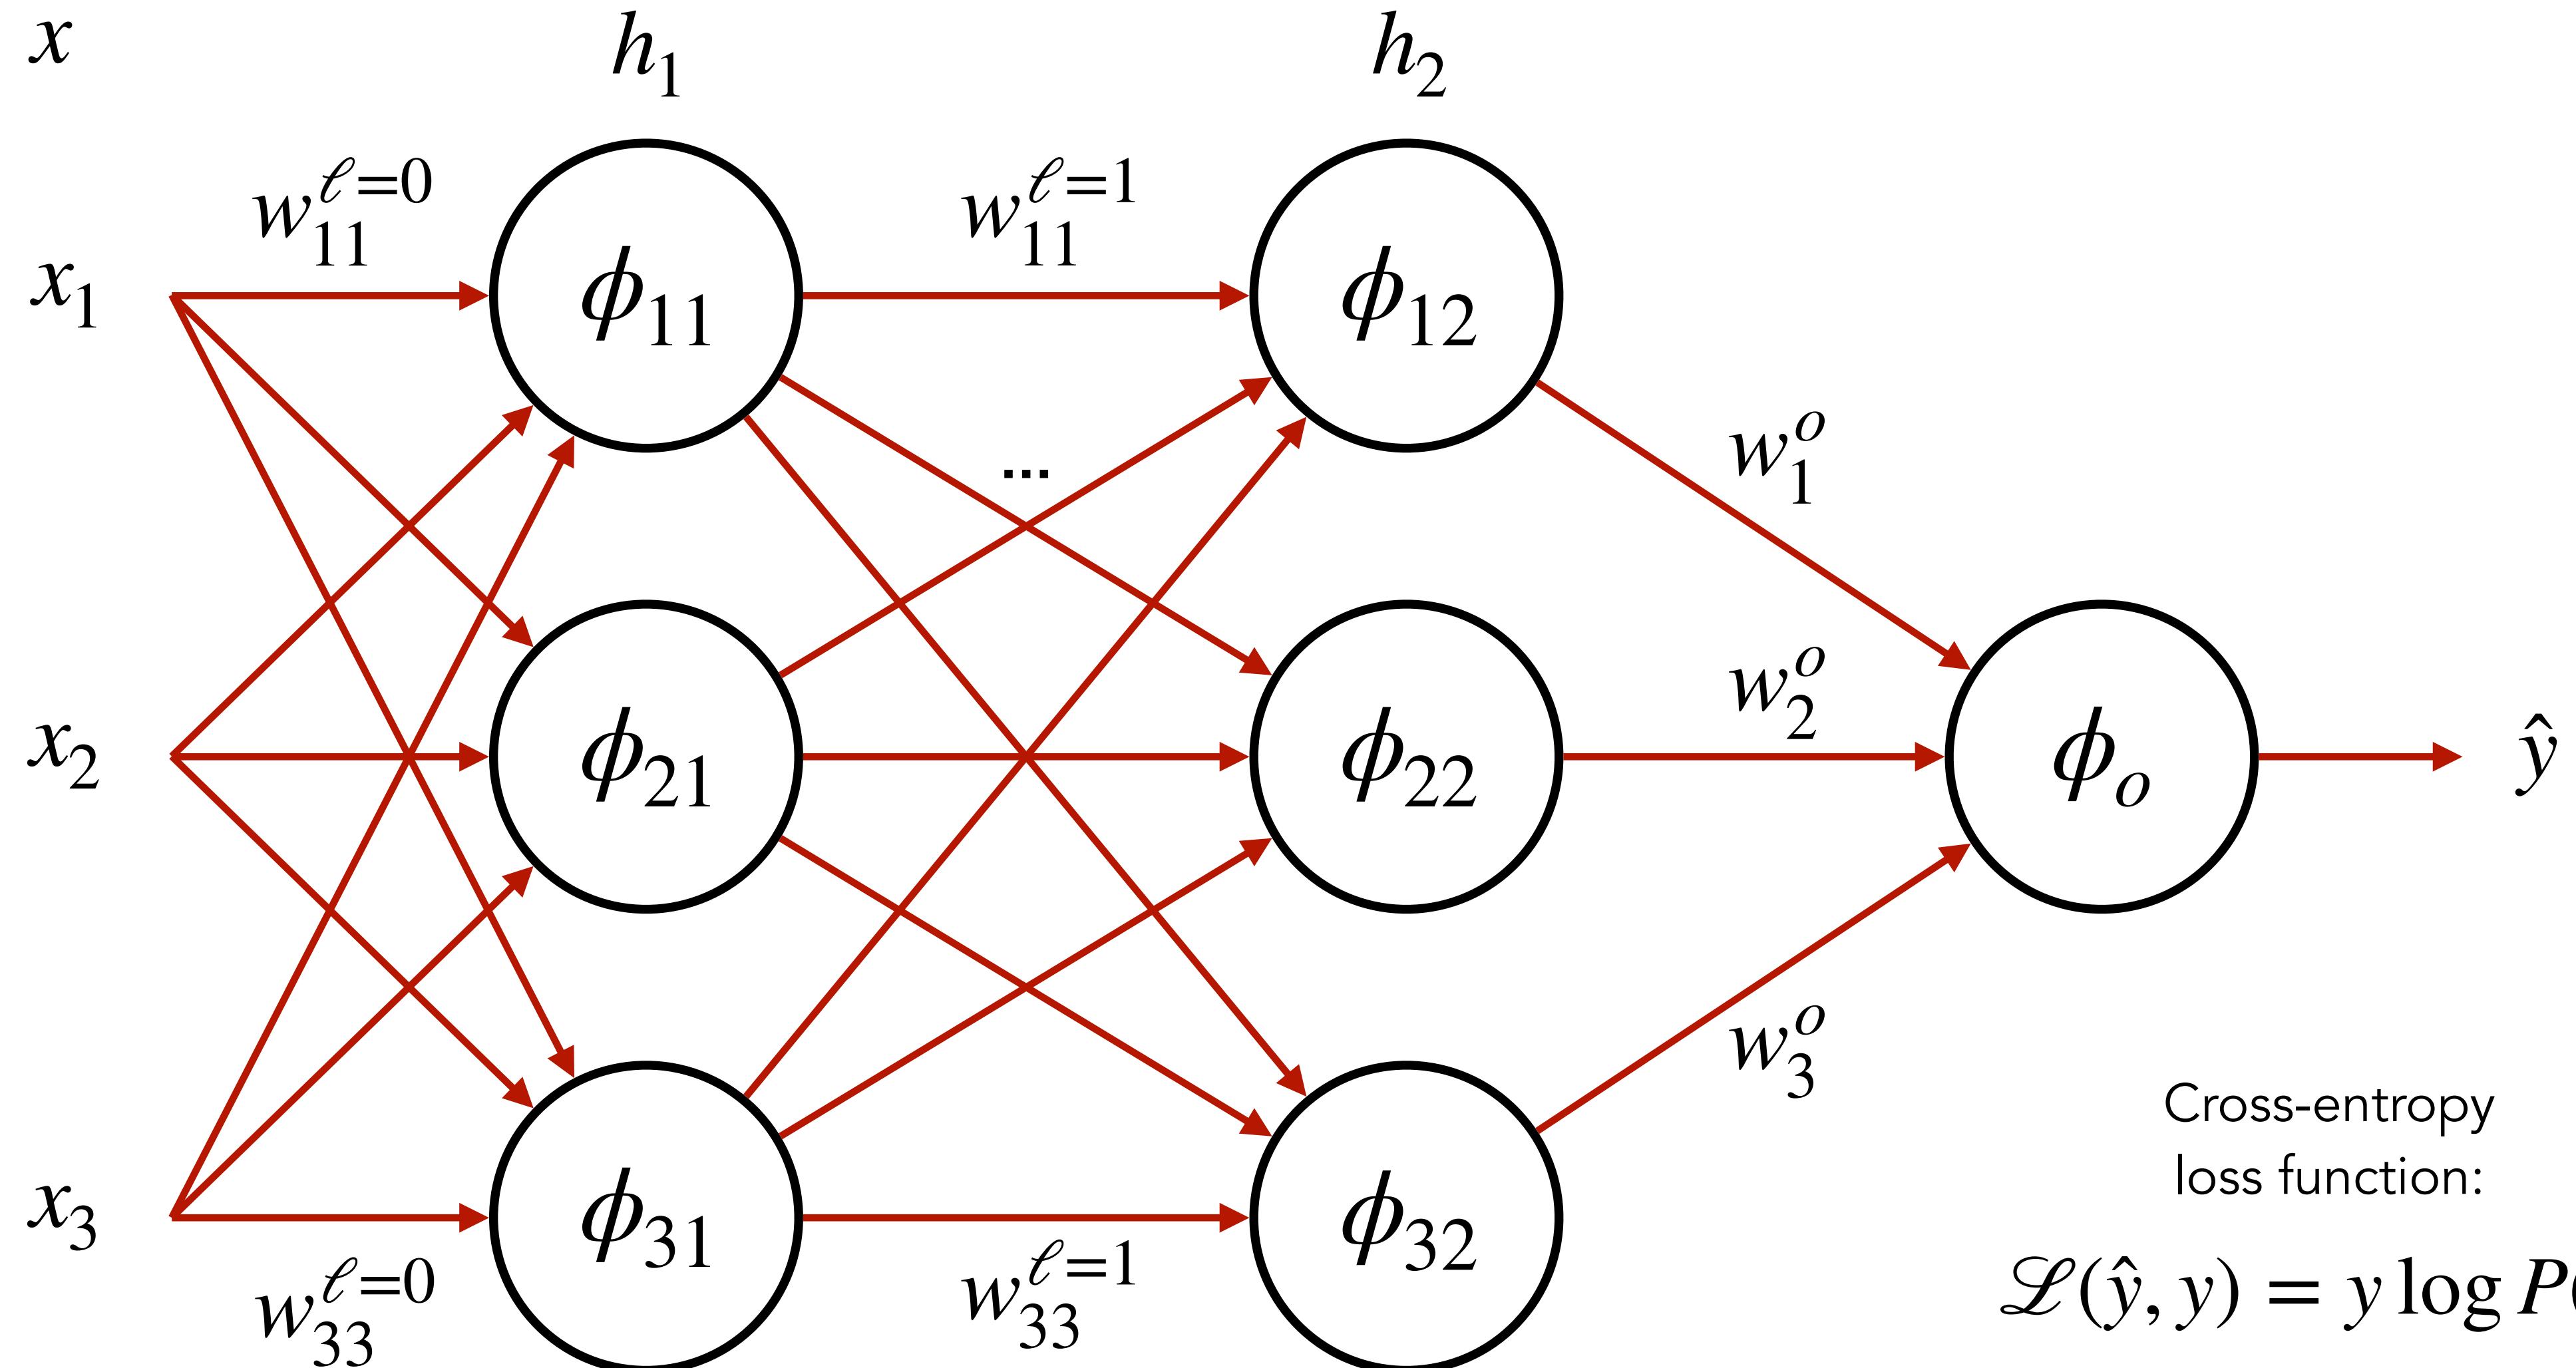
Wikipedia works too!

Naturalism and decision for the majority of Arab countries' capitalide was grounded by the Irish language by [[John Clair]], [[An Imperial Japanese Revolt]], associated with Guangzham's sovereignty. His generals were the powerful ruler of the Portugal in the [[Protestant Immineners]], which could be said to be directly in Cantonese Communication, which followed a ceremony and set inspired prison, training. The emperor travelled back to [[Antioch, Perth, October 25|21]] to note, the Kingdom of Costa Rica, unsuccessful fashioned the [[Thrales]], [[Cynth's Dajoard]], known in western [[Scotland]], near Italy to the conquest of India with the conflict. Copyright was the succession of independence in the slop of Syrian influence that was a famous German movement based on a more popular servicious, non-doctrinal and sexual power post. Many governments recognize the military housing of the [[Civil Liberalization and Infantry Resolution 265 National Party in Hungary]], that is sympathetic to be to the [[Punjab Resolution]] (PJS) [<http://www.humah.yahoo.com/guardian.cfm/7754800786d17551963s89.htm>] Official economics Adjoint for the Nazism, Montgomery was swear to advance to the resources for those Socialism's rule, was starting to signing a major tripad of aid exile.]

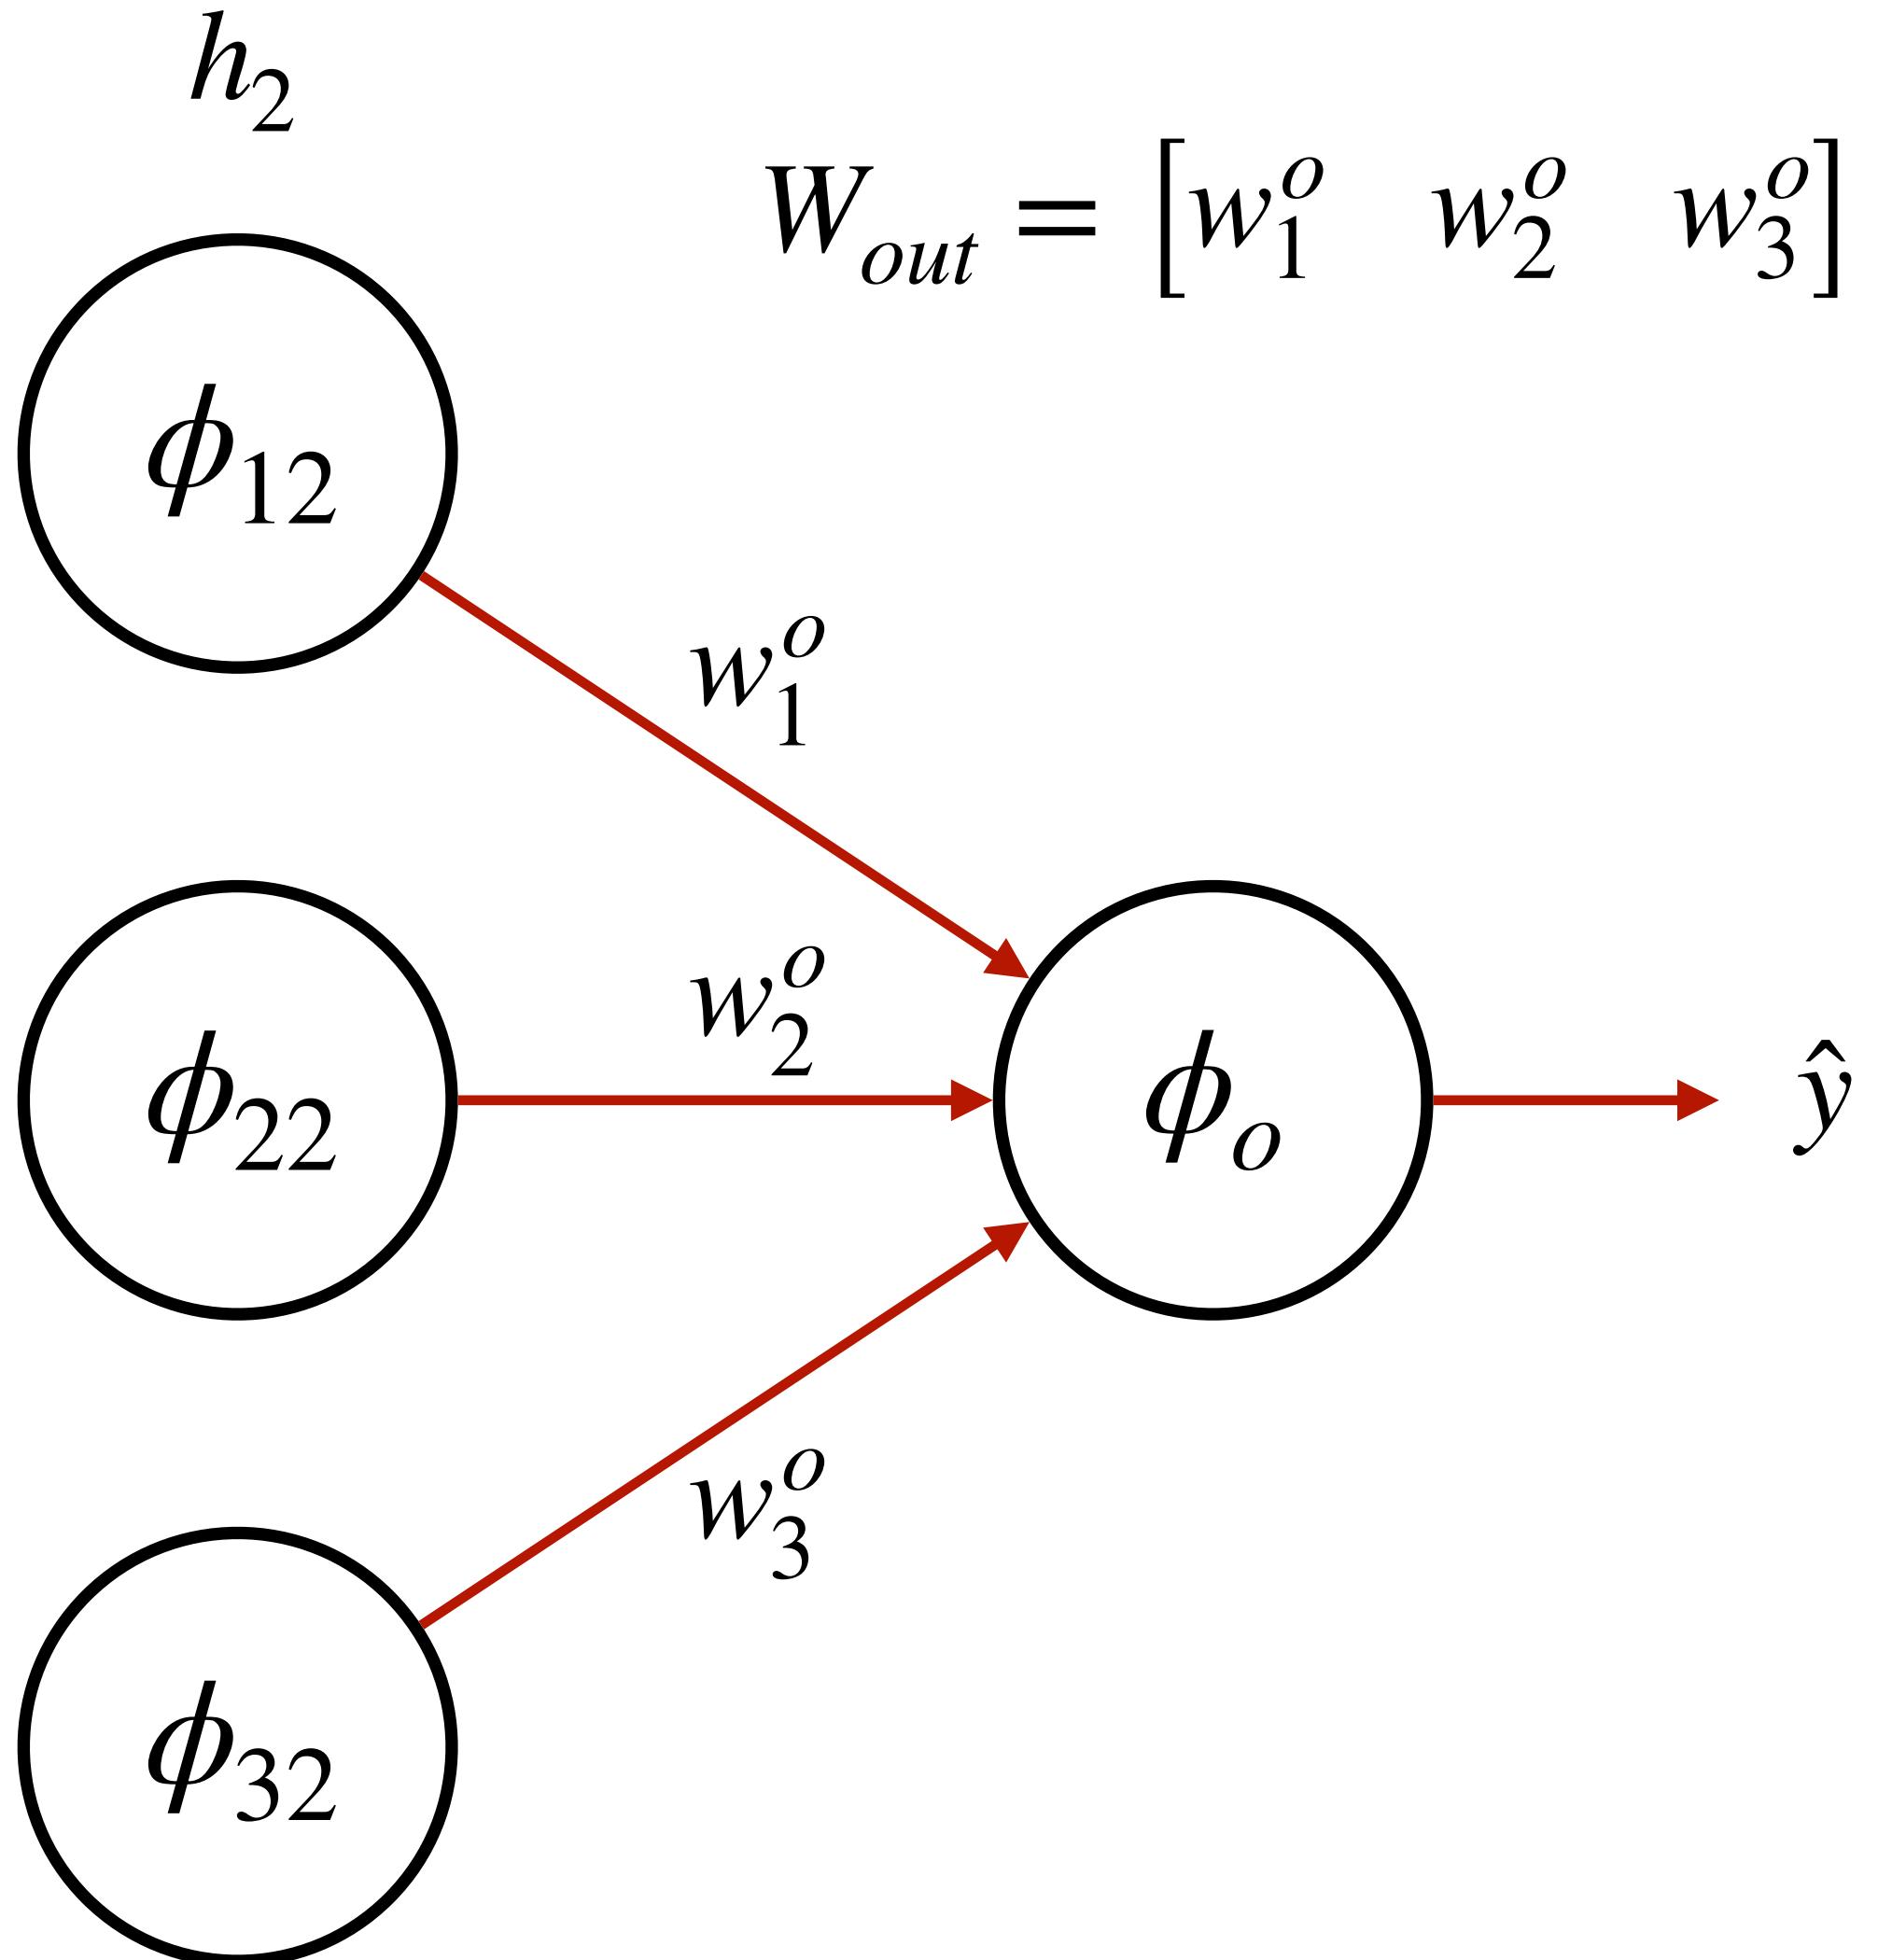
Classical RNN: Elman Network



Backpropagation Review: FFNs



Backpropagation Review: FFNs



$$W_{out} = [w_1^o \quad w_2^o \quad w_3^o]$$

$$\mathcal{L}(\hat{y}, y) = -y \log P(\hat{y})$$

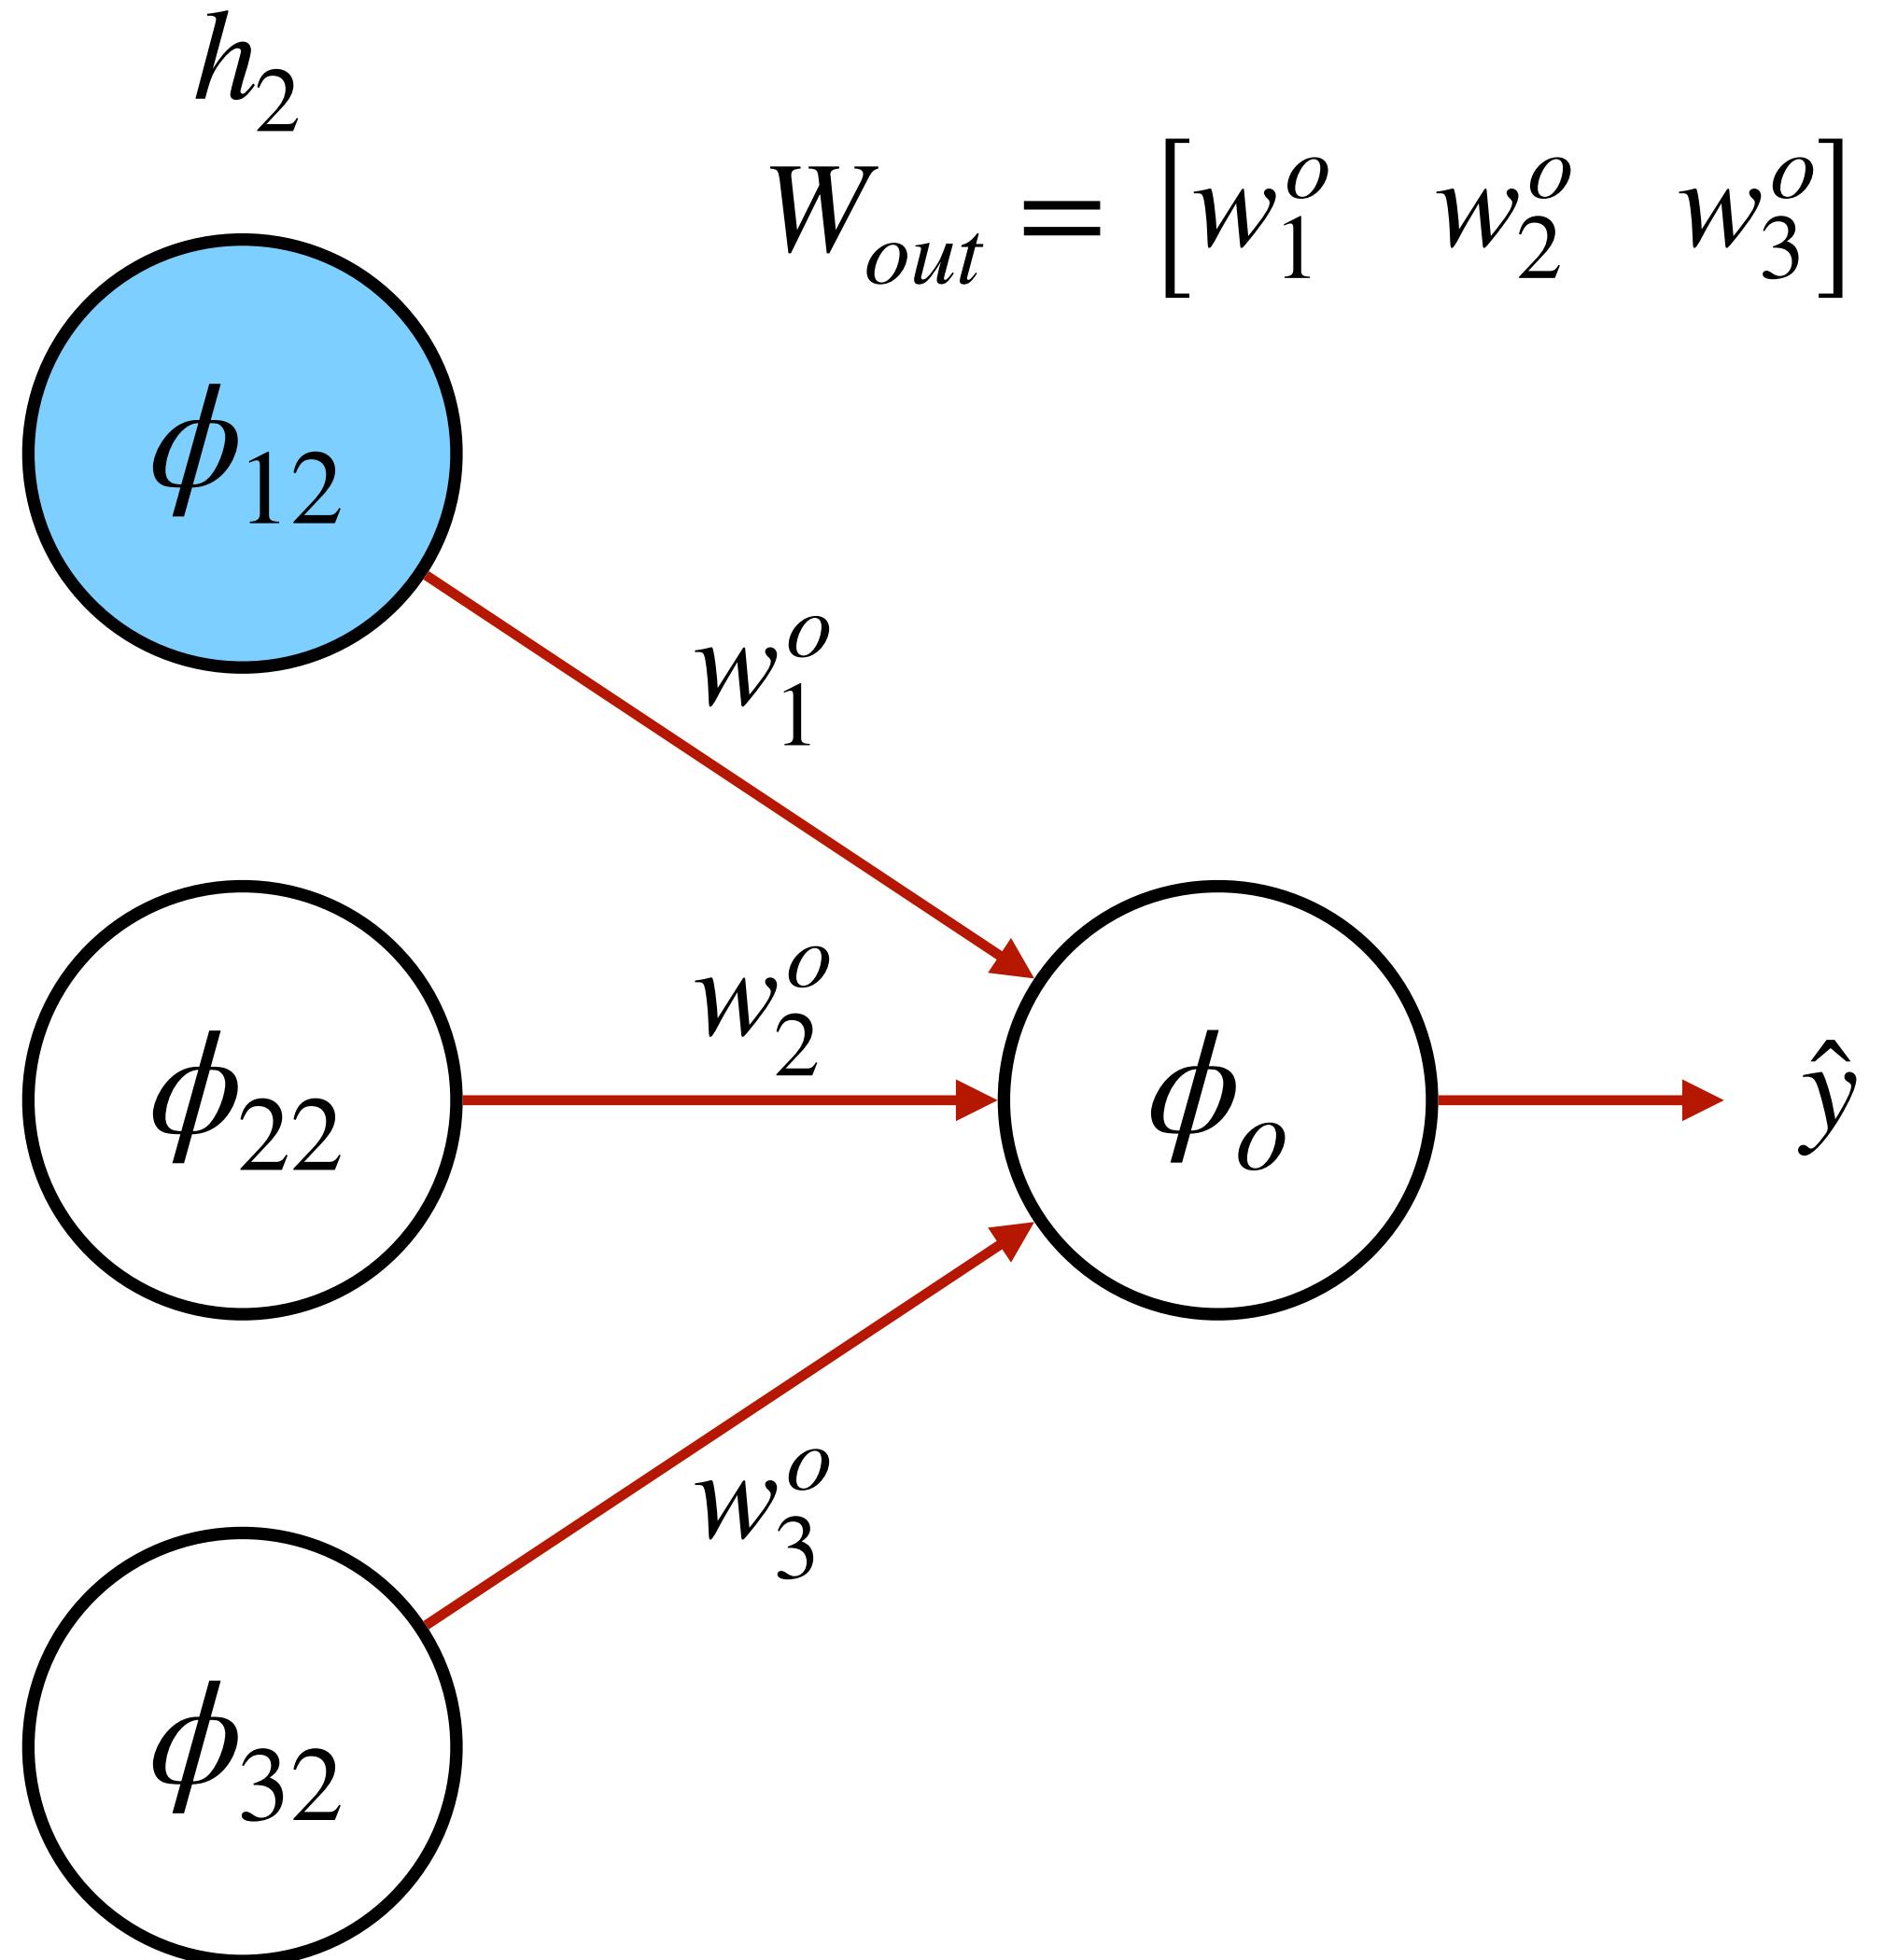
$$P(\hat{y}) = \phi_o(u)$$

$$u = W_{out}h_2 = w_1^o \times \phi_{12}(.) + w_2^o \times \phi_{22}(.) + w_3^o \times \phi_{32}(.)$$

$$\frac{\partial \mathcal{L}(\hat{y}, y)}{\partial \phi_{12}(.)} = \frac{\partial \mathcal{L}(\hat{y}, y)}{\partial \hat{y}} \frac{\partial \hat{y}}{\partial u} \frac{\partial u}{\partial \phi_{12}(.)}$$

$$= \frac{\partial \mathcal{L}(\hat{y}, y)}{\partial \hat{y}} \frac{\partial \phi_o(u)}{\partial u} w_1^o$$

Backpropagation Review: FFNs



$$\mathcal{L}(\hat{y}, y) = -y \log P(\hat{y})$$

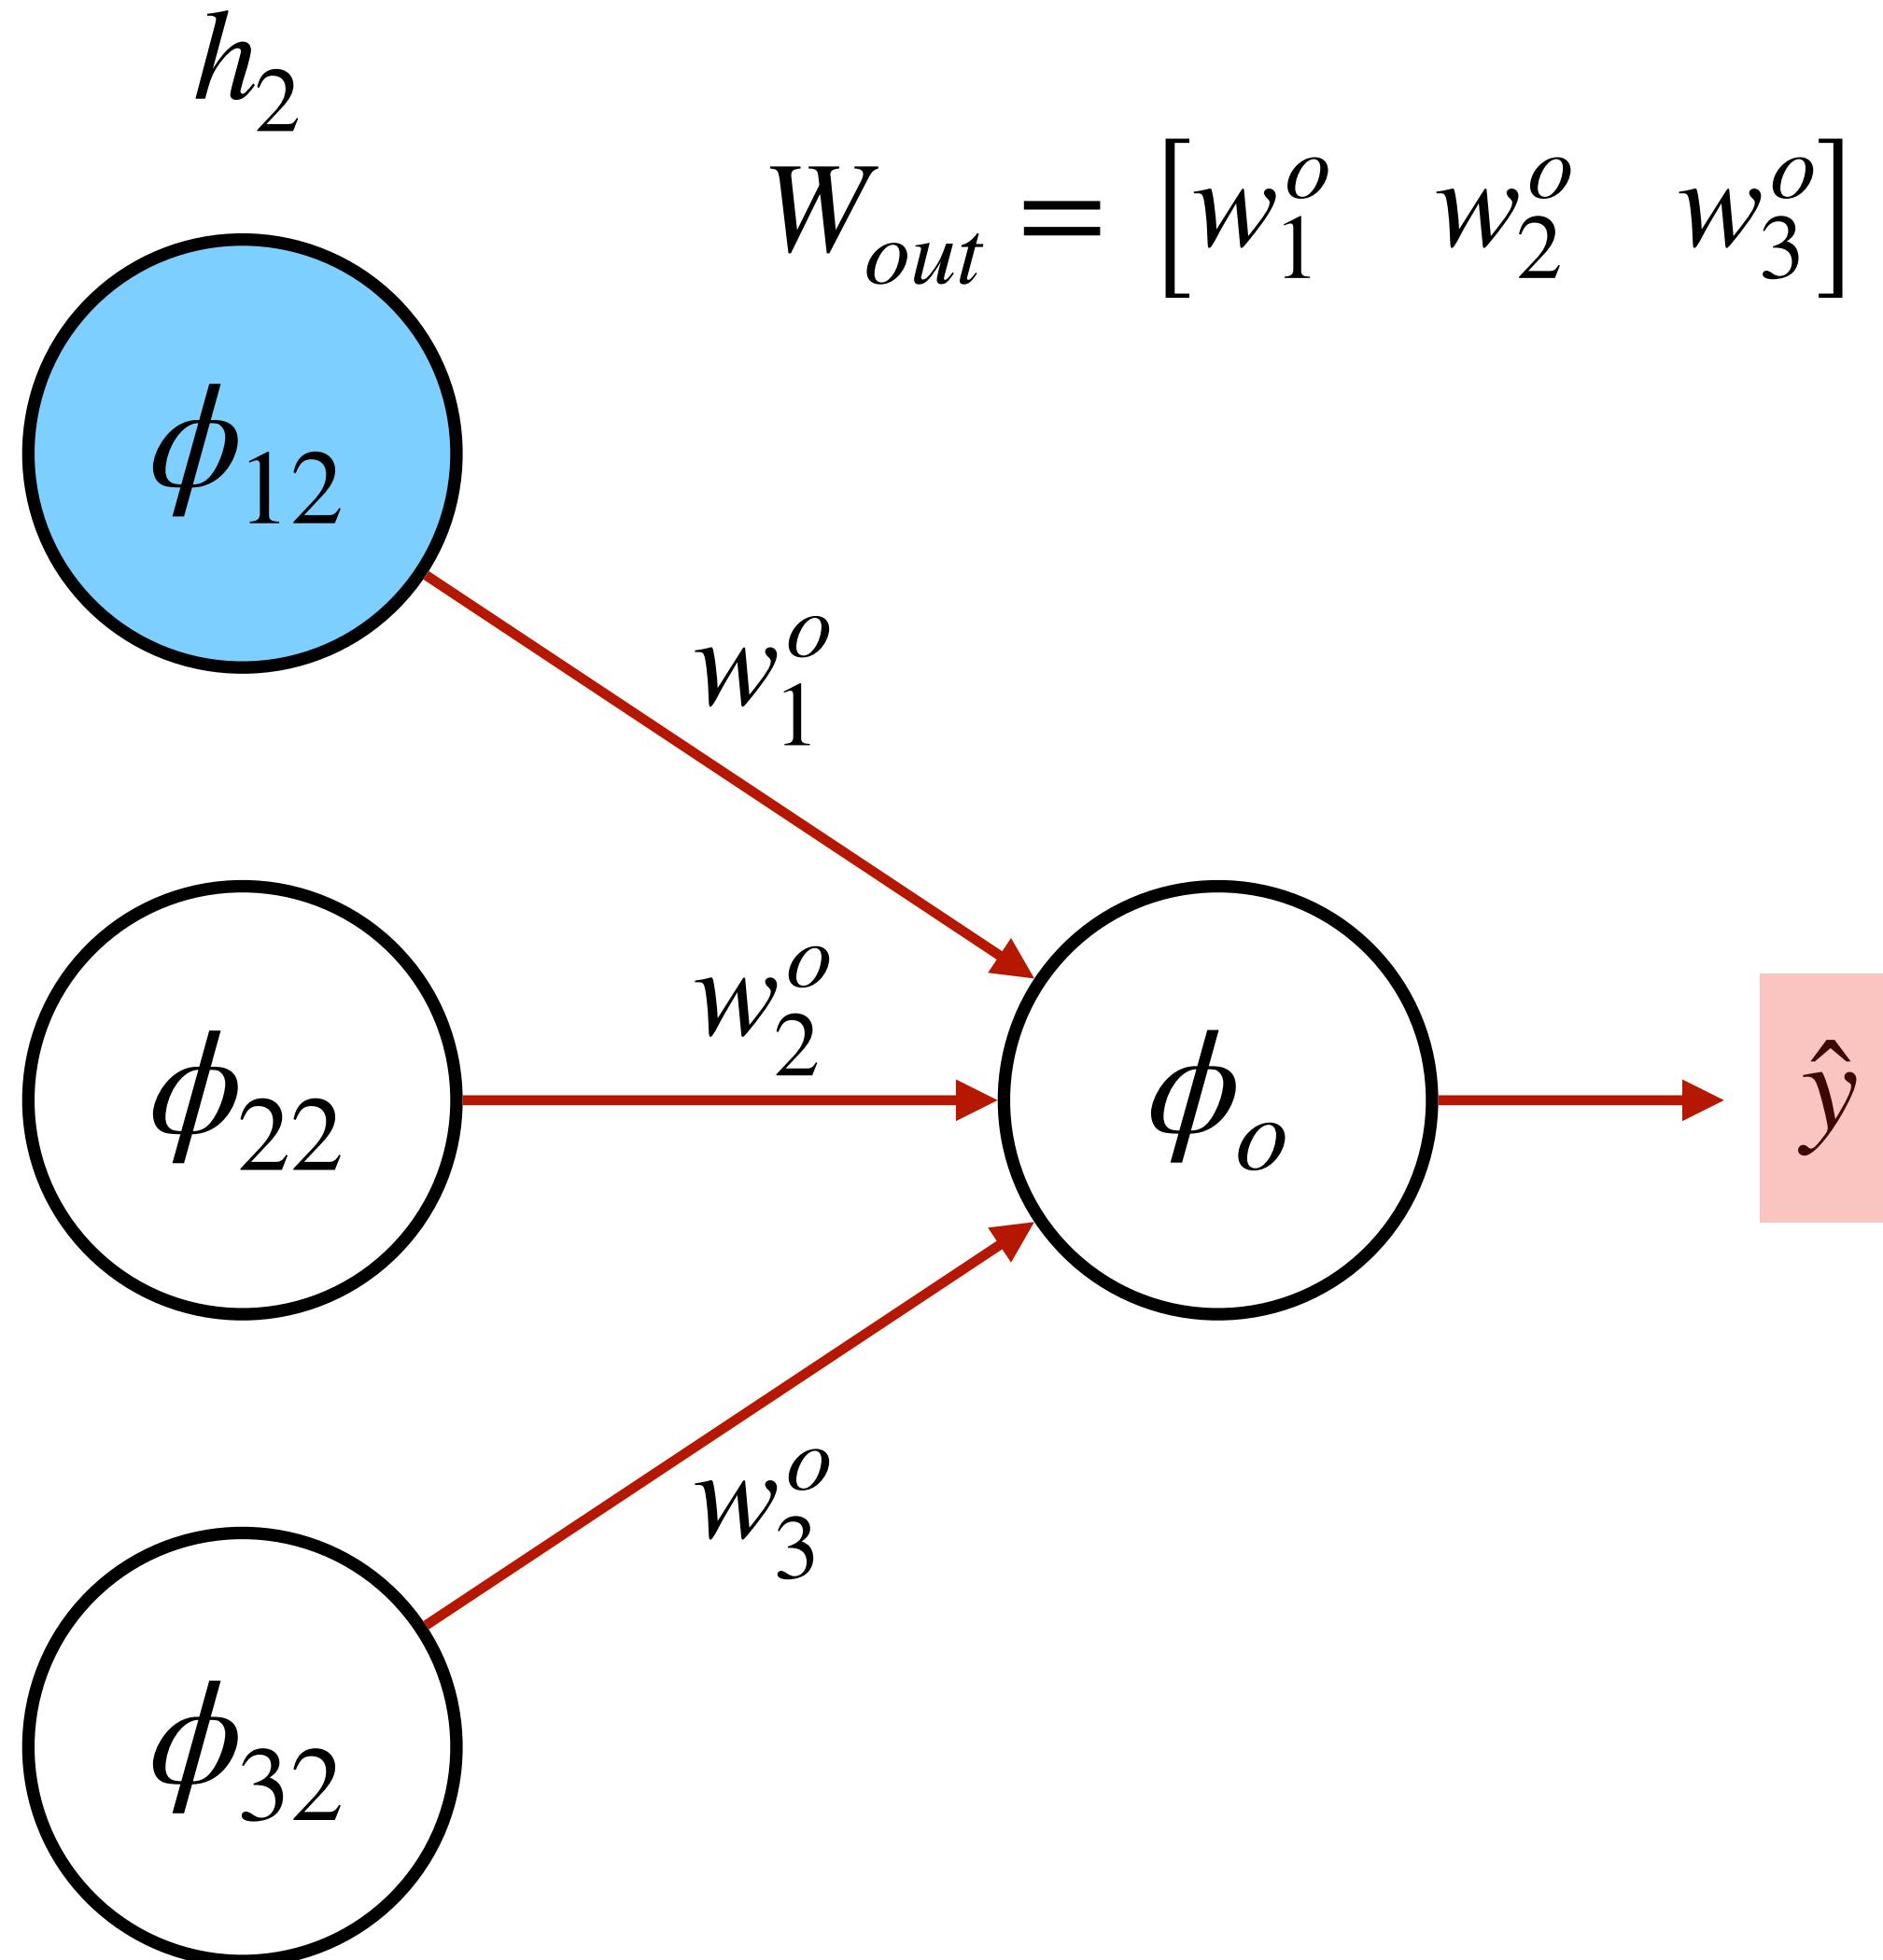
$$P(\hat{y}) = \phi_o(u)$$

$$u = W_{out}h_2 = w_1^o \times \phi_{12}(.) + w_2^o \times \phi_{22}(.) + w_3^o \times \phi_{32}(.)$$

$$\frac{\partial \mathcal{L}(\hat{y}, y)}{\partial \phi_{12}(.)} = \frac{\partial \mathcal{L}(\hat{y}, y)}{\partial \hat{y}} \frac{\partial \hat{y}}{\partial u} \frac{\partial u}{\partial \phi_{12}(.)}$$

$$= \frac{\partial \mathcal{L}(\hat{y}, y)}{\partial \hat{y}} \frac{\partial \phi_o(u)}{\partial u} w_1^o$$

Backpropagation Review: FFNs



$$W_{out} = [w_1^o \quad w_2^o \quad w_3^o]$$

$$\mathcal{L}(\hat{y}, y) = -y \log P(\hat{y})$$

$$P(\hat{y}) = \phi_o(u)$$

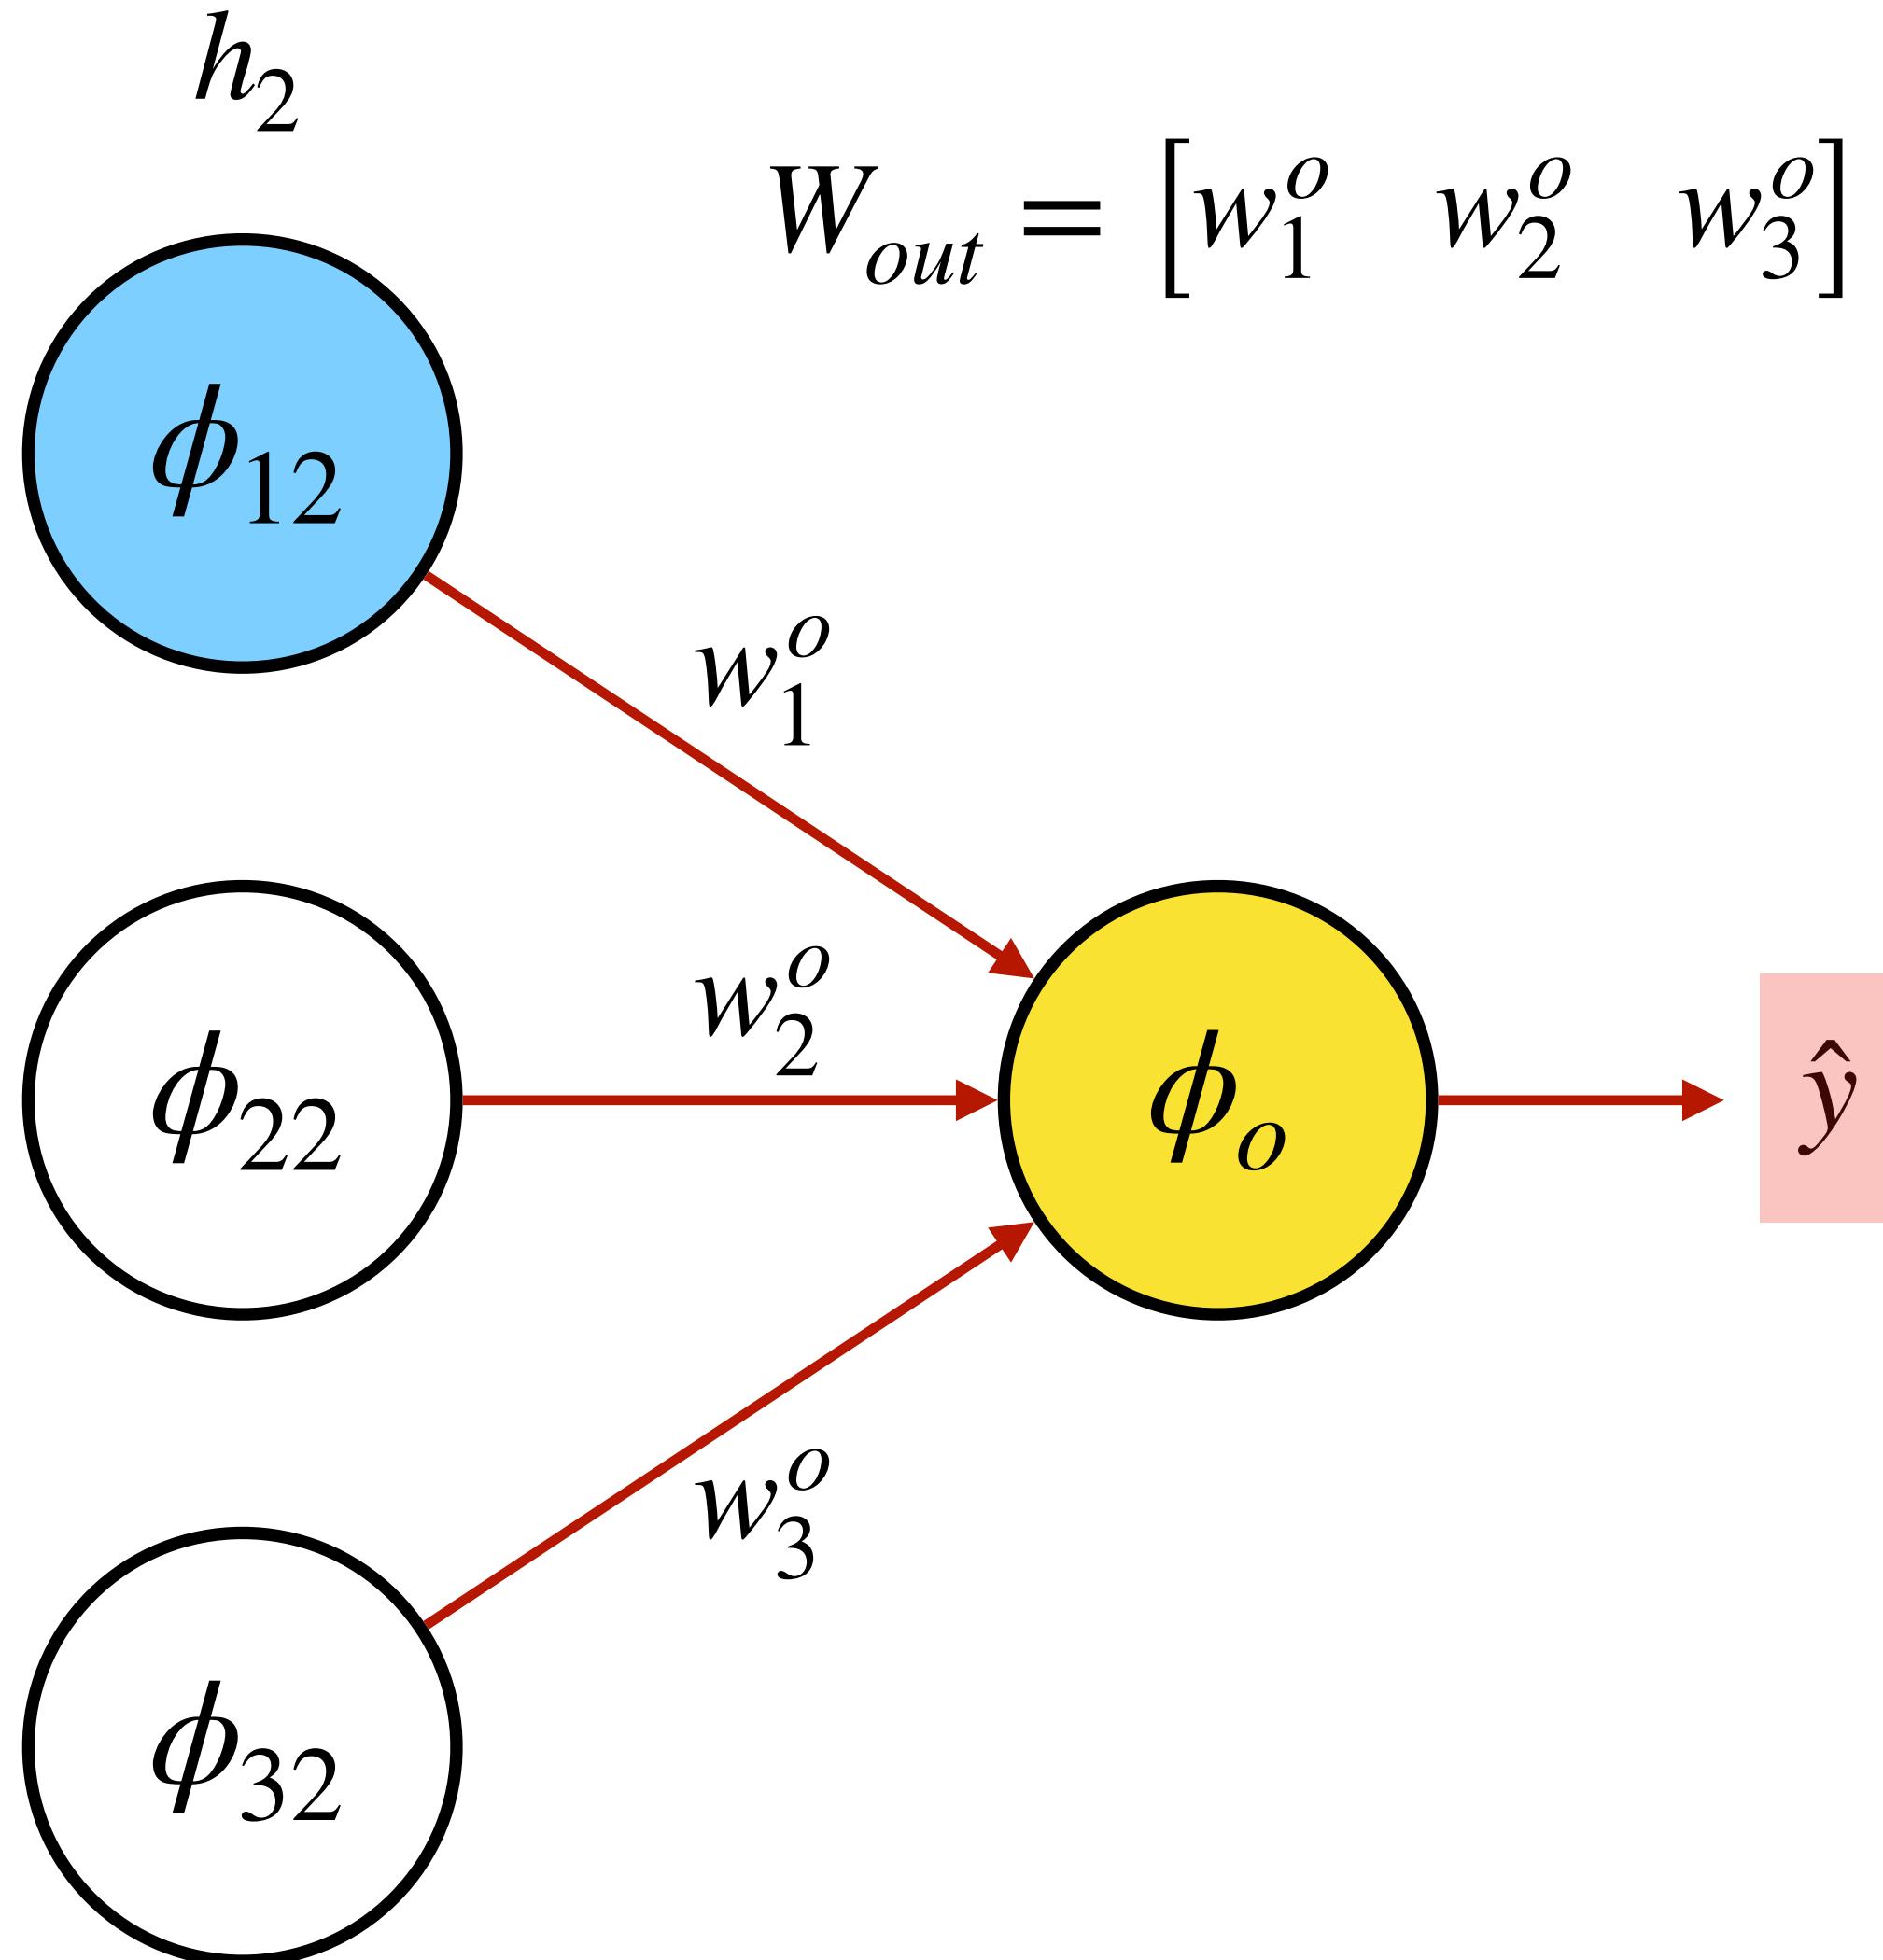
$$u = W_{out}h_2 = w_1^o \times \phi_{12}(.) + w_2^o \times \phi_{22}(.) + w_3^o \times \phi_{32}(.)$$

$$\frac{\partial \mathcal{L}(\hat{y}, y)}{\partial \phi_{12}(.)} = \frac{\partial \mathcal{L}(\hat{y}, y)}{\partial \hat{y}} \frac{\partial \hat{y}}{\partial u} \frac{\partial u}{\partial \phi_{12}(.)}$$

$$= \frac{\partial \mathcal{L}(\hat{y}, y)}{\partial \hat{y}} \frac{\partial \phi_o(u)}{\partial u} w_1^o$$

Depends on label y

Backpropagation Review: FFNs



$$W_{out} = [w_1^o \quad w_2^o \quad w_3^o]$$

$$\mathcal{L}(\hat{y}, y) = -y \log P(\hat{y})$$

$$P(\hat{y}) = \phi_o(u)$$

$$u = W_{out}h_2 = w_1^o \times \phi_{12}(.) + w_2^o \times \phi_{22}(.) + w_3^o \times \phi_{32}(.)$$

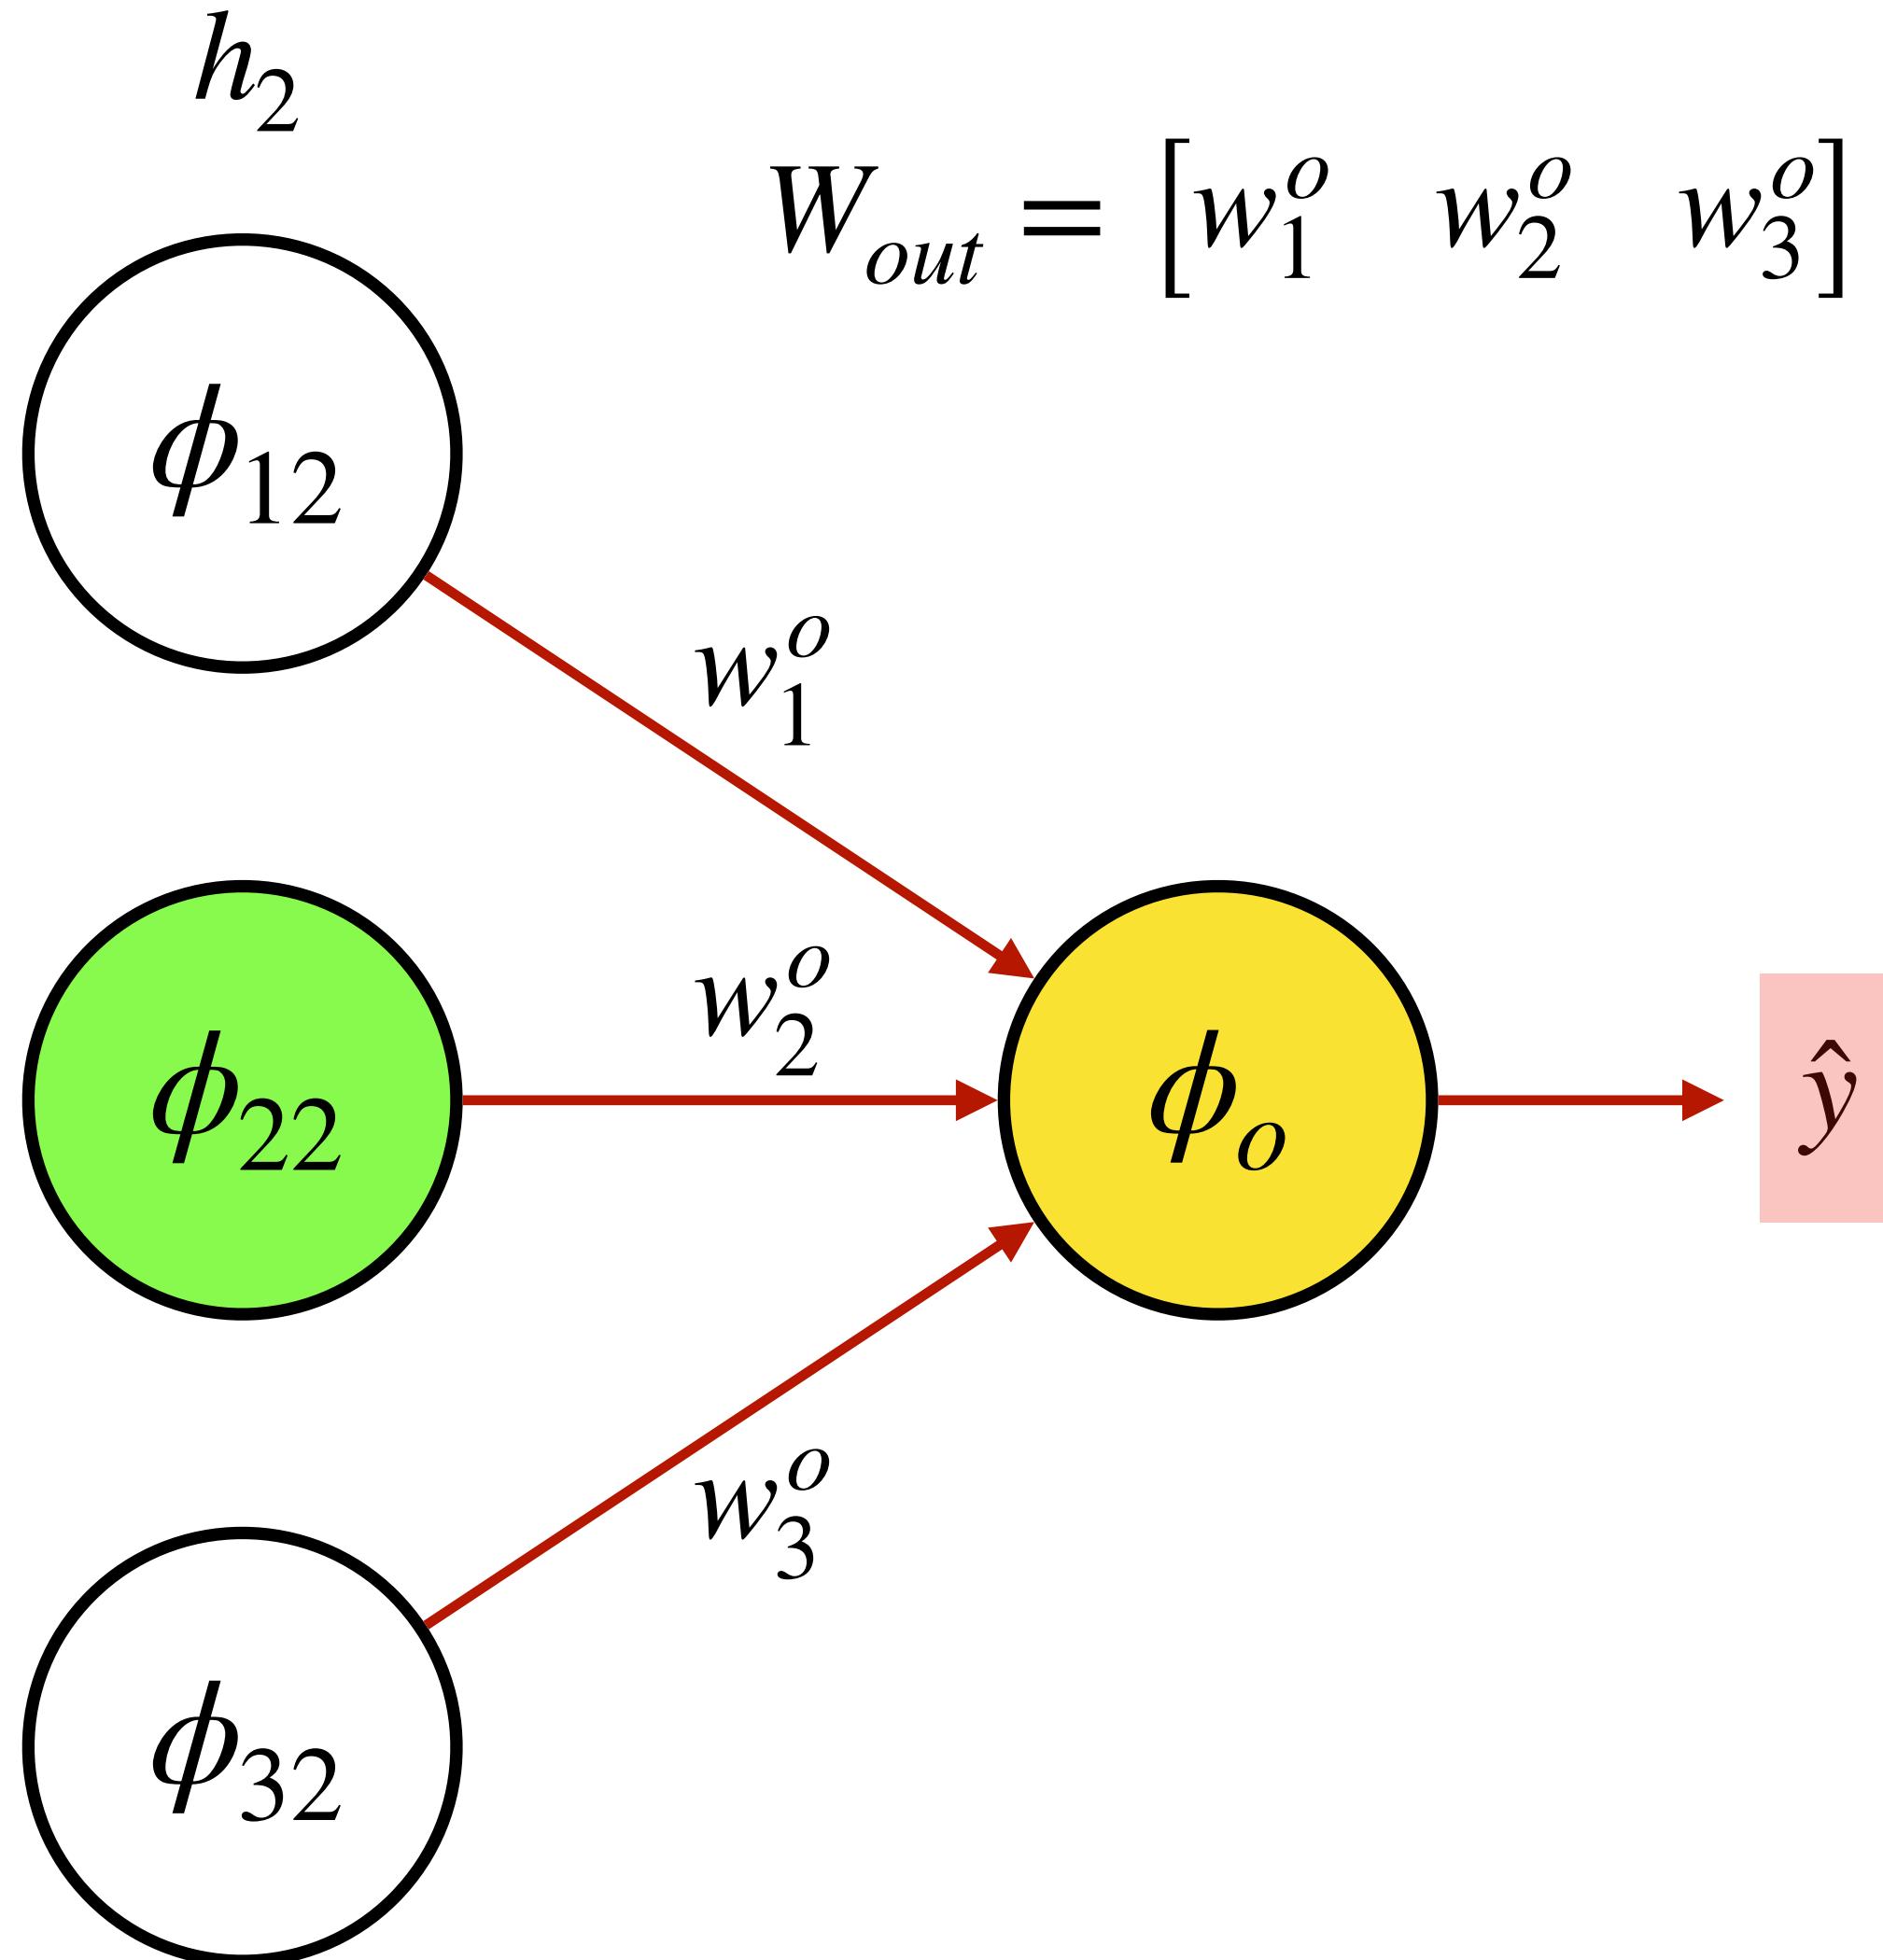
$$\frac{\partial \mathcal{L}(\hat{y}, y)}{\partial \phi_{12}(.)} = \frac{\partial \mathcal{L}(\hat{y}, y)}{\partial \hat{y}} \frac{\partial \hat{y}}{\partial u} \frac{\partial u}{\partial \phi_{12}(.)}$$

$$= \frac{\partial \mathcal{L}(\hat{y}, y)}{\partial \hat{y}} \frac{\partial \phi_o(u)}{\partial u} w_1^o$$

Depends on label y

Depends on ϕ_o

Backpropagation Review: FFNs



$$\mathcal{L}(\hat{y}, y) = -y \log P(\hat{y})$$

$$P(\hat{y}) = \phi_o(u)$$

$$u = W_{out}h_2 = w_1^o \times \phi_{12}(.) + w_2^o \times \phi_{22}(.) + w_3^o \times \phi_{32}(.)$$

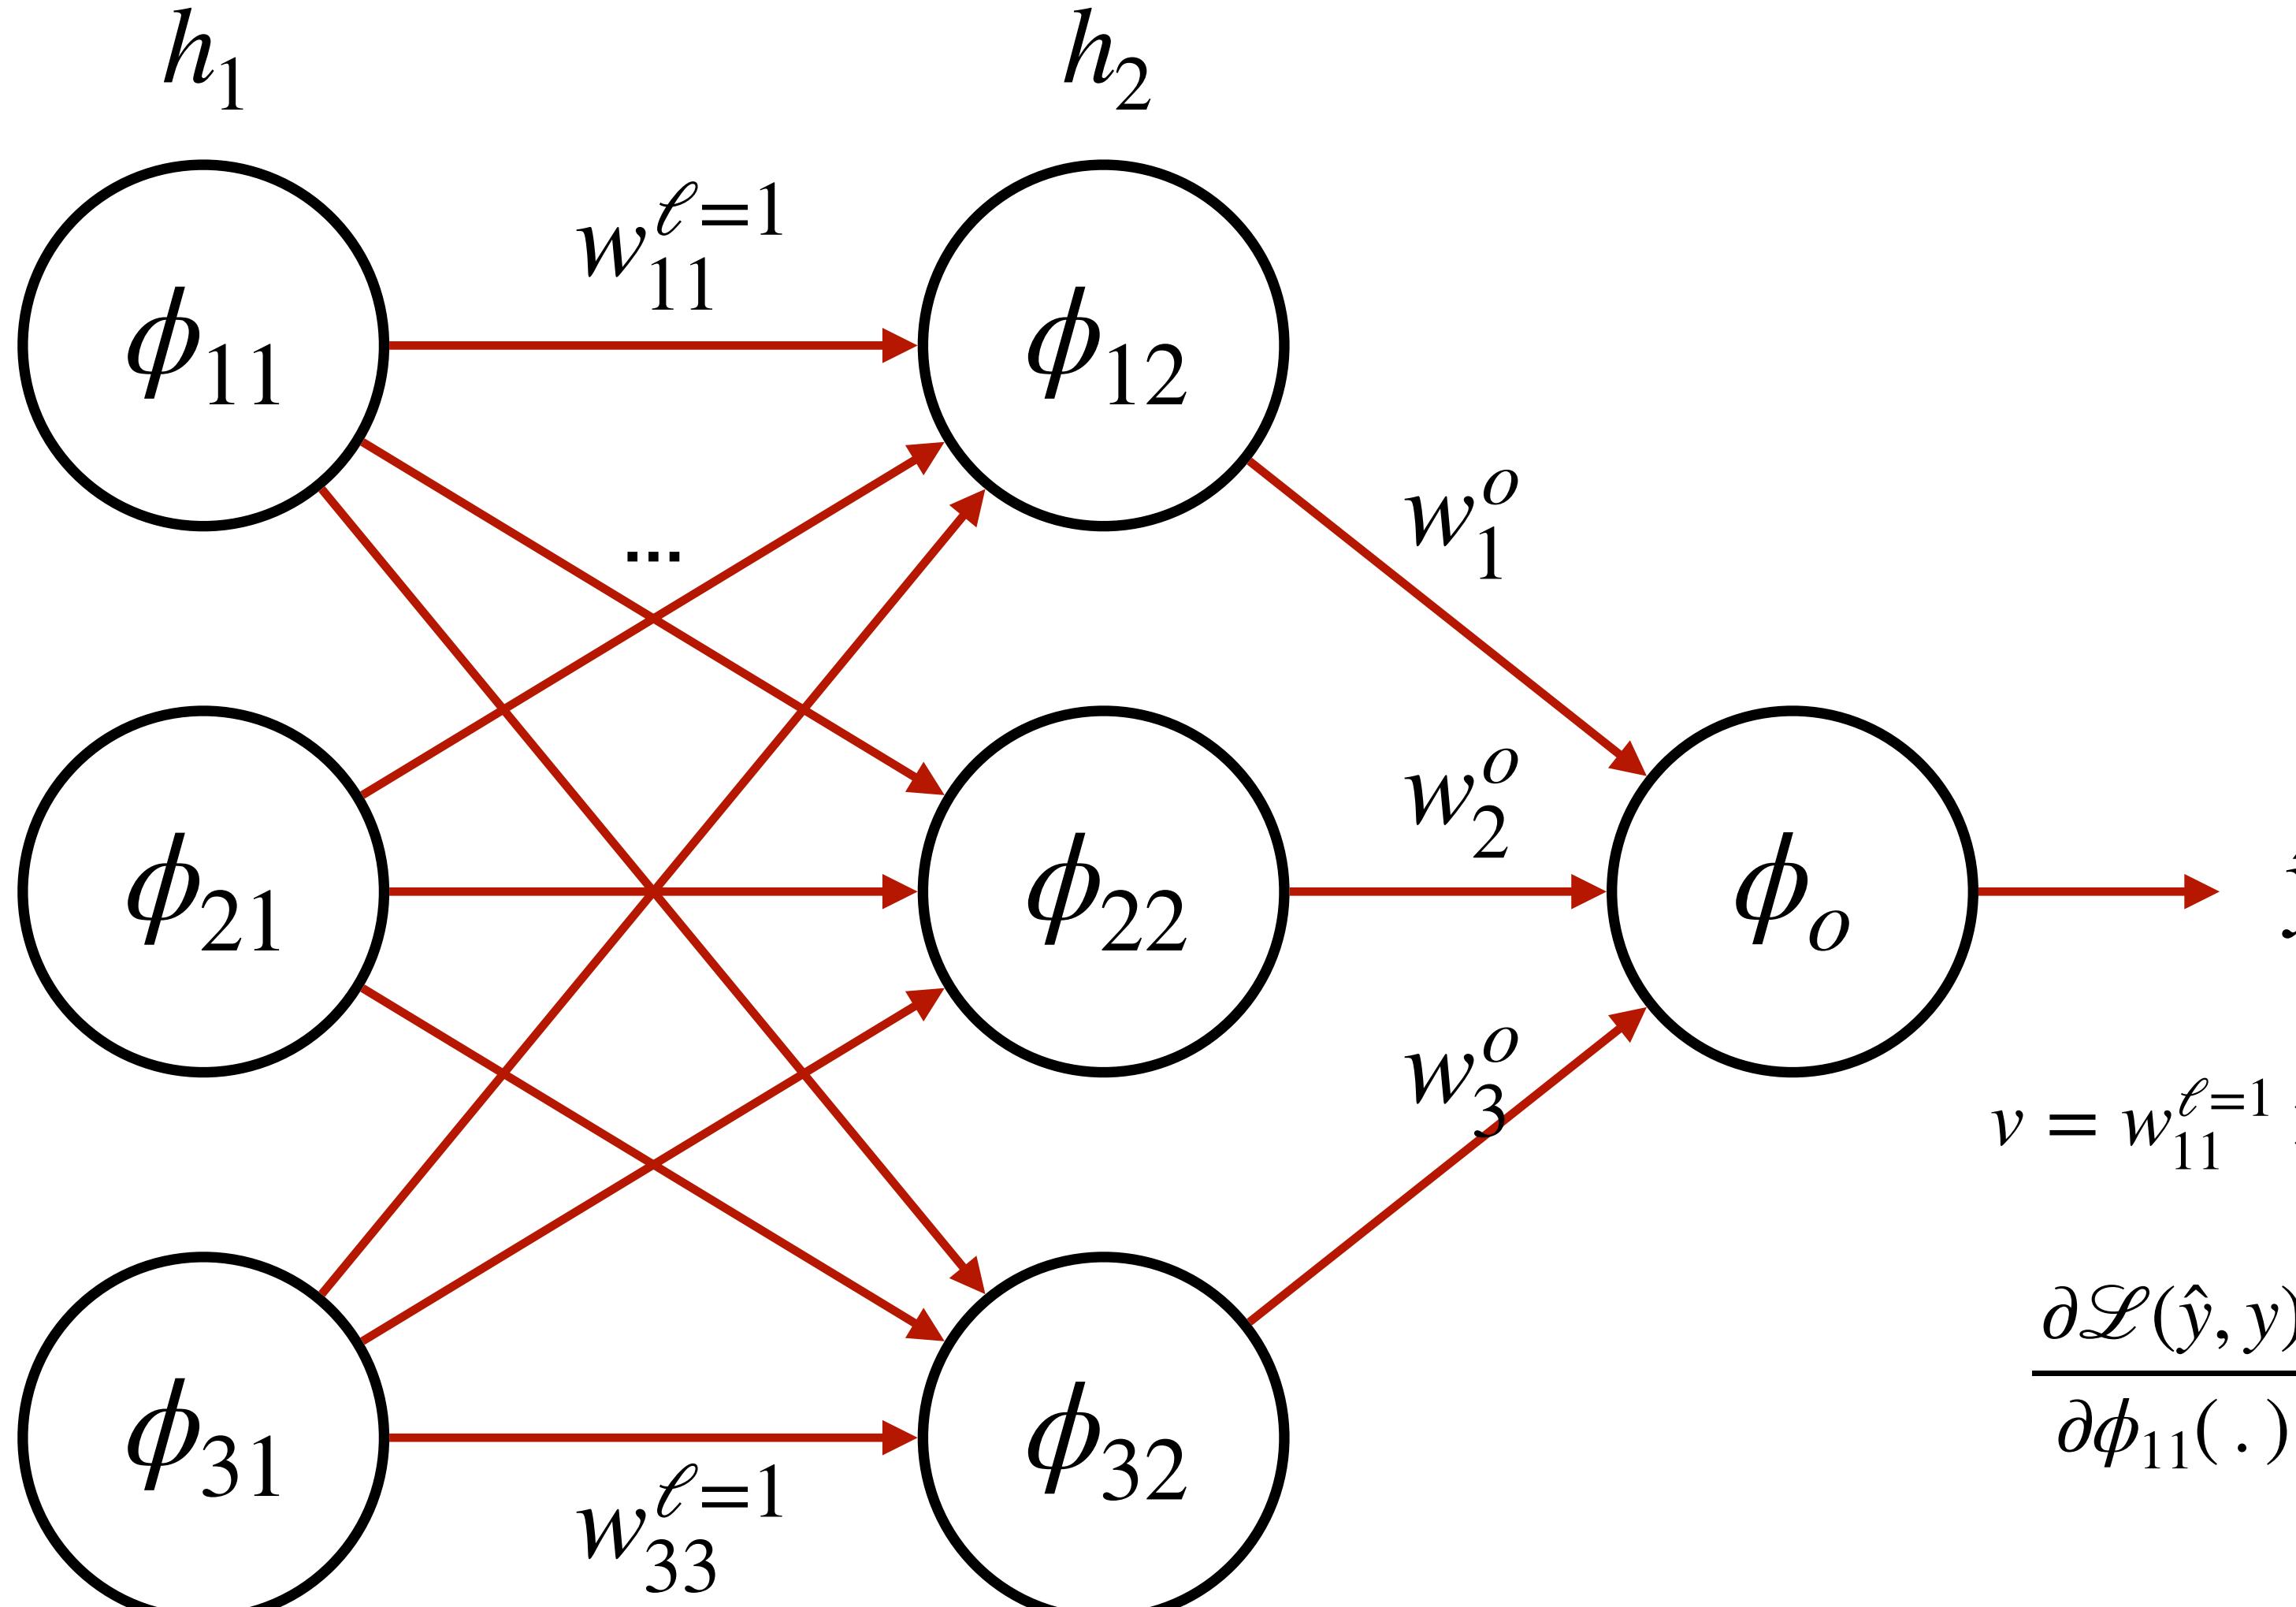
$$\frac{\partial \mathcal{L}(\hat{y}, y)}{\partial \phi_{22}(.)} = \frac{\partial \mathcal{L}(\hat{y}, y)}{\partial \hat{y}} \frac{\partial \hat{y}}{\partial u} \frac{\partial u}{\partial \phi_{22}(.)}$$

$$= \frac{\partial \mathcal{L}(\hat{y}, y)}{\partial \hat{y}} \frac{\partial \phi_o(u)}{\partial u} w_2^o$$

Depends on label y

Depends on ϕ_o

Backpropagation Review: FFNs



$$\frac{\partial \mathcal{L}(\hat{y}, y)}{\partial \phi_{12}(\cdot)} = \frac{\partial \mathcal{L}(\hat{y}, y)}{\partial \hat{y}} \frac{\partial \hat{y}}{\partial u} \frac{\partial u}{\partial \phi_{12}(\cdot)}$$

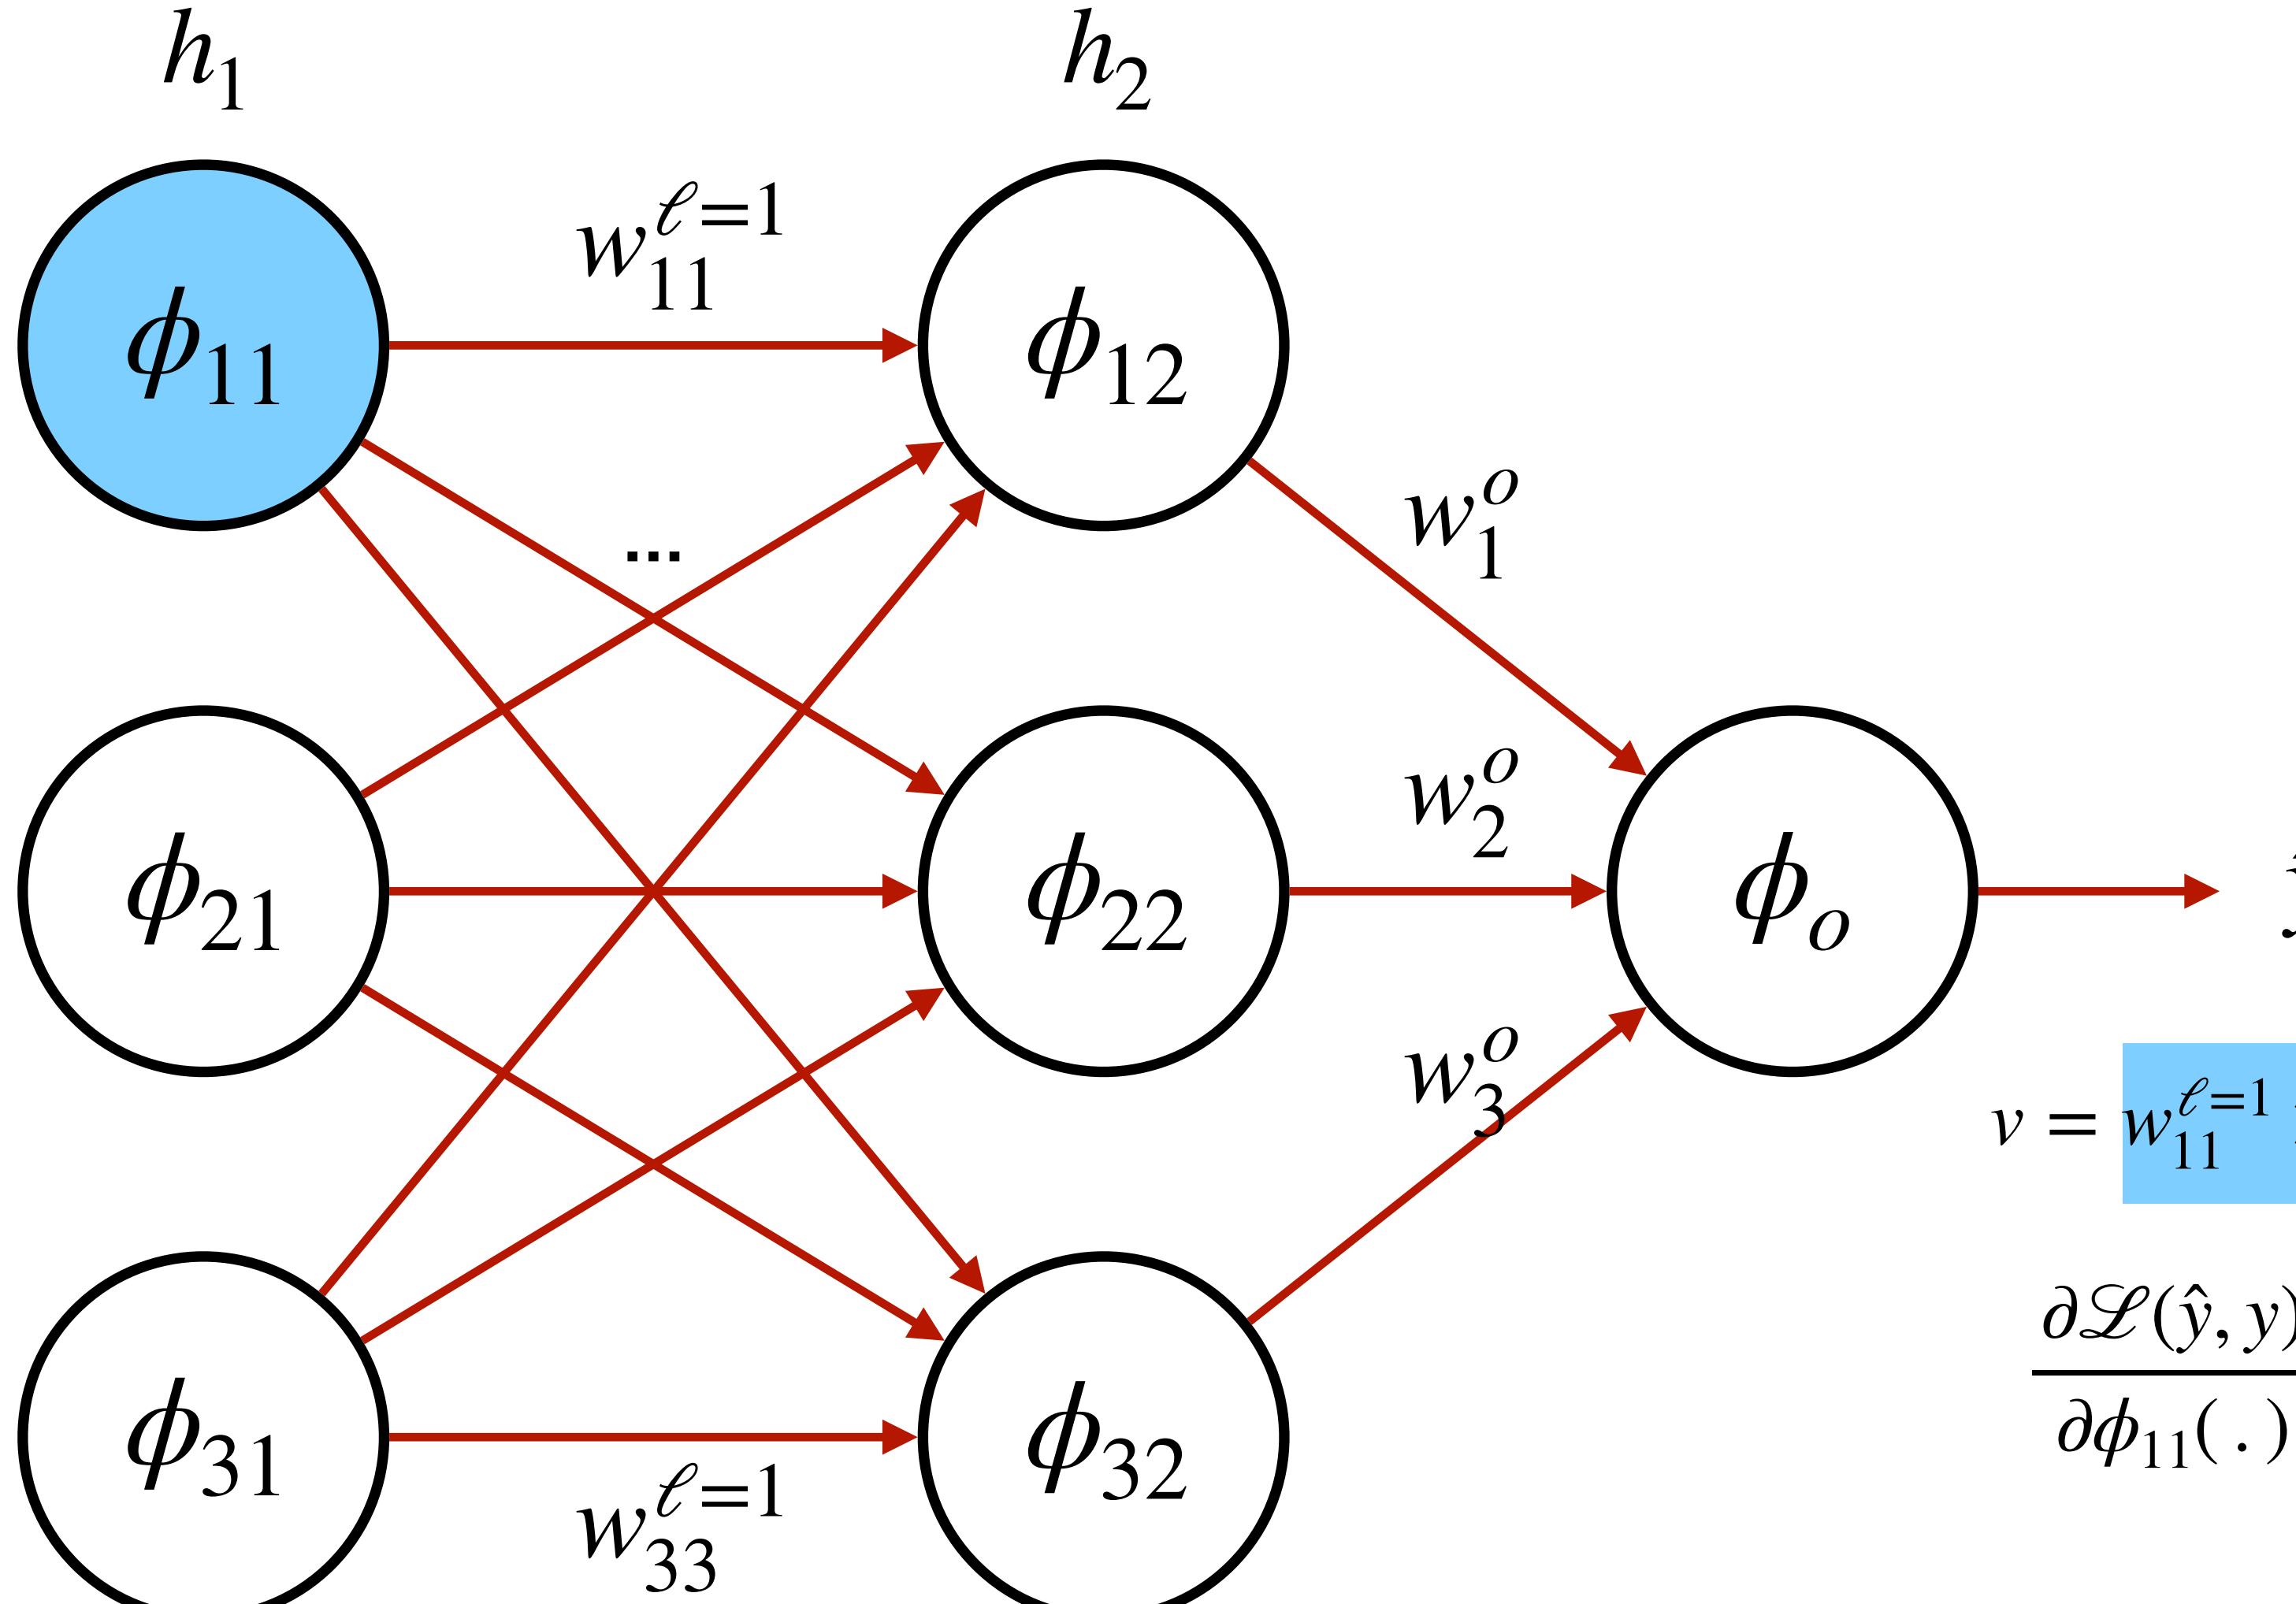
$$= \frac{\partial \mathcal{L}(\hat{y}, y)}{\partial \hat{y}} \frac{\partial \phi_o(u)}{\partial u} w_1^o$$

$$v = w_{11}^{\ell=1} \times \phi_{11}(\cdot) + w_{21}^{\ell=1} \times \phi_{21}(\cdot) + w_{31}^{\ell=1} \times \phi_{31}(\cdot)$$

$$\frac{\partial \mathcal{L}(\hat{y}, y)}{\partial \phi_{11}(\cdot)} = \frac{\partial \mathcal{L}(\hat{y}, y)}{\partial \hat{y}} \frac{\partial \hat{y}}{\partial u} \frac{\partial u}{\partial \phi_{12}(v)} \frac{\partial \phi_{12}(v)}{\partial v} \frac{\partial v}{\partial \phi_{11}(\cdot)}$$

$$= \frac{\partial \mathcal{L}(\hat{y}, y)}{\partial \hat{y}} \frac{\partial \phi_o(u)}{\partial u} w_1^o \frac{\partial \phi_{12}(v)}{\partial v} w_{11}^{\ell=1}$$

Backpropagation Review: FFNs



$$\frac{\partial \mathcal{L}(\hat{y}, y)}{\partial \phi_{12}(\cdot)} = \frac{\partial \mathcal{L}(\hat{y}, y)}{\partial \hat{y}} \frac{\partial \hat{y}}{\partial u} \frac{\partial u}{\partial \phi_{12}(\cdot)}$$

$$= \frac{\partial \mathcal{L}(\hat{y}, y)}{\partial \hat{y}} \frac{\partial \phi_o(u)}{\partial u} w_1^o$$

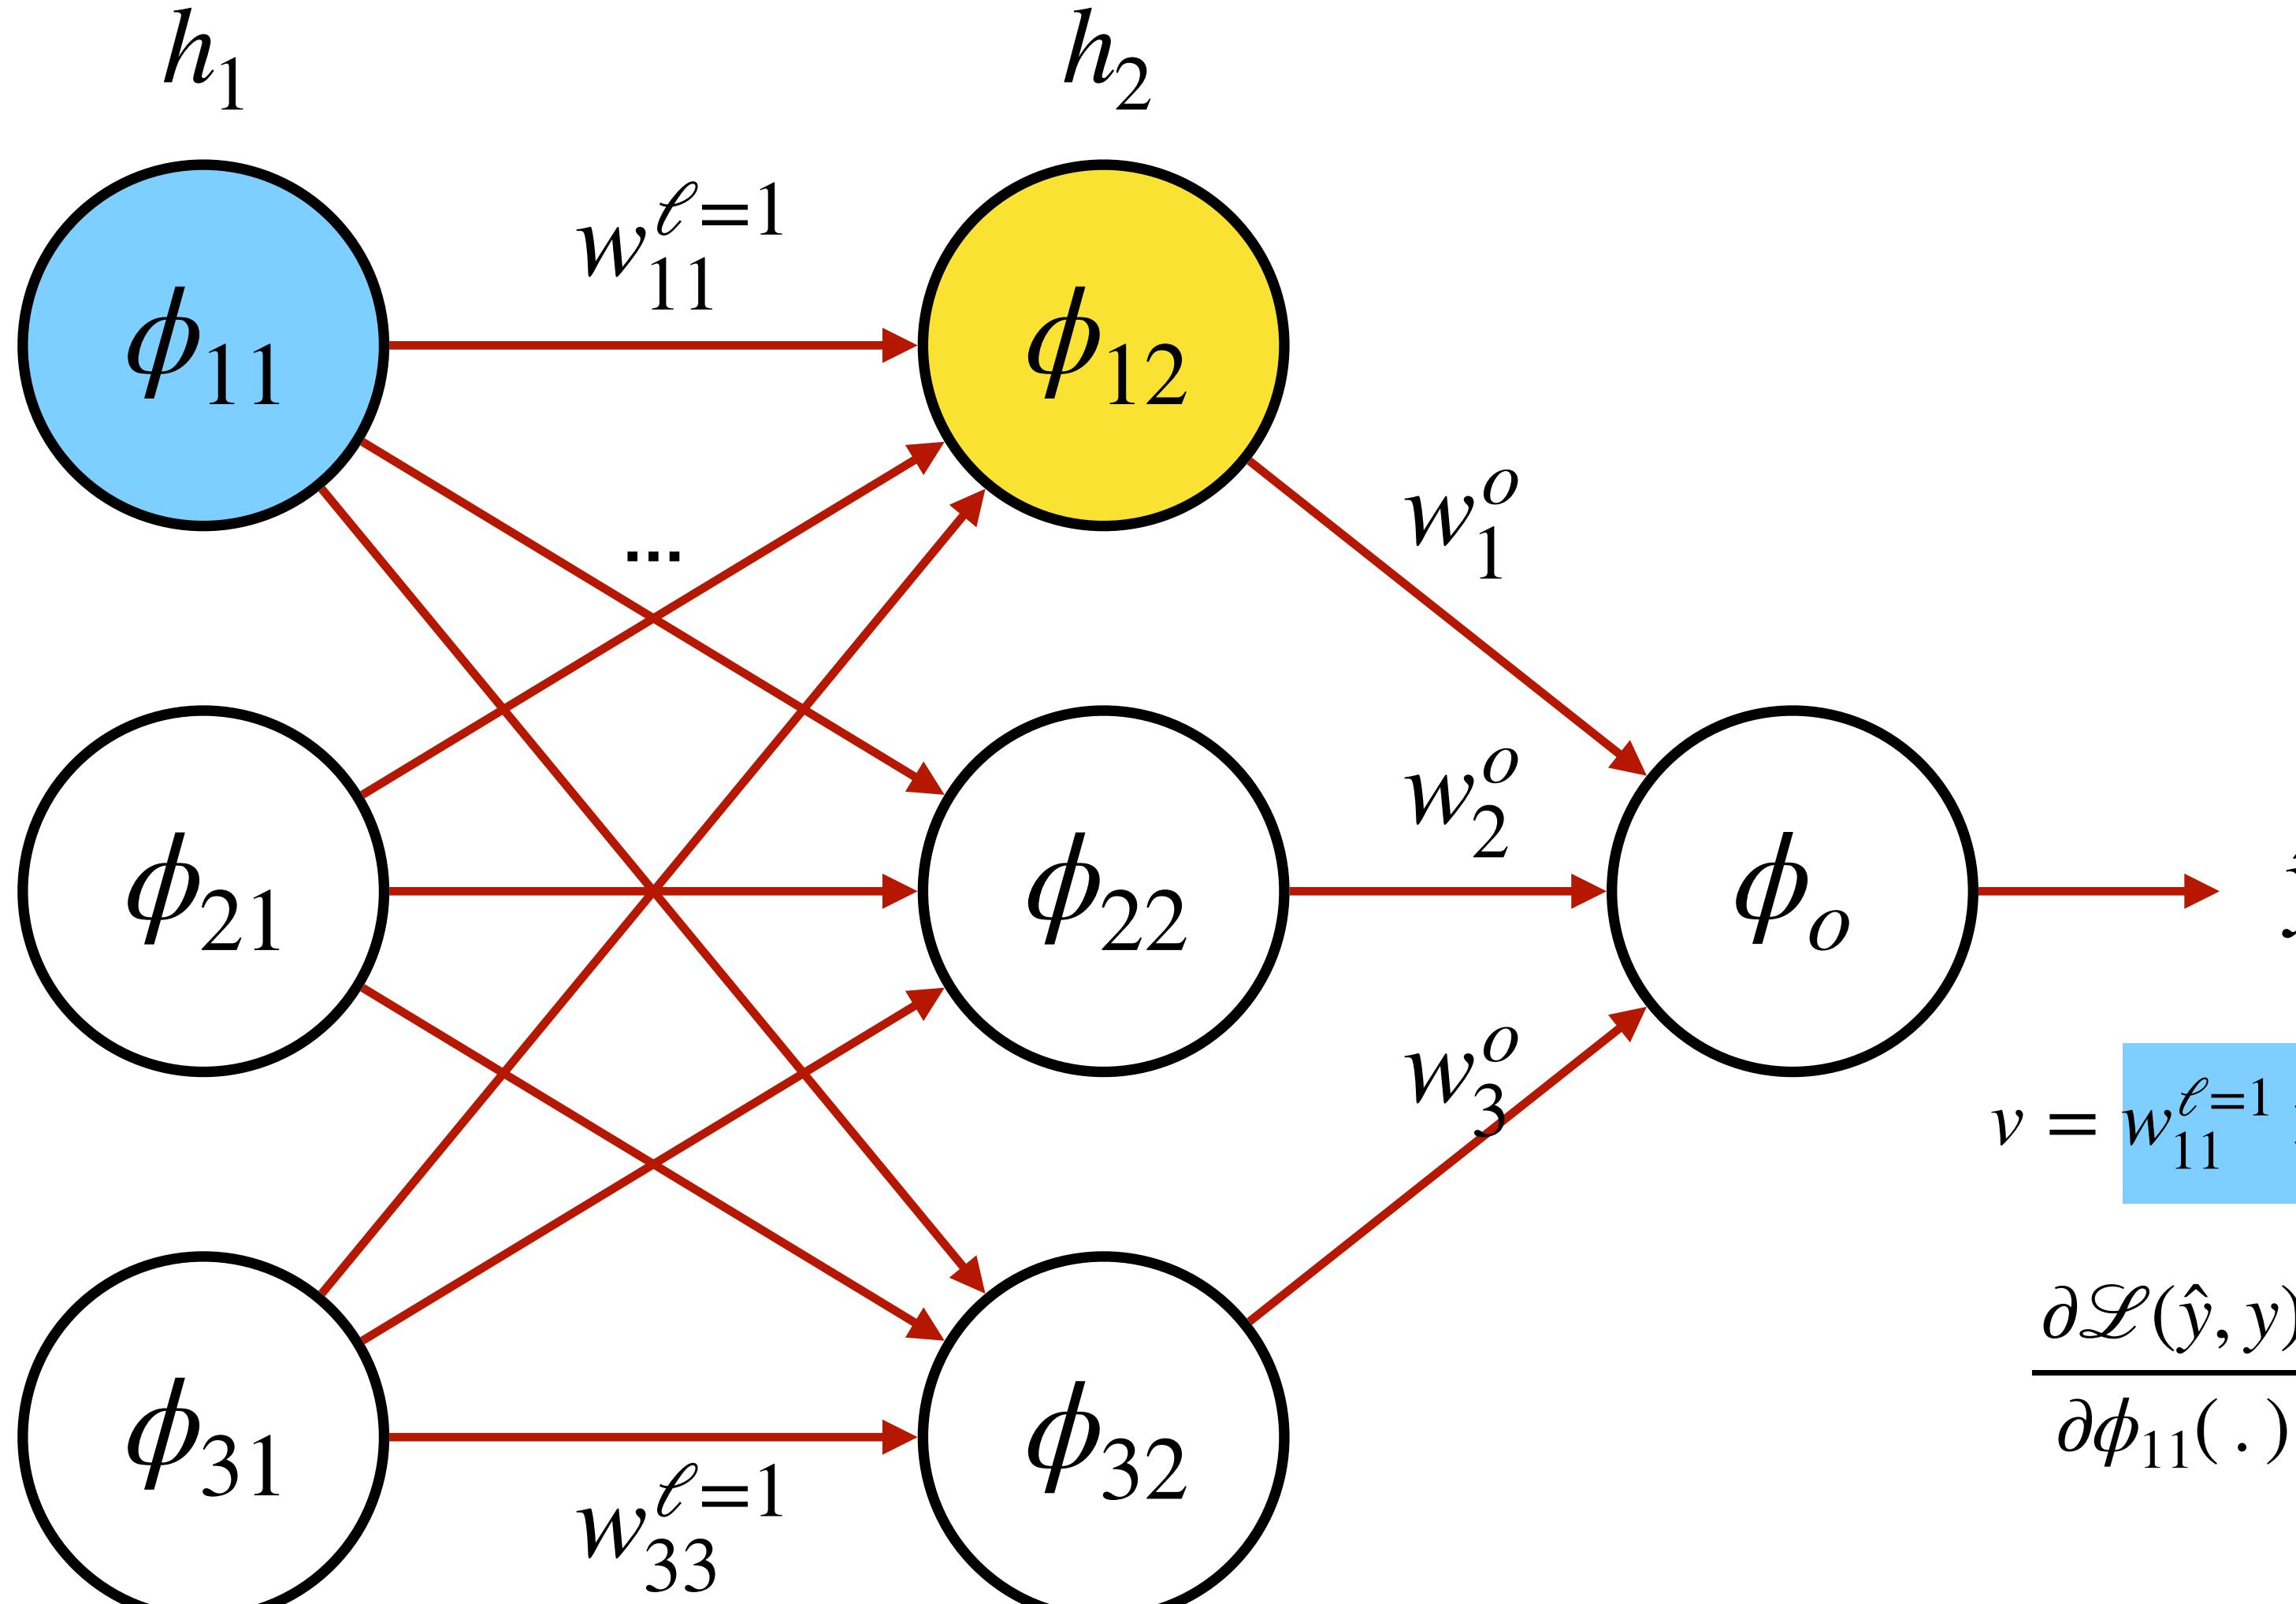
$$\hat{y}$$

$$v = w_{11}^{\ell=1} \times \phi_{11}(\cdot) + w_{21}^{\ell=1} \times \phi_{21}(\cdot) + w_{31}^{\ell=1} \times \phi_{31}(\cdot)$$

$$\frac{\partial \mathcal{L}(\hat{y}, y)}{\partial \phi_{11}(\cdot)} = \frac{\partial \mathcal{L}(\hat{y}, y)}{\partial \hat{y}} \frac{\partial \hat{y}}{\partial u} \frac{\partial u}{\partial \phi_{12}(v)} \frac{\partial \phi_{12}(v)}{\partial v} \frac{\partial v}{\partial \phi_{11}(\cdot)}$$

$$= \frac{\partial \mathcal{L}(\hat{y}, y)}{\partial \hat{y}} \frac{\partial \phi_o(u)}{\partial u} w_1^o \frac{\partial \phi_{12}(v)}{\partial v} w_{11}^{\ell=1}$$

Backpropagation Review: FFNs



$$\frac{\partial \mathcal{L}(\hat{y}, y)}{\partial \phi_{12}(.)} = \frac{\partial \mathcal{L}(\hat{y}, y)}{\partial \hat{y}} \frac{\partial \hat{y}}{\partial u} \frac{\partial u}{\partial \phi_{12}(.)}$$

$$= \frac{\partial \mathcal{L}(\hat{y}, y)}{\partial \hat{y}} \frac{\partial \phi_o(u)}{\partial u} w_1^o$$

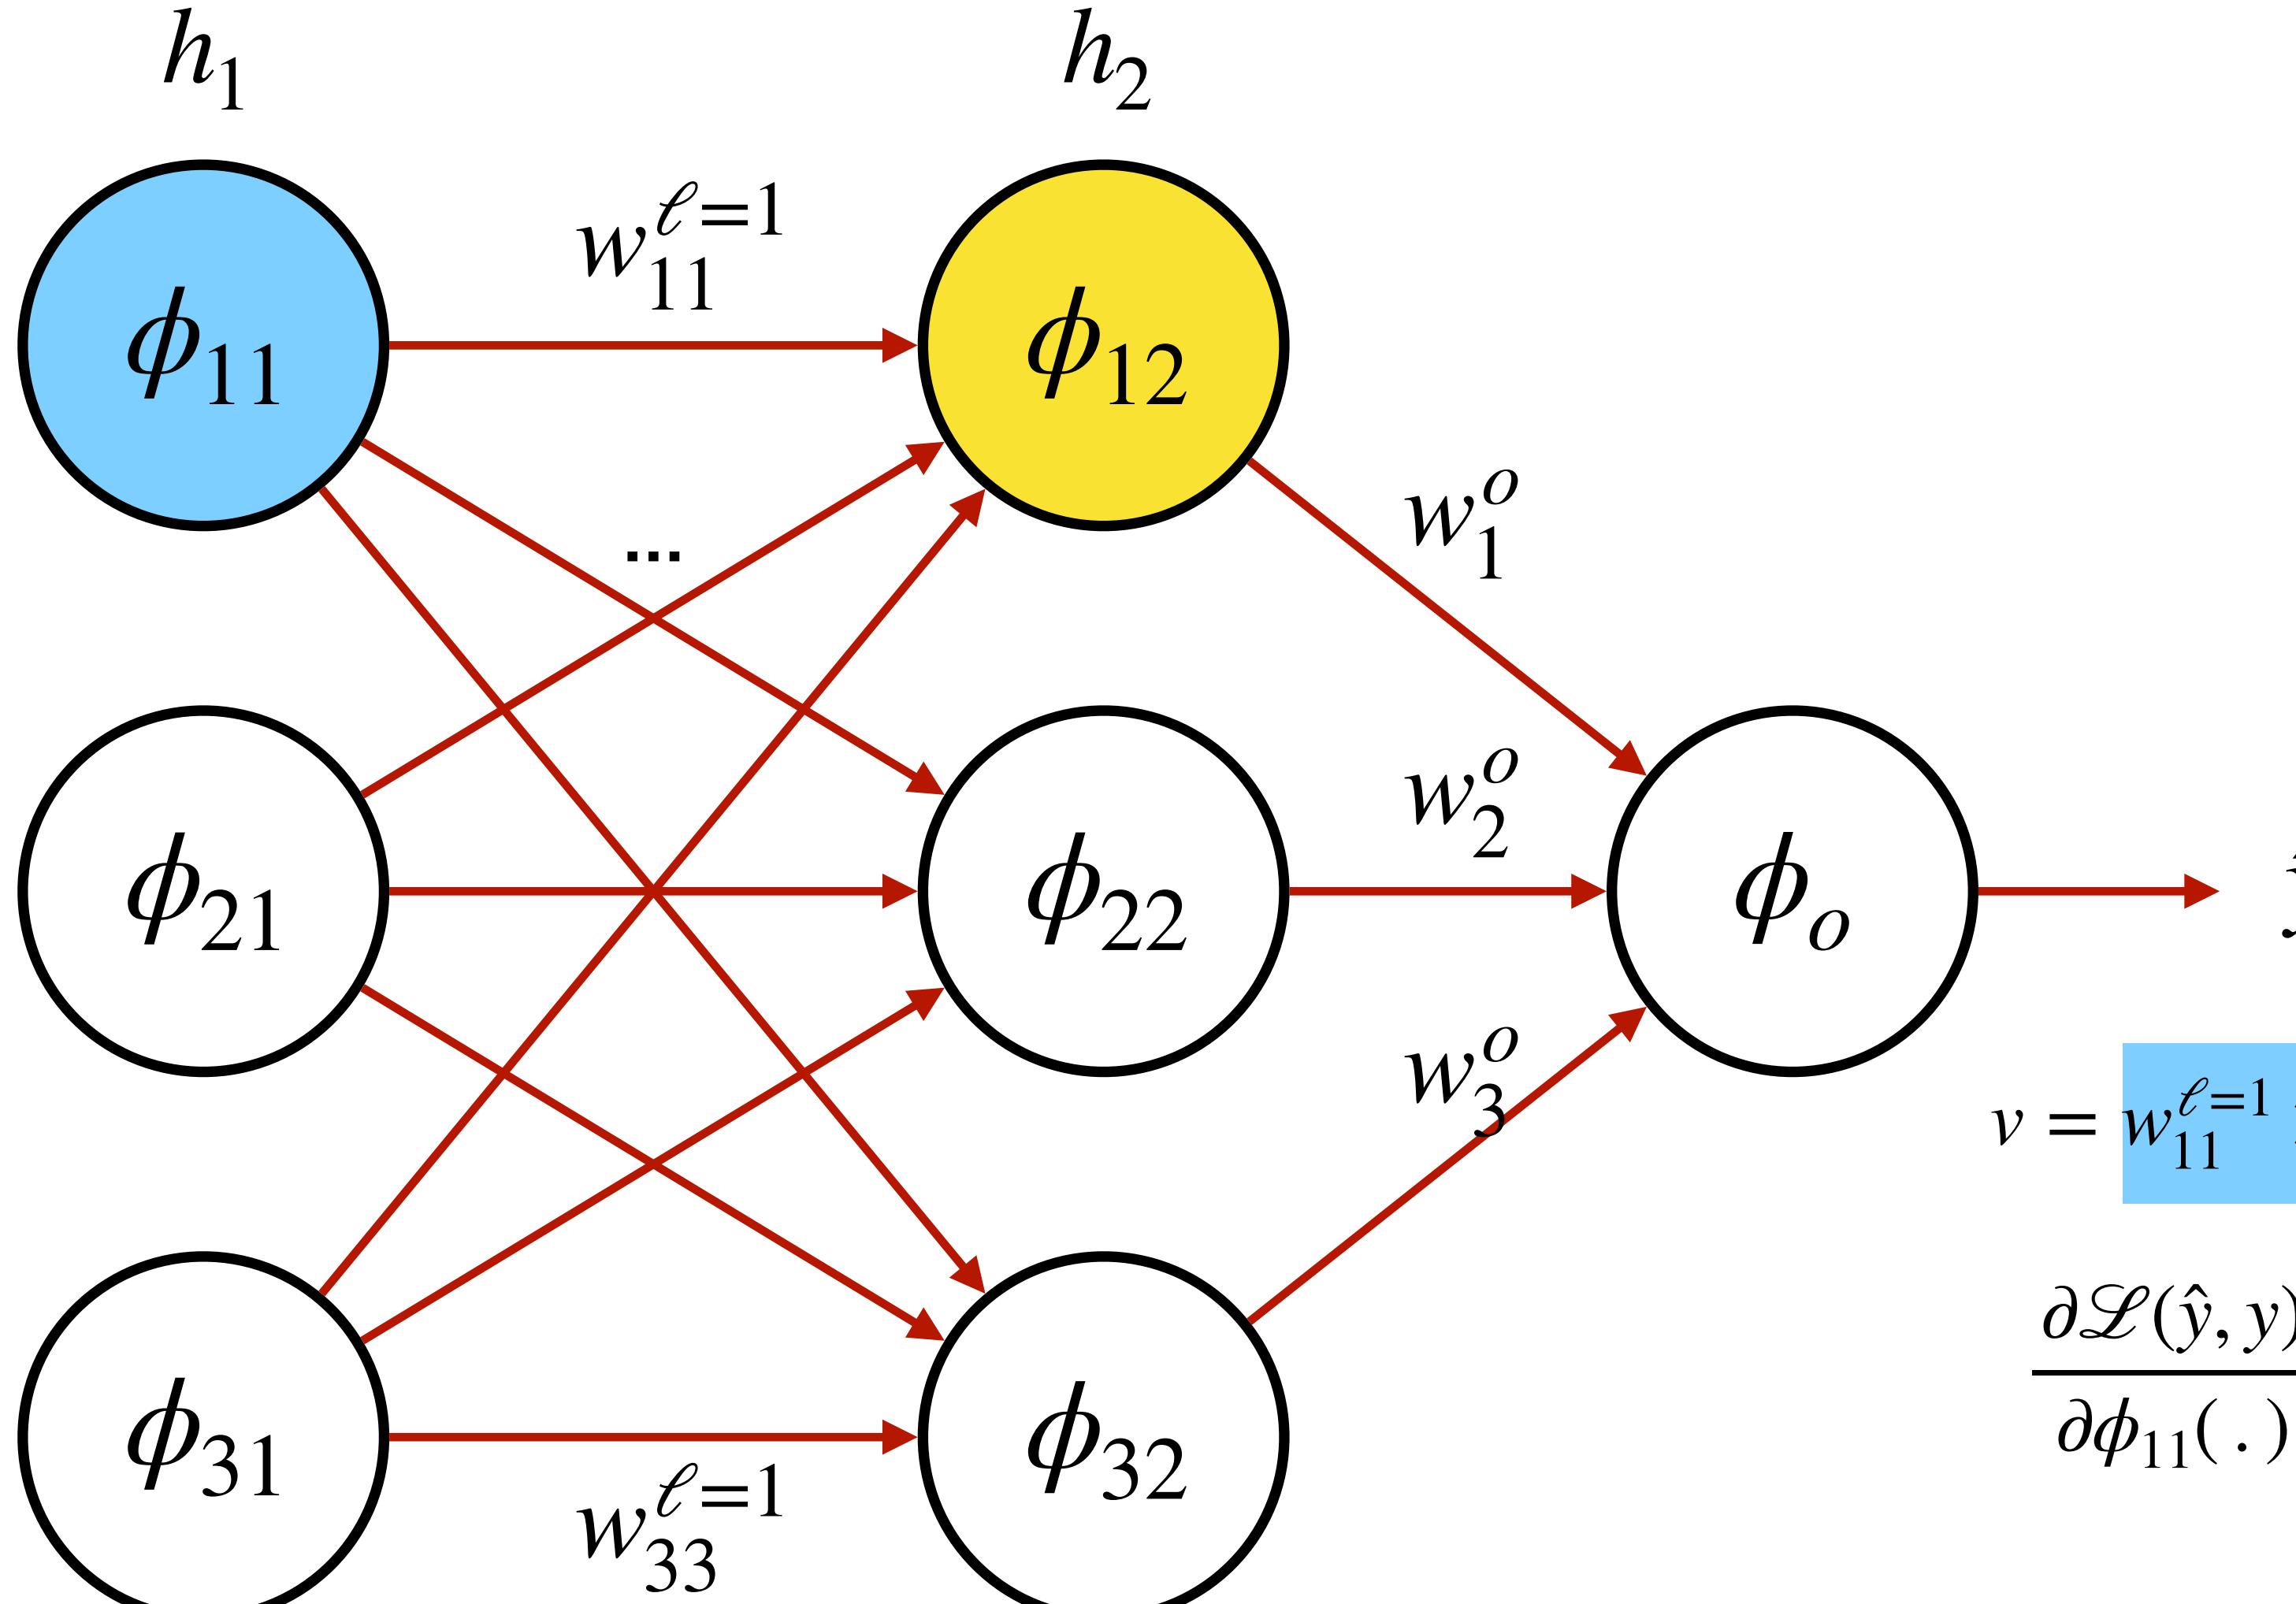
Depends on ϕ_{12}

$$v = w_{11}^{\ell=1} \times \phi_{11}(.) + w_{21}^{\ell=1} \times \phi_{21}(.) + w_{31}^{\ell=1} \times \phi_{31}(.)$$

$$\frac{\partial \mathcal{L}(\hat{y}, y)}{\partial \phi_{11}(.)} = \frac{\partial \mathcal{L}(\hat{y}, y)}{\partial \hat{y}} \frac{\partial \hat{y}}{\partial u} \frac{\partial u}{\partial \phi_{12}(v)} \frac{\partial \phi_{12}(v)}{\partial v} \frac{\partial v}{\partial \phi_{11}(.)}$$

$$= \frac{\partial \mathcal{L}(\hat{y}, y)}{\partial \hat{y}} \frac{\partial \phi_o(u)}{\partial u} w_1^o \frac{\partial \phi_{12}(v)}{\partial v} \frac{\partial v}{\partial \phi_{11}(.)} w_{11}^{\ell=1}$$

Backpropagation Review: FFNs



$$\frac{\partial \mathcal{L}(\hat{y}, y)}{\partial \phi_{12}(\cdot)} = \frac{\partial \mathcal{L}(\hat{y}, y)}{\partial \hat{y}} \frac{\partial \hat{y}}{\partial u} \frac{\partial u}{\partial \phi_{12}(\cdot)}$$

$$= \frac{\partial \mathcal{L}(\hat{y}, y)}{\partial \hat{y}} \frac{\partial \phi_o(u)}{\partial u} w_1^o$$

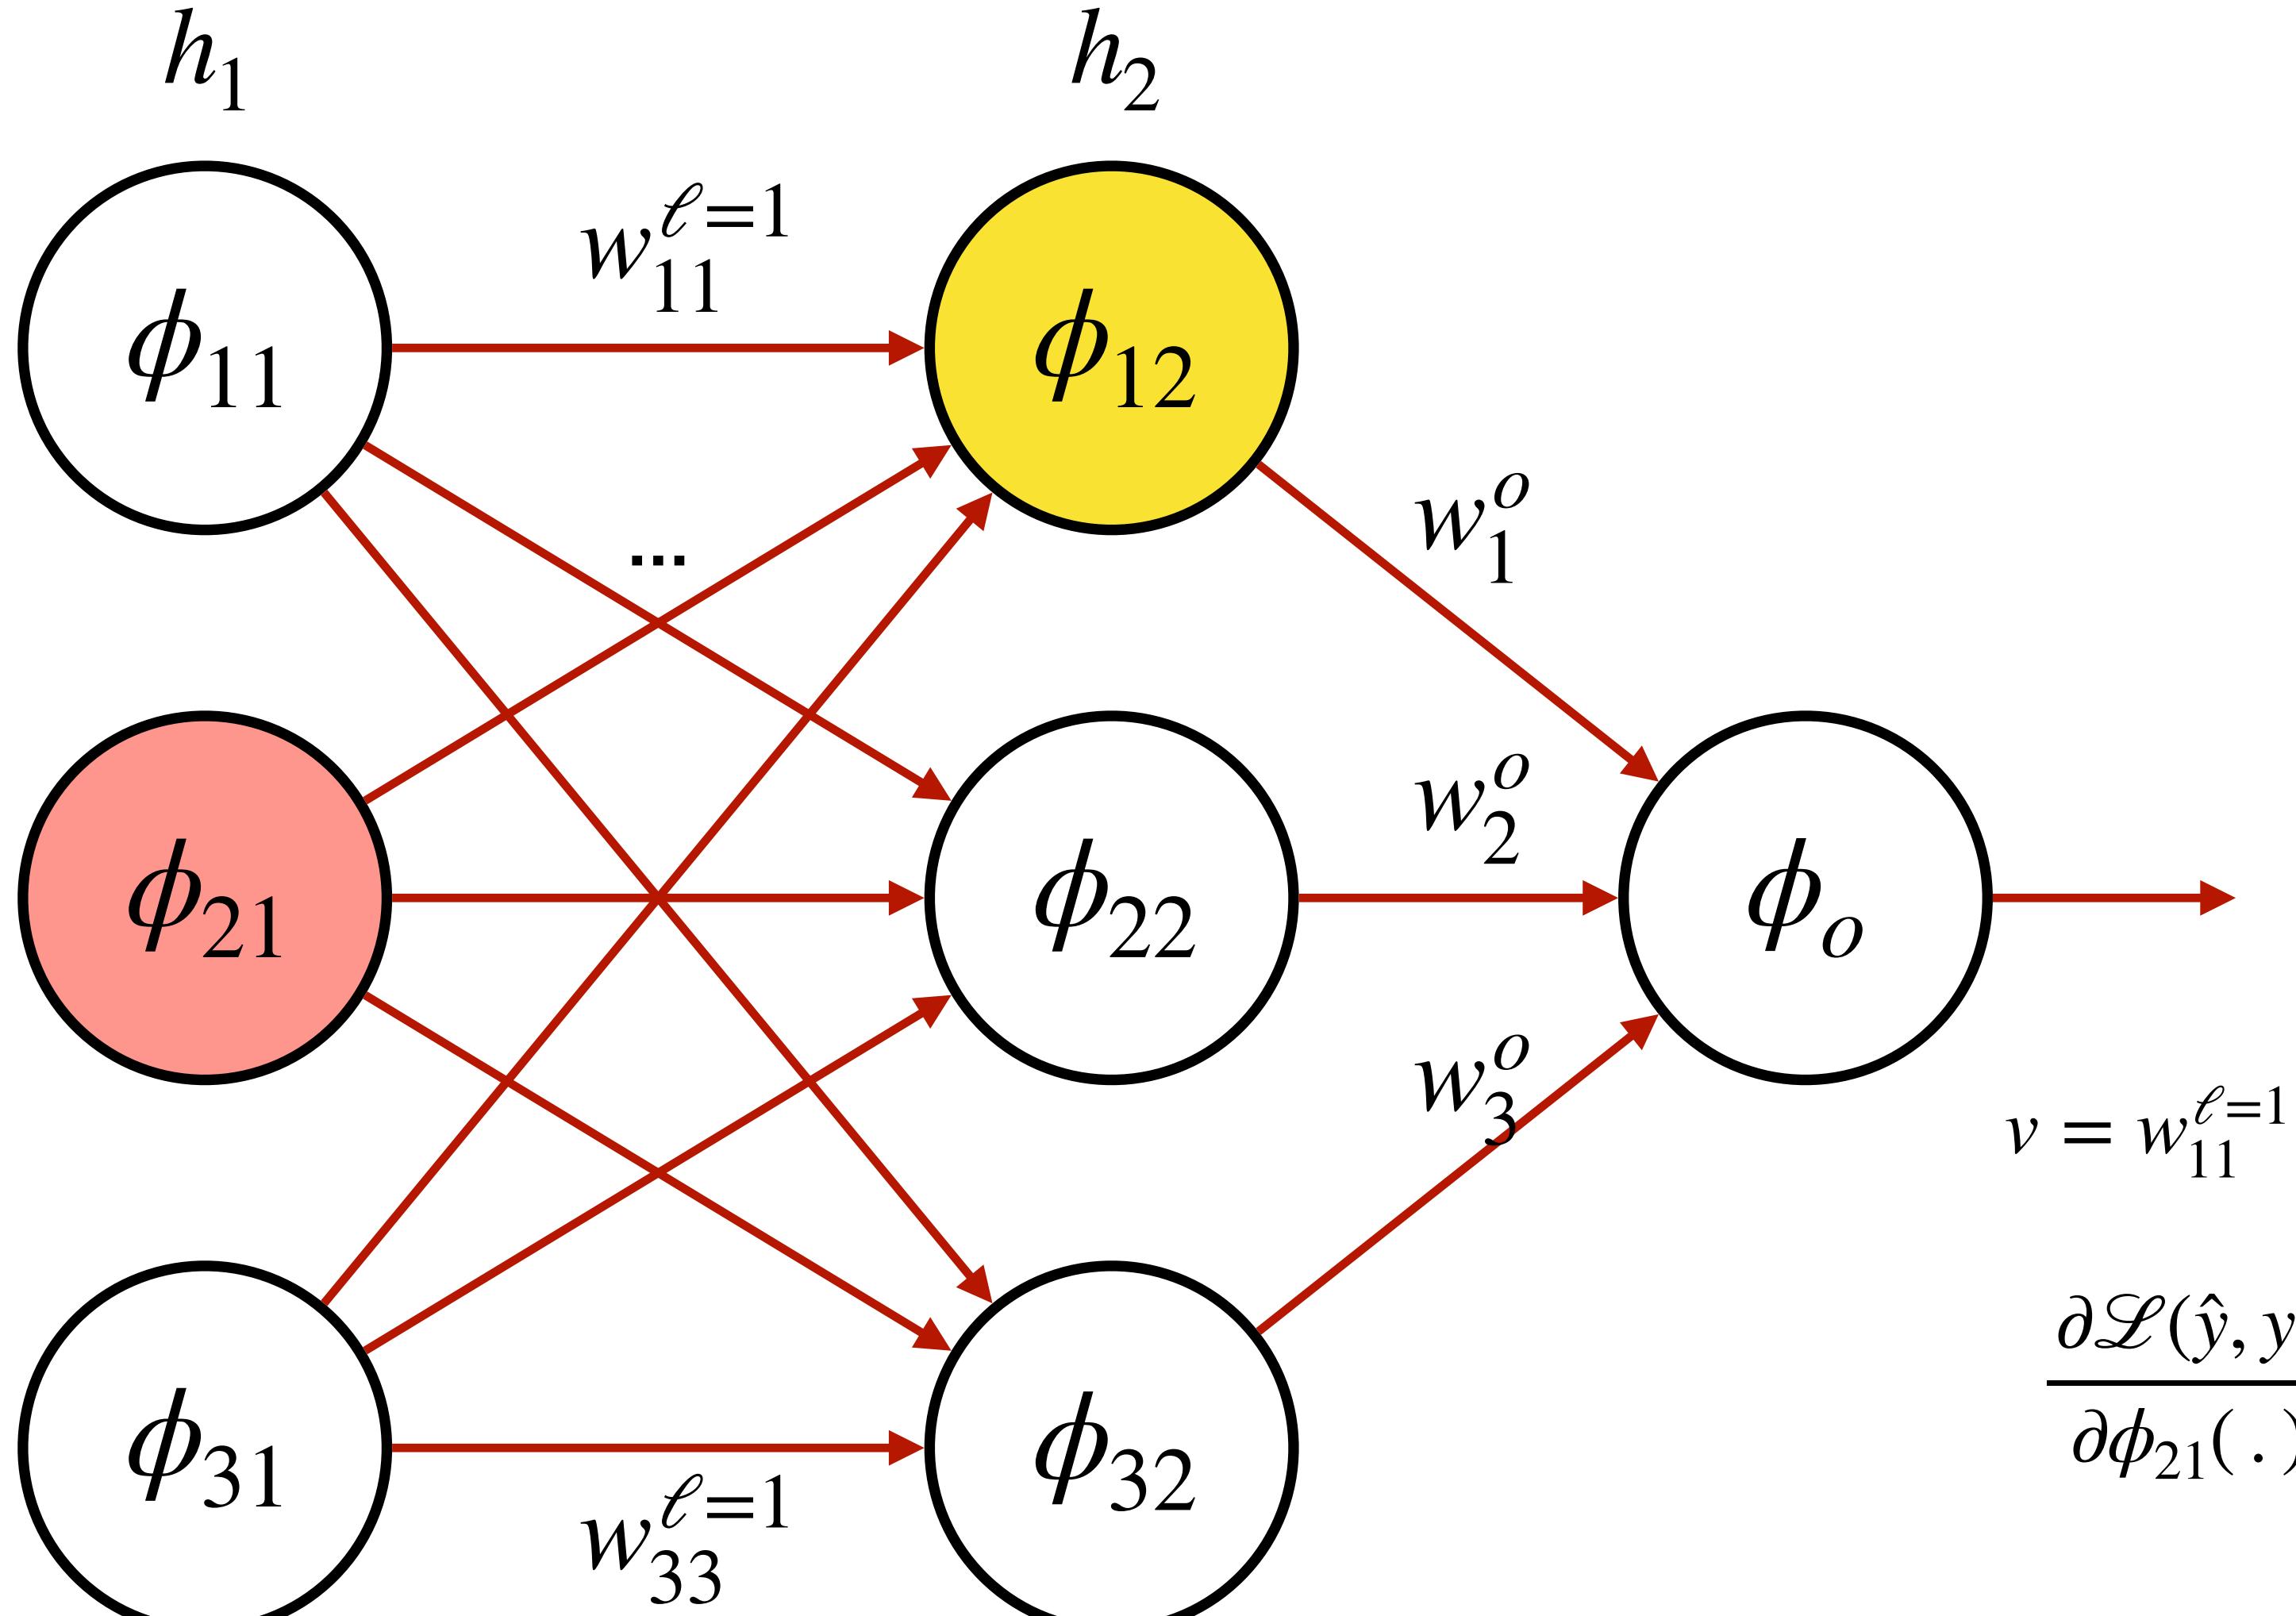
Depends on ϕ_{12}

$$v = w_{11}^{\ell=1} \times \phi_{11}(\cdot) + w_{21}^{\ell=1} \times \phi_{21}(\cdot) + w_{31}^{\ell=1} \times \phi_{31}(\cdot)$$

$$\frac{\partial \mathcal{L}(\hat{y}, y)}{\partial \phi_{11}(\cdot)} = \frac{\partial \mathcal{L}(\hat{y}, y)}{\partial \hat{y}} \frac{\partial \hat{y}}{\partial u} \frac{\partial u}{\partial \phi_{12}(v)} \frac{\partial \phi_{12}(v)}{\partial v} \frac{\partial v}{\partial \phi_{11}(\cdot)}$$

$$= \frac{\partial \mathcal{L}(\hat{y}, y)}{\partial \hat{y}} \frac{\partial \phi_o(u)}{\partial u} w_1^o \frac{\partial \phi_{12}(v)}{\partial v} w_{11}^{\ell=1}$$

Backpropagation Review: FFNs



$$\frac{\partial \mathcal{L}(\hat{y}, y)}{\partial \phi_{12}(\cdot)} = \frac{\partial \mathcal{L}(\hat{y}, y)}{\partial \hat{y}} \frac{\partial \hat{y}}{\partial u} \frac{\partial u}{\partial \phi_{12}(\cdot)}$$

$$= \frac{\partial \mathcal{L}(\hat{y}, y)}{\partial \hat{y}} \frac{\partial \phi_o(u)}{\partial u} w_1^o$$

Depends on ϕ_{12}

$$v = w_{11}^{\ell=1} \times \phi_{11}(\cdot) + w_{21}^{\ell=1} \times \phi_{21}(\cdot) + w_{31}^{\ell=1} \times \phi_{31}(\cdot)$$

$$\frac{\partial \mathcal{L}(\hat{y}, y)}{\partial \phi_{21}(\cdot)} = \frac{\partial \mathcal{L}(\hat{y}, y)}{\partial \hat{y}} \frac{\partial \hat{y}}{\partial u} \frac{\partial u}{\partial \phi_{12}(v)} \frac{\partial \phi_{12}(v)}{\partial v} \frac{\partial v}{\partial \phi_{21}(\cdot)}$$

$$= \frac{\partial \mathcal{L}(\hat{y}, y)}{\partial \hat{y}} \frac{\partial \phi_o(u)}{\partial u} w_1^o \frac{\partial \phi_{12}(v)}{\partial v} w_{21}^{\ell=1}$$

Question

**How would we extend backpropagation
to a recurrent neural network?**

Let's find out next class!

Recap

- Neural language models allow us to *share information* among similar sequences by learning neural representations that similarly represent them
- **Problem:** Fixed context language models can only process a limited window of the word history at a time
- **Solution:** recurrent neural networks can **theoretically** learn to model an **unbounded context length**
- **Next Class:** Training recurrent neural networks, associated challenges and mitigations through gated recurrent networks